

# Relational Data Analysis With Hive

MSBA 6330 Prof Liu

# Relational Data Analysis With Hive

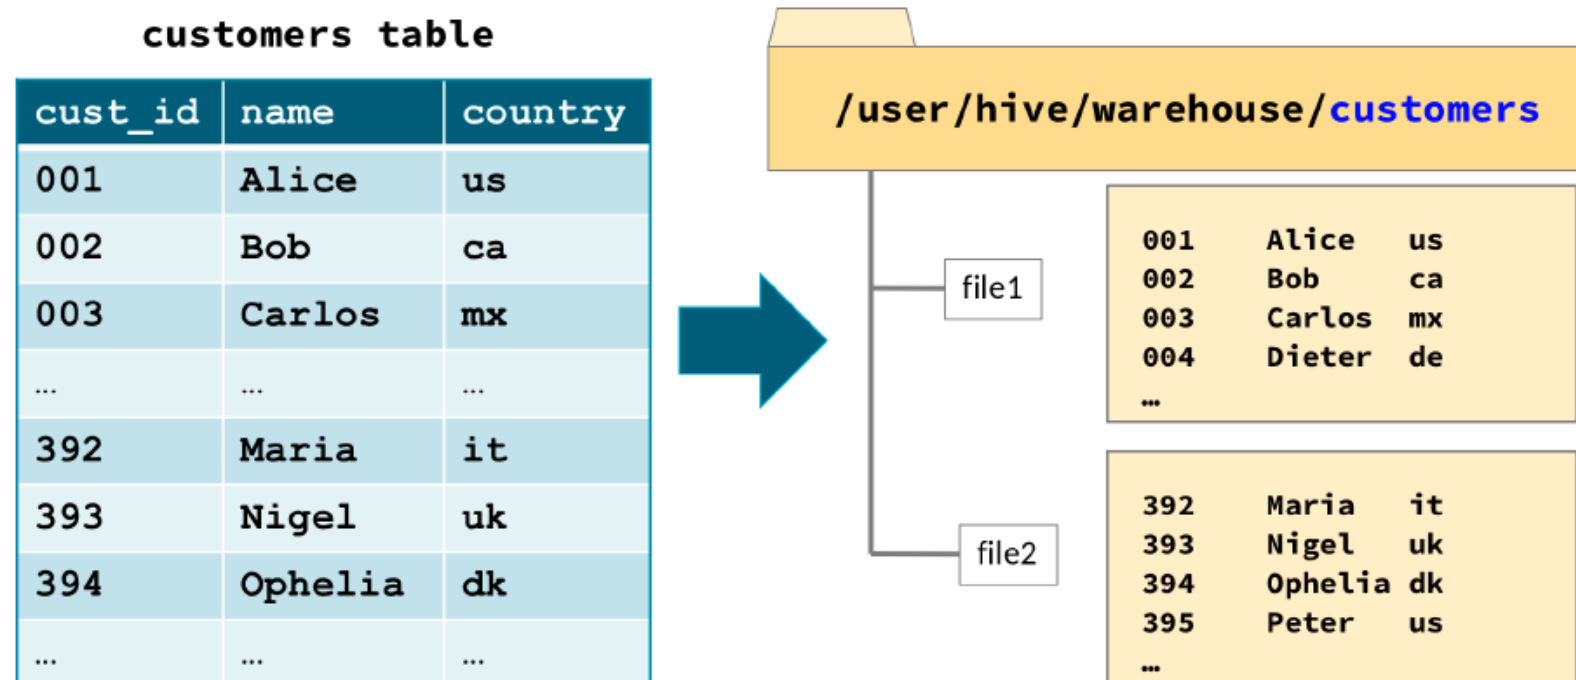
- In this chapter, you will learn
  - How to explore databases and tables in Hive
  - How HiveQL syntax compares to SQL
  - Which data types Hive supports
  - Which types of join operations Hive supports and how to use them

Relational Data Analysis With Hive

# **HIVE DATABASES AND TABLES**

# Hive Tables

- By default, Hive stores data for managed tables in the HDFS directory `/user/hive/warehouse`
  - Each table's data is stored in a subdirectory named after the table
  - A table's directory may contain multiple files
  - External tables can be stored elsewhere (including in the cloud).



# Hive Databases

- Each Hive table belongs to a specific database
  - If you don't specify a database, the table belongs to Hive's `default` database (not recommended, especially for large organizations)
- Please note that many small tables or lots of small partitions lead to small files in HDFS, which is not optimal.

# Exploring Hive Databases And Tables (1 Of 2)

- See which databases are available with the **SHOW DATABASES** command

```
SHOW DATABASES;  
accounting  
default
```

- Switch between databases with the **USE** command

```
SELECT * FROM customers;      -- customers in the default database  
USE dualcore;                 -- Switch databases.  
SELECT * FROM customers;      -- customers in dualcore  
SELECT * FROM sales.customers; -- customers in dualcore
```

All Hive keywords are case-insensitive, including the names of Hive operators and functions.

- See which tables the current database contains with the **SHOW TABLES** command

```
USE dualcore;  
SHOW TABLES;  
customers  
Employees ...
```

# Exploring Hive Databases And Tables (2 Of 2)

- See the basic structure for a table with the DESCRIBE command

```
> DESCRIBE orders;      -- Provide the fully qualified name
order_id
int cust_id
int order_date
timestamp
```

- **DESCRIBE FORMATTED** provides even more detailed information for those with advanced requirements

```
> DESCRIBE FORMATTED orders;
# col_name data_type          comment
order_id      int             None
cust_id       int             None
order_date    timestamp       None
# Detailed Table Information ...
```

Relational Data Analysis With Hive

# HIVEQL SYNTAX



# An Introduction To HiveQL

- HiveQL is Hive's query language
  - Based on a subset of SQL-92, plus Hive-specific extensions
- Some limitations compared to 'standard' SQL
  - Some features are not supported
    - e.g. Updating or deleting individual records (not available before Hive v0.14)
  - Others are only partially implemented
    - Include joins on non-equality conditions

```
JOIN ON customers.id = orders.id      -- This is supported
JOIN ON customers.id > orders.id      -- This is not supported
JOIN ON customers.id <> orders.id      -- Nor is this
```

# SQL support

- Semantics: Similar to MySQL
  - Select
  - Group by: Hive requires the group-by field to be among the selected fields.
  - Limit
  - Order by: Hive requires the order-by field to be among the selected fields
  - Where
  - UNION [ALL]
- Windowing /analytics functions (0.11+):
  - lead/lag/first\_value/last\_value
  - over/window/partition by/cube/rollup
  - rank(), row\_number(), dense\_rank(), cume\_dist(), percent\_rank(), ntile()

# Hive functions

- Many functions are similar to MySQL ([complete list](#))

- **keyword/function/identifier names are not case sensitive.**

<code>cast(&lt;expr&gt; as &lt;type&gt;):cast('1' as int)</code>	<code>rlike(regex)</code> – regular expression like.
<code>length(s)</code>	<code>to_date(s)</code>
<code>concat(s1,s2,s3,...)</code>	<code>year(d)</code>
<code>concat_ws(separator,s1,s2,s3,...)</code>	<code>month(d)</code>
<code>substr(s, start, length)</code>	<code>day(d)</code>
<code>upper(s)/ucase(s),</code>	<code>from_unixtime(i)</code>
<code>trim(s), ltrim(s)</code>	<code>size(Map or Array)</code>
<code>regexp_replace(s, regex, replacement)</code>	<code>rand()</code>
<code>repeat(s,n)</code>	<code>round(d)</code>
<code>split(s, pattern)</code>	<code>floor(d)</code>
<code>instr(str, substr)</code>	<code>ceil(d)</code>

- But, **string comparisons are case-sensitive**

```
SELECT * FROM customers WHERE state
      IN ('CA', 'OR', 'WA', 'NV', 'AZ');
```

# Subqueries In Hive

- It supports subqueries in the FROM and WHERE clauses
  - SELECT ... FROM (subquery) name ...
  - SELECT ... FROM ... WHERE EXISTS (subquery)
  - SELECT ... FROM ... WHERE x IN (subquery)

```
SELECT prod_id, brand, name
FROM (SELECT *
      FROM products
      WHERE (price - cost) / price > 0.65
      ORDER BY price DESC
      LIMIT 10) high_profits -- Mandatory subquery name
WHERE price > 1000          -- Price in cents
ORDER BY brand, name;
```

- Support for correlated subqueries is limited.
  - E.g. cannot be used in aggregations or conditional statements.

Relational Data Analysis With Hive

# DATA TYPES

# Hive Data Types

- Hive supports more than a dozen types
  - Most are similar to ones found in relational databases
  - Hive also supports three complex types
- Use the DESCRIBE command to see a table's column types

```
> DESCRIBE products;  
prod_id      int  
brand        string  
name         string  
price        int  
cost         int  
shipping_wt  int
```

# Hive Integer Types

- Integer types are appropriate for whole (signed) numbers
  - Both positive and negative values allowed

Name	Size	Range	Example
TINYINT	1 Byte	-128 - 127	17
SMALLINT	2 Bytes	-32,768 - 32,767	5842
INT	4 Bytes	-2,147,483,648 - 2,147,483,647	84127213
BIGINT	8 Bytes	~-9.2 quintillion - ~ 9.2 quintillion	632197432180964

- The default type for literal values is INT
- Best Practice:
  - Use the smallest type capable of doing the job

# Hive Decimal Types

- Float/double for floating point numbers
  - Caution: Avoid using when exact values are required!
    - So a float value entered as 3.1 might actually be stored as 3.10000000000000012
- Decimal for precise decimal numbers (e.g. money)

Name	Description	Example
FLOAT	Decimals	3.14159
DOUBLE	Very precise decimals	3.14159265358979323846
DECIMAL(p,s)	Controls scale/precision of a number	100.45 (p=5, s=2)



## Other Simple (Scalar) Types In Hive

- Hive can also store several other types of information

Name	Description	Example
STRING	Character sequence	Betty F. Smith
CHAR(n)	Fixed-length character sequence	Hive _ _ (n=6)
VARCHAR(n)	Variable length character sequence	Hive (n=10)
BOOLEAN	True or False	TRUE
TIMESTAMP	Instant in time (UTC)	2013-06-14 16:51:05
BINARY	Raw bytes (Like VARBINARY in SQL)	N/A

# Complex column types in Hive

- Hive also has a few complex data types
  - These are capable of holding multiple values

Name	Description & how to Define	Stored Data (suppose \$ is the collection item delimiter)	Access members
ARRAY	Ordered list of values, all of the same type, e.g. <code>departments array&lt;string&gt;</code>	<code>finance\$marketing\$hr</code>	<code>departments[0]</code>
MAP	Key/value pairs, each of the same type e.g. <code>prices map&lt;string,int&gt;</code>	<code>shoe#50\$shirt#75</code>	<code>prices['shirt']</code>
STRUCT	Named fields, of possibly mixed types e.g. <code>addr struct&lt;city:string, state:string, zip:int&gt;</code>	<code>Minneapolis\$MN\$55455</code>	<code>addr.city</code>

- Complex data types violate the “normal form”, but offer fast data access**
  - They are often desirable in Hadoop/Hive because they eliminate the need for big joins

Relational Data Analysis With Hive

# JOINING DATASETS

# Joins In Hive

- Hive supports several types of joins
  - Inner joins
  - Outer joins (Left, Right, and Full)
  - CROSS joins (supported in Hive 0.10 and later)
  - Left semi joins
- Only equality conditions are allowed in joins (equi-joins)
  - Valid: `customers.cust_id = orders.cust_id`
  - Invalid: `customers.cust_id <> orders.cust_id`
- For best performance, **list the largest table last in your query**
  - `Small_table JOIN big_table`

# Join Syntax

- Hive requires the following syntax for joins

```
SELECT c.cust_id, name, total  
FROM customers c  
JOIN orders o ON (c.cust_id = o.cust_id);
```

- The above example is an inner join (the word “inner” is not required) which emits records only when the join key is found in both tables
- **Implicit inner join syntax is not supported in Hive**

```
SELECT c.cust_id, name, total  
FROM customers c, orders o  
WHERE (c.cust_id = o.cust_id);
```



# Left Outer Join Example

- “OUTER” is required for outer joins in Hive

- Customers Table

cust_id	name	country
a	Alice	us
b	Bob	ca
c	Carlos	mx
d	Dieter	dw

```
SELECT c.cust_id, name, total
FROM customers c
LEFT OUTER JOIN orders o
ON (c.cust_id = o.cust_id);
```

- Orders Table

order_id	cust_id	total
1	a	1539
2	c	1871
3	a	6532
4	b	1456
5	z	2137

cust_id	name	total
a	Alice	1539
a	Alice	6352
b	Bob	1456
c	Carlos	1871
d	Dieter	NULL

# Full Outer Join Example

## • Customers Table

cust_id	name	country
a	Alice	us
b	Bob	ca
c	Carlos	mx
d	Dieter	dw

## • Orders Table

order_id	cust_id	total
1	a	1539
2	c	1871
3	a	6532
4	b	1456
5	z	2137

## • Code

```
SELECT c.cust_id, name, total
FROM customers c
FULL OUTER JOIN orders o
ON (c.cust_id = o.cust_id);
```

## • Result

cust_id	name	total
a	Alice	1539
a	Alice	6352
b	Bob	1456
c	Carlos	1871
d	Dieter	NULL
NULL	NULL	2137

# Essential Points

- Every Hive table belongs to exactly one database
  - The SHOW DATABASES command lists databases
  - The USE command switches the active database
  - The SHOW TABLES command lists all tables in a database
- Every column in a Hive table has an associated data type
  - Most simple column types are similar to SQL
  - Hive also supports a few complex types
- HiveQL syntax is familiar to those who know SQL
  - A subset of SQL-92, plus Hive-specific extensions
  - Supports inner, outer, and Left semi joins
  - Many SQL functions are built into Hive