

Optimize Hive Query Performance

MSBA 6330 Prof Liu

Hive Optimization

- In this chapter, you will learn
 - Use faster engines for Hive
 - How to choose storage formats for Hive tables
 - How to partition tables to reduce amount of data read for a query
 - How to understand and write better performing Hive queries
 - Why/how to use bucketing

Optimize Hive Query Performance

USE FASTER PROCESSING ENGINES

Uses a Faster Execution Engine

- Hive uses MapReduce (MR) engine by default, but it can also leverage faster engines
 - Tez: `set hive.execution.engine=tez;`
 - Tez also offers a customizable execution architecture, permitting dynamic performance optimizations, and dramatically improving the speed of execution.
 - Spark: `set hive.execution.engine=spark;`
 - In memory computing, more customizable execution architecture.
 - Expect a long delay as Spark initializes after you submit the first query
 - Subsequent queries run without a delay
 - Hive on Spark requires more memory on cluster than Hive on MapReduce
- To see the current engine, use `SET hive.execution.engine;`

Optimize Hive Query Performance

USE FASTER STORAGE FORMATS

Choose a File Format

- Hive supports multiple storage formats

```
CREATE TABLE tablename (colname DATATYPE, ...)
ROW FORMAT DELIMITED FIELDS TERMINATED BY char
STORED AS format;
```

- Format Options
 - TEXTFILE
 - SEQUENCEFILE
 - AVRO
 - PARQUET
 - RCFILE
 - ORCFILE

Considerations for Choosing a File Format

- Hadoop and its ecosystem support many file formats
 - You can ingest data in one format and convert to another as needed
- Selecting the format for your dataset involves several considerations
 - Ingest pattern
 - Tool compatibility
 - Expected lifetime
 - Storage and performance requirements

TEXTFILE

- Text files are the most basic file type in Hadoop
 - Can be read or written from virtually any programming language
 - Comma- and tab-delimited files are compatible with many applications
- Text files are human-readable
 - All values are represented as strings
 - Useful when debugging
- At scale, this format is inefficient
 - Representing numeric values as strings wastes storage space
 - Difficult to represent binary data such as images
 - Often resort to techniques such as Base64 encoding
 - Conversion to/from native types adds performance penalty
- **Verdict: Good interoperability, but poor performance**

How many bytes to store
108125150 as text (9 bytes)
or integer (4 bytes)?

SequenceFile Format

- SequenceFiles store key-value pairs in a **binary** container format
 - Less verbose and more efficient than text files
 - Capable of storing binary data such as images
 - Format is **Java-specific** and tightly coupled to Hadoop
 - Use internally for temporal outputs of mappers
 - Support splittable compression
- **Verdict: Good performance, but poor interoperability**

SequenceFile File Layout

| | | | | | | | | |
|------|-----|-------|-----|-------|-----|-------|-----|-------|
| Data | Key | Value | Key | Value | Key | Value | Key | Value |
|------|-----|-------|-----|-------|-----|-------|-----|-------|

Apache Avro File Format



- Apache Avro is an efficient data serialization framework
- Avro also defines a **binary** data file format for storing Avro records
 - Similar to SequenceFile format
- Efficient storage due to optimized binary encoding
 - Support compression
- Widely supported throughout the Hadoop ecosystem
 - Can also be used outside of Hadoop
- Ideal for long-term storage of important data
 - Many languages can read and write Avro files
 - Embeds schema in the file, so will always be readable
 - Schema evolution can accommodate changes
- **Verdict: Excellent interoperability and performance**
 - **Best choice for general-purpose storage in Hadoop**

Columnar formats

- Organize data storage on disk by column, rather than by row
 - Very efficient when selecting only a subset of a table's column
 - Including Apache Parquet, RCFile, and ORCFile

| Organization of data in traditional row-based formats | | | |
|---|-------|-----------|------------|
| id | name | city | occupation |
| 1 | Alice | Palo Alto | Accountant |
| 2 | Bob | Sunnyvale | Accountant |
| 3 | Bob | Palo Alto | Dentist |
| 4 | Bob | Palo Alto | Manager |
| 5 | Carol | Palo Alto | Manager |

Row-based storage to disk

| | | | | | | |
|---|-------|-----------|------------|---|-----|--|
| 1 | Alice | Palo Alto | Accountant | 2 | Bob | |
|---|-------|-----------|------------|---|-----|--|

| Organization of data in columnar formats | | | |
|--|-------|-----------|------------|
| id | name | city | occupation |
| 1 | Alice | Palo Alto | Accountant |
| 2 | Bob | Sunnyvale | Accountant |
| 3 | Bob | Palo Alto | Dentist |
| 4 | Bob | Palo Alto | Manager |
| 5 | Carol | Palo Alto | Manager |

Column-based storage to disk

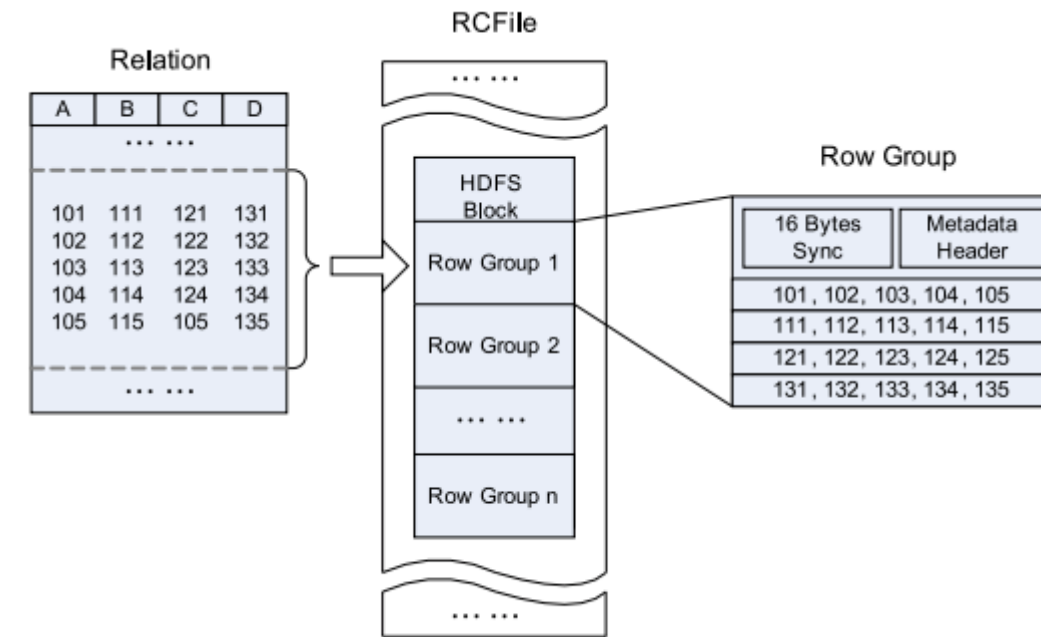
| | | | | | |
|-------|-----|------|-------|-----|-----------|
| Alice | Bob | (x3) | Carol | ... | Palo Alto |
|-------|-----|------|-------|-----|-----------|

Columnar File Format: Apache Parquet


- Apache Parquet is an open source columnar format
 - Originally developed by engineers at Cloudera and Twitter
 - Now an Apache Software Foundation project
 - Supported in MapReduce, Hive, Pig, Impala, Spark, and others
 - Schema is embedded in the file
 - Stores **binary-encoded** records
- Uses advanced optimizations described in [Google's Dremel paper](#)
 - Reduces storage space
 - Increases performance
- Most efficient when adding many records at once
 - Some optimizations rely on identifying repeated patterns
- **Verdict: Excellent interoperability and performance**
 - Good choice for column-based access patterns

Columnar File Formats: RCFile and ORCFile

- RCFile
 - Splits data horizontally into row groups, each row group saves data in a columnar format
 - All data stored as strings (inefficient)
 - **Verdict: Poor performance and limited interoperability**
- ORCFile – Optimized RCFile
 - An improved version of RCFile
 - Currently supported in Hive, Spark, but limited in Impala
 - **Verdict: Good performance but limited interoperability**



HortonWorks (2013) Study of Storage Format and File Sizes



Convert between formats

- Load data in the original format, then use INSERT INTO TABLE to convert it into a new format.
- Create a new table stored in Parquet format

```
CREATE TABLE order_details_parquet (  
  order_id INT,  
  prod_id INT)  
STORED AS PARQUET;
```

- Load data from another table into a Parquet table

```
INSERT OVERWRITE TABLE order_details_parquet  
SELECT * FROM order_details;
```

Optimize Hive Query Performance

TABLE PARTITIONS

Motivation for Partitions

- When you perform lots of repetitive FILTER based queries and the FILTER conditions lend themselves to partitioning

```
SELECT event_type, COUNT(event_type)
FROM call_log
WHERE call_date = '2013-06-03'
GROUP BY event_type;
```

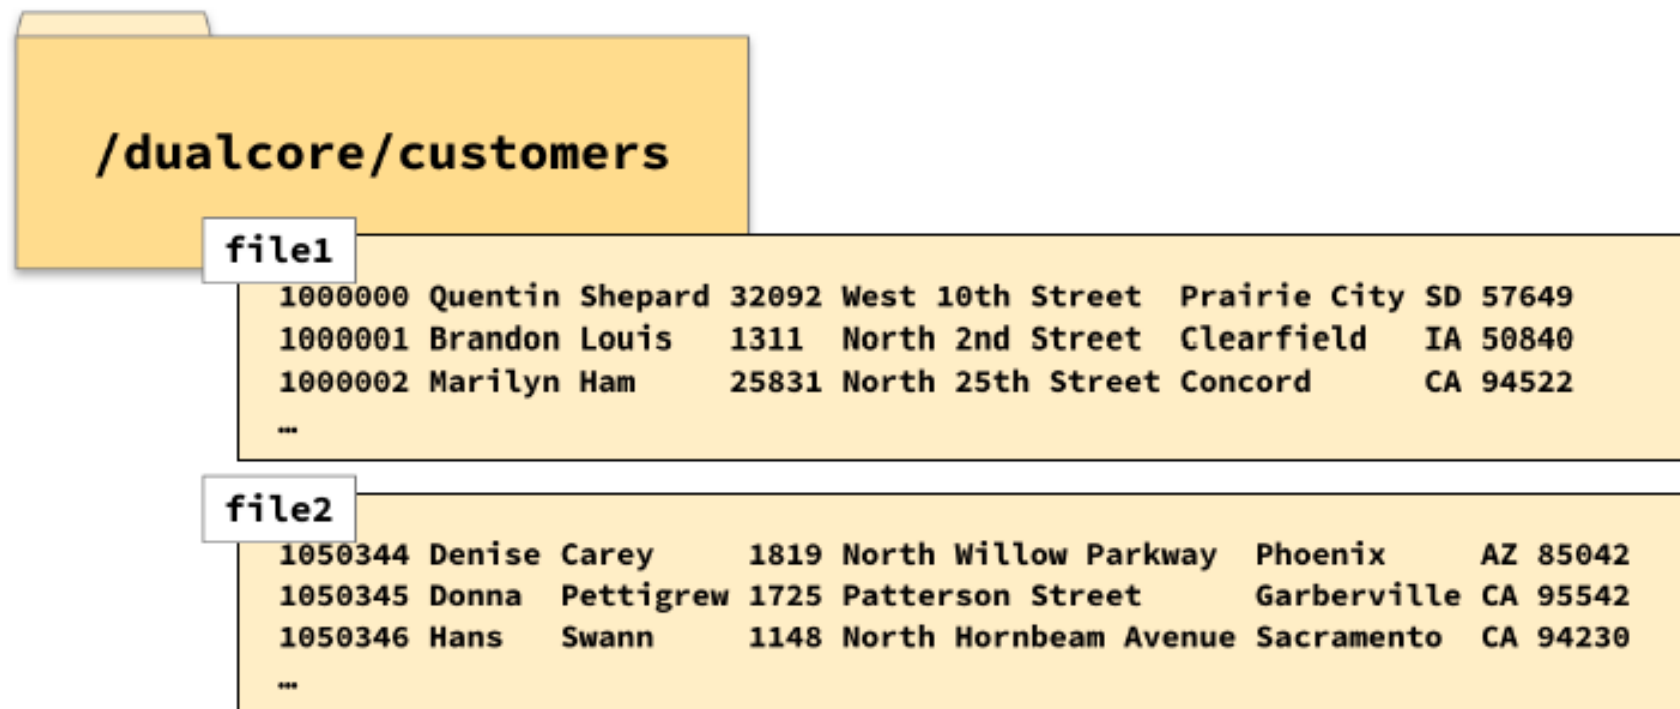
- It does not make sense to read and scan the whole dataset when you only need a small part of it.
 - An appropriate table partition allows Hive to fetch only a portion of the data.

Table Partitioning

- By default, all data files for a table are stored in a single directory
 - All files in the directory are read during a query
- Partitioning subdivides the data
 - Data is physically divided during loading, based on values from one or more columns
- Speeds up queries that filter on partition columns
 - Only the files containing the selected data need to be read
 - Does not prevent you from running queries that span multiple partitions

Example: Partitioning Customers by State

- `customers` is a non-partitioned table
- Data files are stored in a single directory
- All files are scanned for every query



Creating a Partitioned Table

- Using PARTITIONED BY

```
CREATE EXTERNAL TABLE customers (  
  cust_id INT,  
  fname STRING,  
  lname STRING,  
  address STRING,  
  city STRING,  
  state STRING,  
  zipcode STRING)  
ROW FORMAT DELIMITED FIELDS TERMINATED  
BY '\t'  
LOCATION '/dualcore/customers';
```

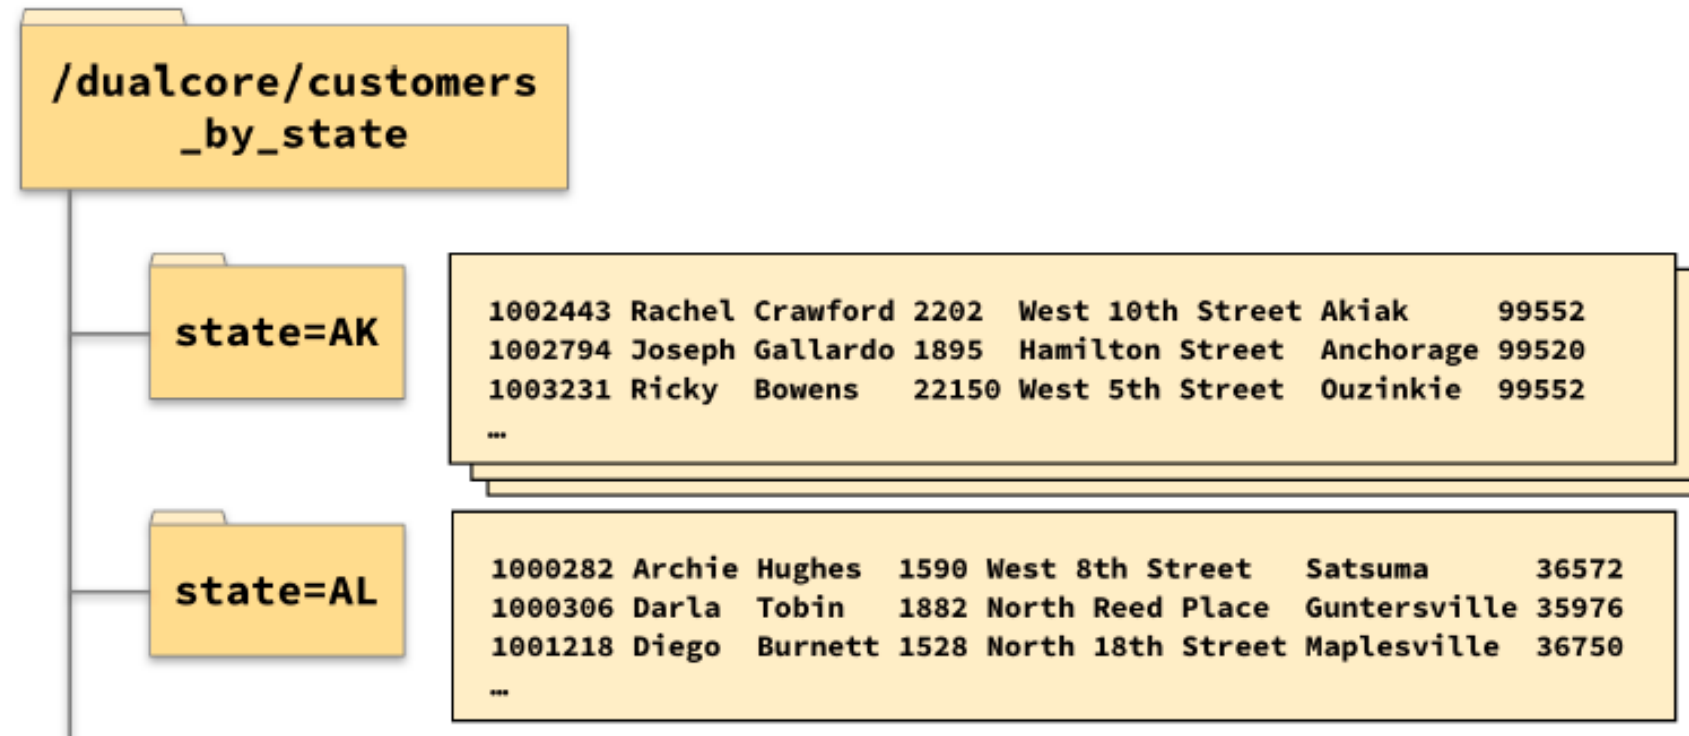
Non-partitioned

```
CREATE EXTERNAL TABLE customers_by_state(  
  cust_id INT,  
  fname STRING,  
  lname STRING,  
  address STRING,  
  city STRING,  
  zipcode STRING)  
PARTITIONED BY (state STRING)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY  
'\t'  
LOCATION '/dualcore/customers';
```

Partitioned

Partitioning File Structure

- Partitioned tables store data in subdirectories
 - Queries that filter on partitioned fields limit amount of data read



Partition fields

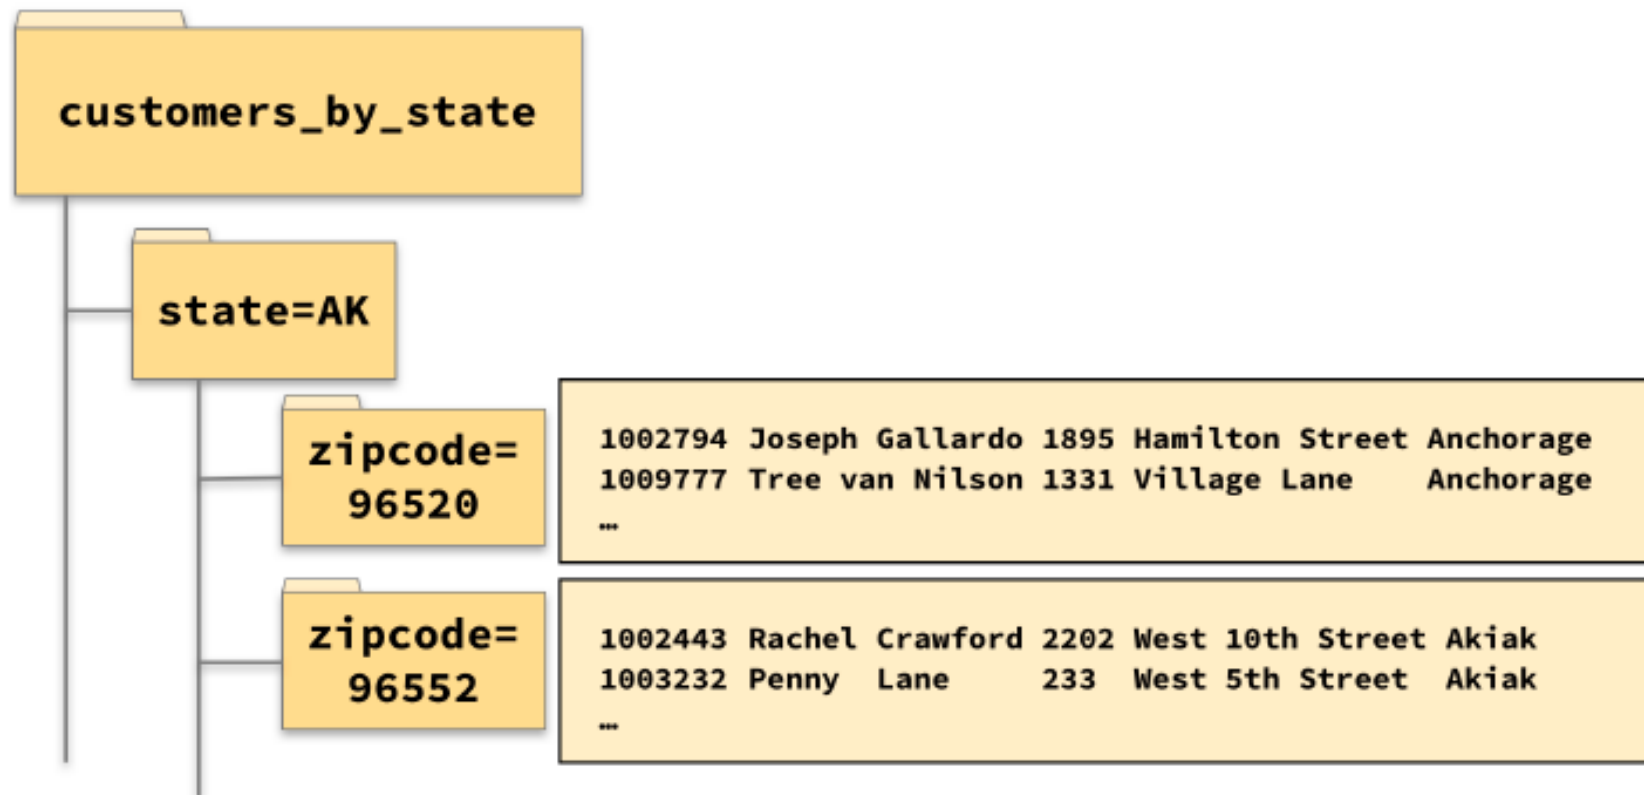
- The partition field(s) is removed from the table, since it is redundant, given the directory name in which the data is stored
- Each subdirectory is not limited to a single file
- A partition field is a *virtual field*. Column values are not stored in the files but the column is displayed if you DESCRIBE the table

```
DESCRIBE customers_by_state;
+-----+-----+-----+
| name      | ty      | comment |
+-----+-----+-----+
| cust_id   | int     |          |
| fname     | string  |          |
| lname     | string  |          |
| address   | string  |          |
| city      | string  |          |
| zipcode   | string  |          |
| state     | string  |          |
+-----+-----+-----+
```

Nested Partitions

- You can also create nested partitions

... PARTITIONED BY (state STRING, zipcode STRING)



When to Use Partitioning

- Use partitioning for tables when
 - Reading the entire dataset takes too long
 - Queries almost always filter on the partition columns
 - There are a reasonable number of different values for partition columns
 - Data generation or ETL process splits data by file or directory names
 - Partition column values are not in the data itself

When *Not* to Use Partitioning

- Avoid partitioning data into numerous small data files
 - Partitioning on columns with too many unique values
- Caution: This can happen easily when using dynamic partitioning!
 - For example, partitioning customers by first name could produce thousands of partitions

Loading Data into a Partitioned Table

- Dynamic partitioning
 - Hive/Impala automatically creates partitions
 - Inserted data is stored in the correct partitions based on column values
- Static partitioning
 - You manually create new partitions using `ADD PARTITION`
 - When loading data, you specify which partition to store it in

Static Partitioning

- With static partitioning, you create new partitions as needed

```
ALTER TABLE customers_by_state  
ADD PARTITION (state='NY');
```

- Adds the partition to the table's metadata if it does not already exist.
- Creates subdirectory state=NY in
/user/hive/warehouse/customers_by_state/

- Then add data one partition at a time, e.g.

```
INSERT OVERWRITE TABLE customers_by_state  
PARTITION(state='NY')  
SELECT cust_id, fname, lname, address,  
city, zipcode FROM customers WHERE state='NY';
```

Dynamic Partitioning

- With dynamic partitioning, you use an INSERT statement to load data
 - The partition column(s) must be included in the PARTITION clause
 - The partition column(s) must be specified last in the SELECT list

```
INSERT OVERWRITE TABLE customers_by_state
    PARTITION(state)
    SELECT cust_id, fname, lname, address, city,
           zipcode, state FROM customers
```

- Hive automatically creates partitions and inserts data into them based on the values of partition column(s)
 - If the partition does not already exist, it will be created
 - If the partition does exist, it will be overwritten

<https://cwiki.apache.org/confluence/display/Hive/Tutorial#Tutorial-Dynamic-PartitionInsert>

Dynamic Partition Inserts (2 Of 3)

- Dynamic partitioning is not enabled by default
 - Enable it by setting these two properties

| Property Name | Value |
|---|------------------------|
| <code>hive.exec.dynamic.partition</code> | <code>true</code> |
| <code>hive.exec.dynamic.partition.mode</code> | <code>nonstrict</code> |

- Remember: Avoid creating an excessive number of partitions
 - This can happen when your data contains many unique values
- Caution: If the partition column has many different values, many partitions will be created
 - Partitioning by date is a popular, but common example

Viewing, Adding, And Removing Partitions

- To view the current partitions in a table

```
SHOW PARTITIONS call_logs;
```

- Use ALTER TABLE to add or drop partitions

```
ALTER TABLE call_logs  
    ADD PARTITION (call_date='2018-06-05');  
ALTER TABLE call_logs  
    DROP PARTITION (call_date='2018-06-06');
```

- Provides a very efficient way to drop data
 - Rather than filtering an entire table, Hive just deletes a subdirectory
- Drops actual data whether or not it is a Hive-managed table.

Optimize Hive Query Performance

UNDERSTAND AND WRITE BETTER PERFORMING HIVE QUERIES

Write better Performing Hive Queries

- Generally speaking, Hive query performance improves if you can
 - Compress intermediate output
 - List columns you need, avoid `SELECT *`
 - Filter early in workflows consisting of several steps
 - e.g. if you can, use `WHERE` in the sub-query instead of in the main query
 - Write better queries that reduce the number of processing steps (e.g. avoid a costly join)
- This requires you to have a better understanding of how Hive executes a query.

Viewing The Execution Plan

- Prefix your query with EXPLAIN to view Hive's execution plan

```
hive> EXPLAIN SELECT zipcode, COUNT(cust_id) AS num  
        FROM customers  
        GROUP BY zipcode;
```

- The output of EXPLAIN can be very long and complex
 - Useful for determine, e.g. how many MapReduce phases it would require?
 - However, fully understanding it requires in-depth knowledge of MapReduce
 - We will cover the basics here

Viewing The Query Plan With EXPLAIN

- Query plan has stages, but stage numbers do not represent the execution sequence
 - They are merely identifiers
 - A root stage always runs first
 - Dependencies define the order of execution
- In the previous example, the stages run in the following order:
 - Stage-1 (MapReduce) runs first
 - Stage-0 (HDFS operation "ListSink") runs next
 - Hive is smart enough to run stages in parallel when there are no dependencies among them

```
STAGE DEPENDENCIES:  
  Stage-1 is a root stage  
  Stage-0 depends on stages: Stage-1
```

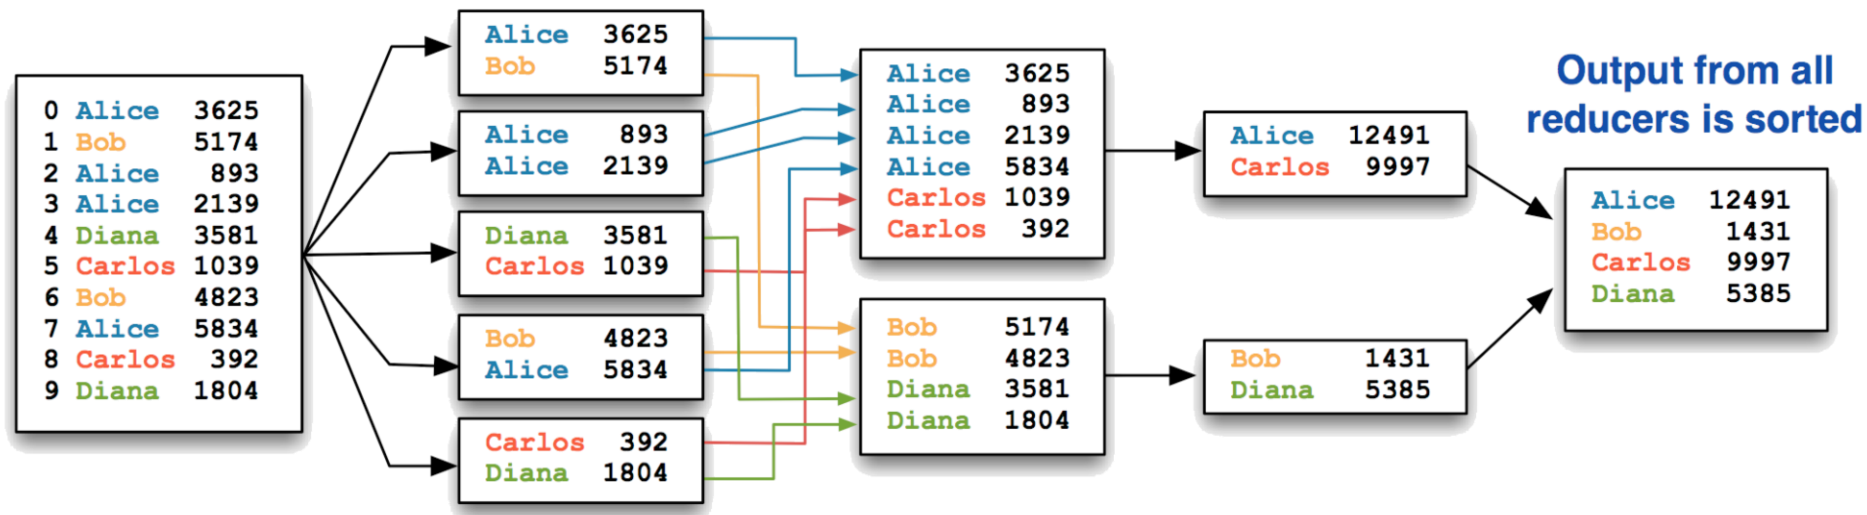
```
STAGE PLANS:  
  Stage: Stage-1  
    Map Reduce  
      Map Operator Tree:...
```

(for more details) <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+Explain>

Sorting Results

- As in SQL, ORDER BY sorts specified fields in HiveQL
 - Consider the result from the following query, utilizing 5 Mappers and 2 Reducers


```
hive> SELECT name, SUM(total)
      FROM order_info
      GROUP BY name
      ORDER BY name;
```



Using SORT BY For Partial Ordering (1 Of 2)

- HiveQL also supports partial ordering via SORT BY
 - Offers much better performance if global order isn't required
 - e.g. You want ordering within each Reducer, but not across multiple Reducers

```
hive> SELECT name, SUM(total)
      FROM order_info
      GROUP BY name
      SORT BY name;
```



Using SORT BY For Partial Ordering (2 Of 2)

- Behind the scene
 - ORDER BY uses a single reducer to ensure global sorting
 - SORT BY sorts the input before feeding them into reducers
 - So that output of each reducer is sorted (partial sorting), but collective output is not.
- SORT BY can improve performance on large queries in 2 ways:
 - Bringing the benefits of parallelism to the reduce phase
 - In some case, it can eliminate the need for a second MapReduce job dedicated to global ordering of results.
- <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+SortBy>

Optimize Hive Query Performance

BUCKETING

Bucketing Data in Hive

- *Bucketing* data is another way of subdividing data
 - Stores data in separate files
 - Divides data into buckets in **an effectively random** way
 - Calculates **hash codes** based on column values
 - Use hash codes to assign records to a bucket
- Goal: Distribute rows across a predefined number of buckets
 - Useful for jobs that need samples of data.
 - Joins may be faster if all tables are bucketed on the join column

Example of Bucketing

- Each bucket should contain roughly 5% of the table's data

```
CREATE TABLE orders_bucketed (order_id INT,  
    cust_id INT, order_date TIMESTAMP)  
    CLUSTERED BY (order_id) INTO 20 BUCKETS;  
  
-- enforce bucketing when inserting data  
SET hive.enforce.bucketing=true;  
INSERT OVERWRITE TABLE orders_bucketed  
    SELECT * FROM orders;  
  
-- sample 10% of data  
SELECT * FROM orders_bucketed  
    TABLESAMPLE (BUCKET 1 OUT OF 10 ON order_id);
```


Essential Points

- Tez/Spark engine can lead to superior performance than the MR engine.
- ORC/PARQUET are significantly faster than the text format.
- Partitioning may reduce the amount of data a query must read
- Understanding Hive execution using the Explain command.
- SORT BY (partial ordering) is faster than ORDER BY
- Bucketing may also improve performance of certain types of queries