




Deep Dive – Amazon Kinesis

Guy Ernest, Solution Architect - Amazon Web Services
@guyernest



Motivation for Real Time Analytics





Deployment & Administration

App Services

Analytics


Compute

Storage

Database

Networking

AWS Global Infrastructure



Amazon Kinesis

Managed Service for Real Time Big Data Processing

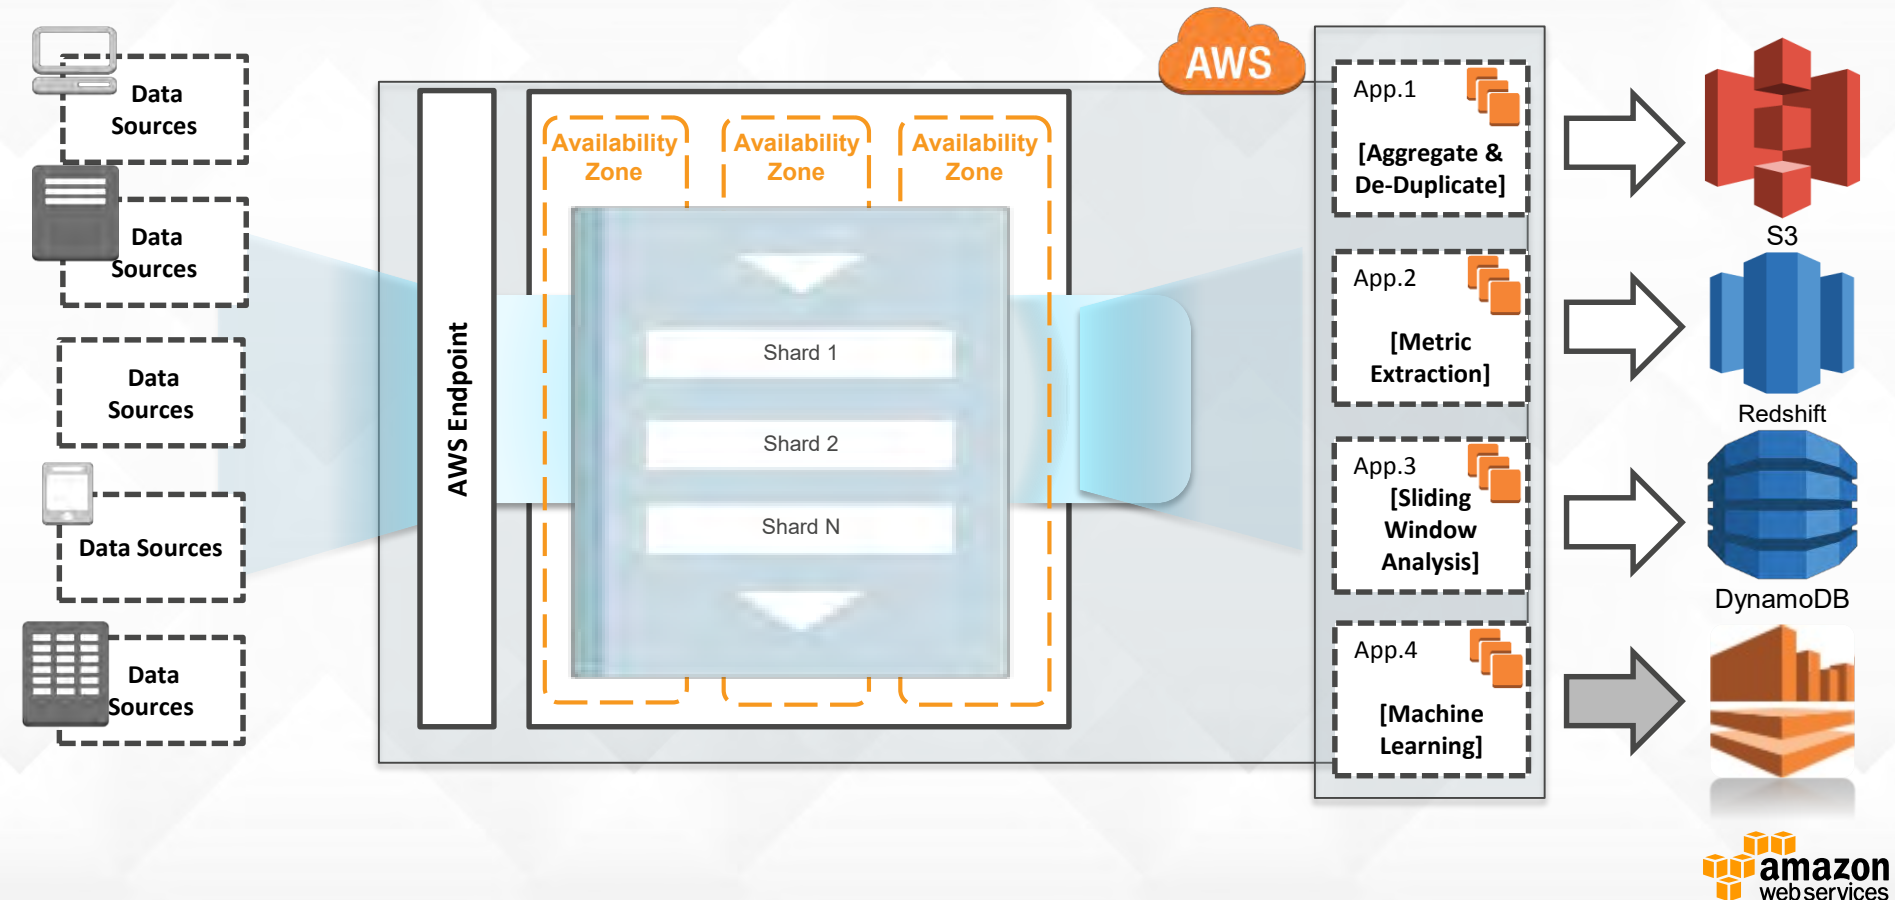
Create Streams to Produce & Consume Data

Elastically Add and Remove Shards for Performance

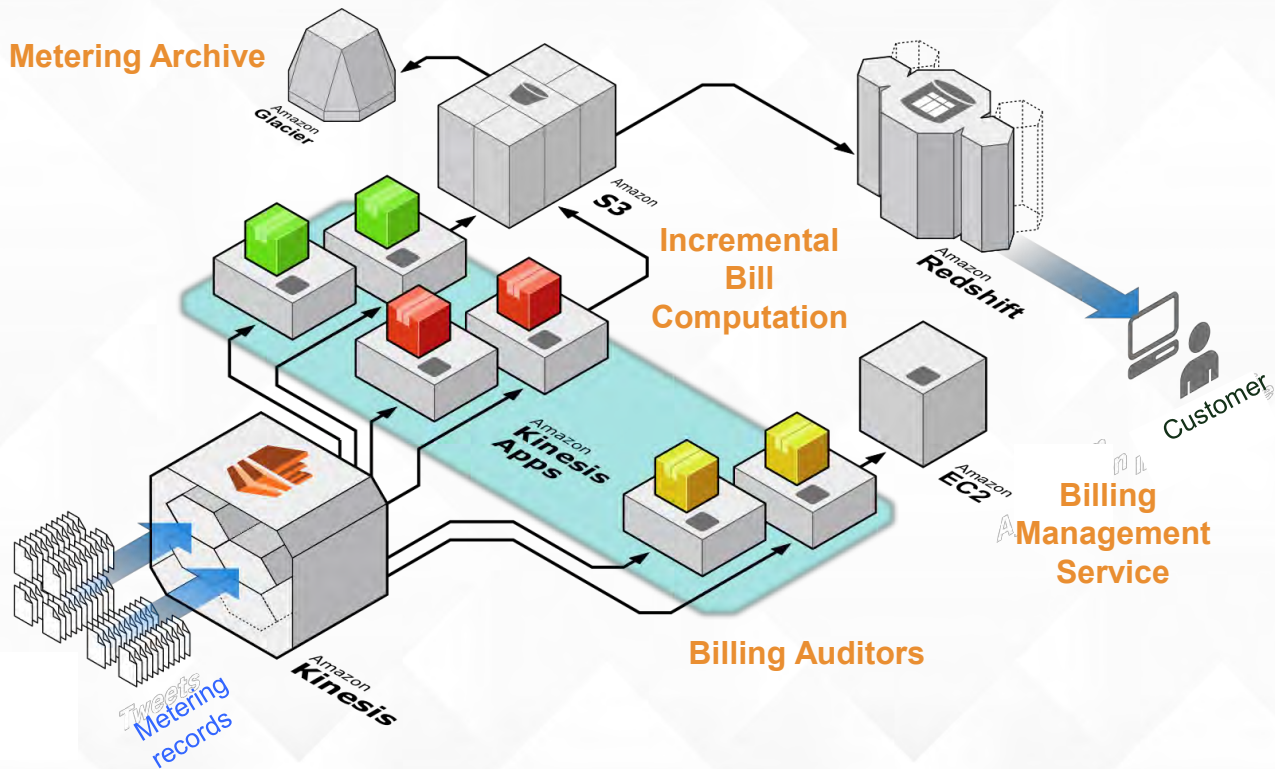
Use Kinesis Worker Library, AWS Lambda, Apache Spark and Apache Storm to Process Data

Integration with S3, Redshift and Dynamo DB

Amazon Kinesis Dataflow



Example Architecture - Metering



Amazon Kinesis Components



Streams

Named Event Streams of Data

All data is stored for 24 hours

Shards

You scale Kinesis streams by adding or removing Shards

Each Shard ingests up to 1MB/sec of data and up to 1000 TPS

Partition Key

Identifier used for Ordered Delivery & Partitioning of Data across Shards

Sequence

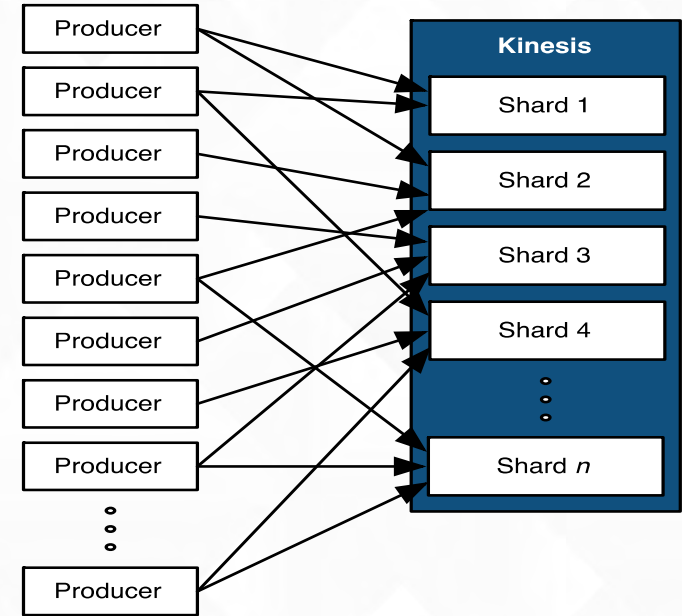
Number of an event as assigned by Kinesis



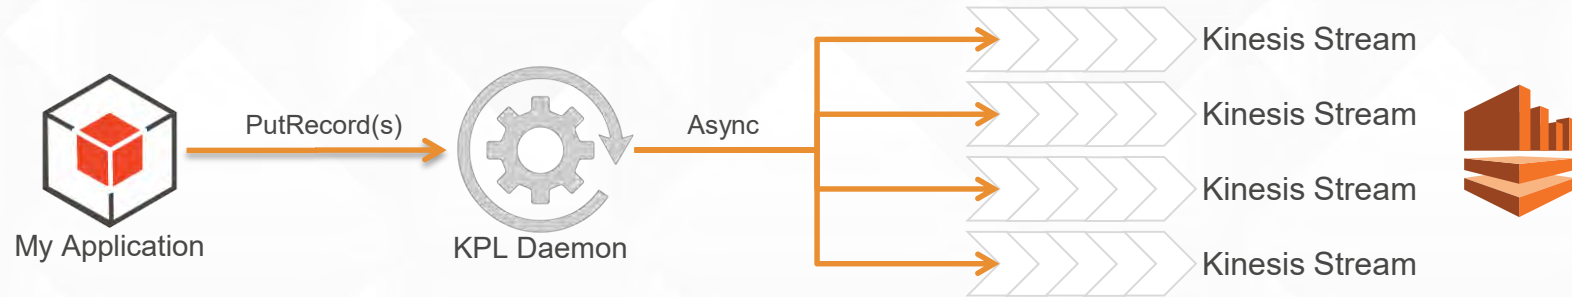
Getting Data In

Kinesis - Ingesting Fast Moving Data

- Producers use a PUT call to store data in a Stream
- A Partition Key is used to distribute the PUTs across Shards
- A unique Sequence # is returned to the Producer for each Event
- Data can be ingested at 1MB/second or 1000 Transactions/second per Shard
- 1MB / Event

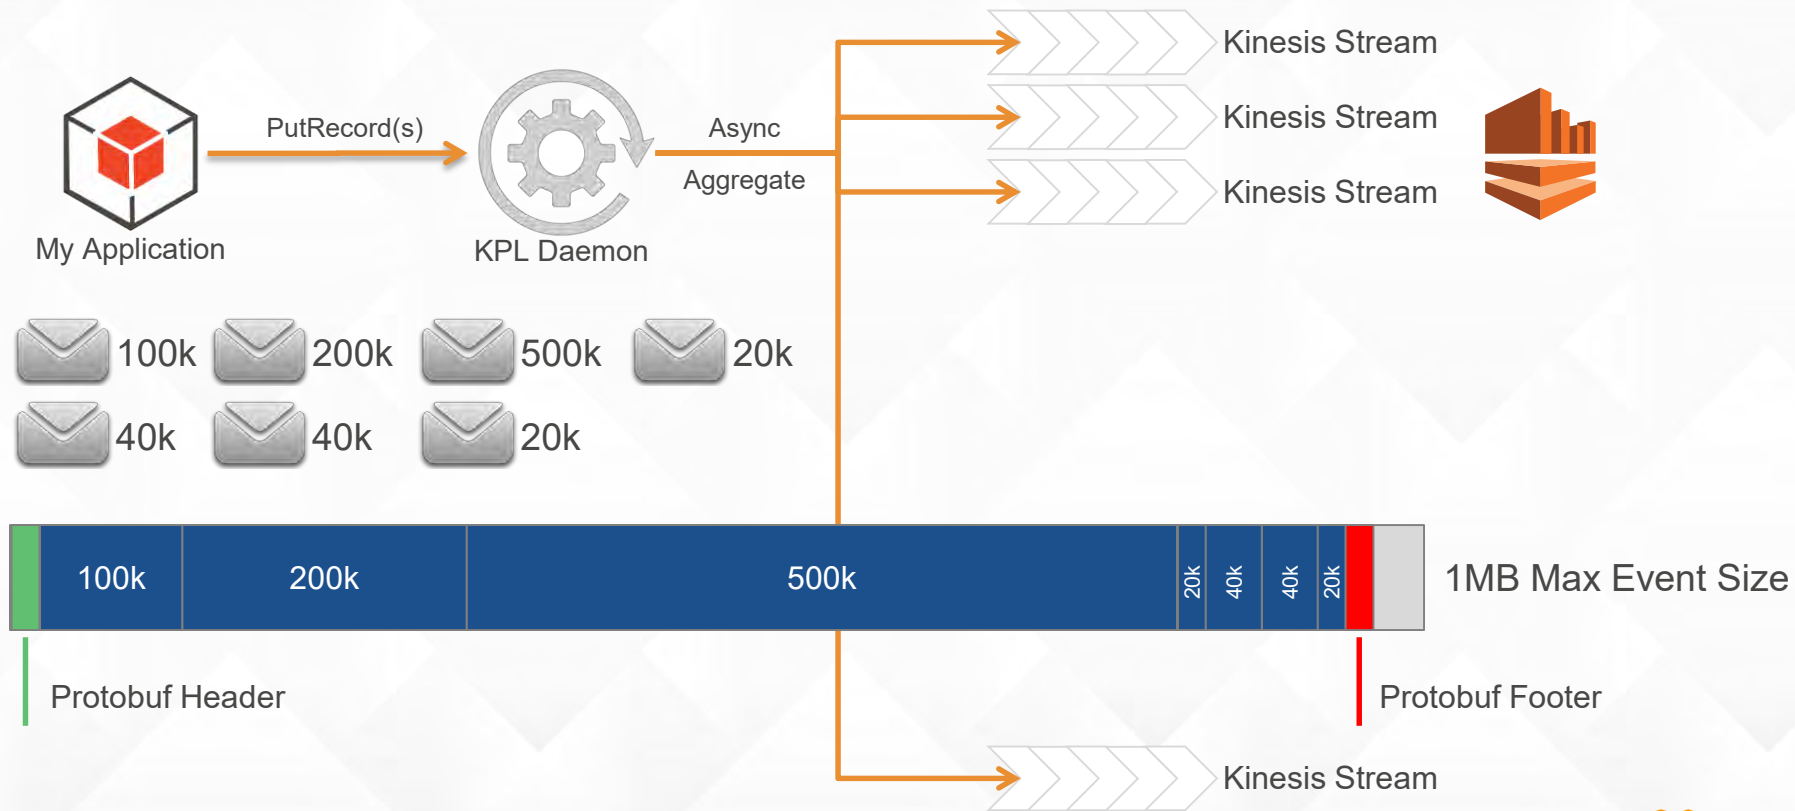


Introducing the Kinesis Producer Library



- Native Code Module to perform efficient writes to Multiple Kinesis Streams
- C++/Boost
- Asynchronous Execution
- Configurable Aggregation of Events

KPL Aggregation



Kinesis Ecosystem - Ingest

Apache Flume

Source & Sink

<https://github.com/pdeyhim/flume-kinesis>

FluentD

Dynamic Partitioning Support

<https://github.com/awslabs/aws-fluent-plugin-kinesis>

Log4J & Log4Net

Included in Kinesis Samples



Best Practices for Partition Key

- Random will give even distribution
- If events should be processed together, choose a relevant high cardinality partition key and monitor shard distribution
- If partial order is important use sequence number



Getting Data Out

Kinesis Client Library

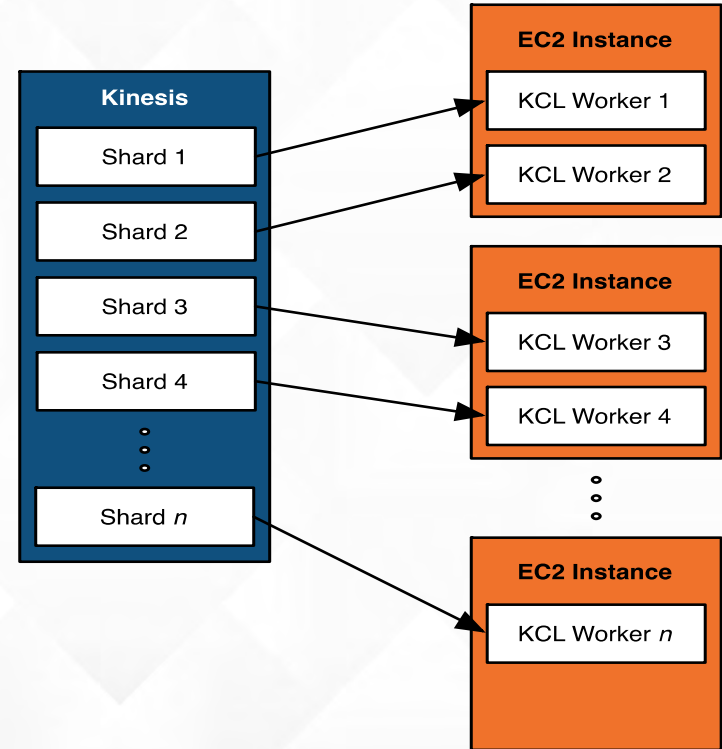
- ❏ KCL Libraries available for Java, Ruby, Node, Go, and a Multi-Lang Implementation with Native Python support
- ❏ All State Management in Dynamo DB



Consuming Data - Kinesis Enabled Applications

Client library for fault-tolerant, at least-once, real-time processing

- 📦 Kinesis Client Library (KCL) simplifies reading from the stream by abstracting your code from individual shards
- 📦 Automatically starts a Worker Thread for each Shard
- 📦 Increases and decreases Thread count as number of Shards changes
- 📦 Uses checkpoints to keep track of a Thread's location in the stream
- 📦 Restarts Threads & Workers if they fail



Kinesis Connectors

Analytics Tooling Integration (github.com/aws-labs/amazon-kinesis-connectors)
S3

Batch Write Files for Archive into S3

Sequence Based File Naming

Redshift

Once Written to S3, Load to Redshift

Manifest Support

User Defined Transformers

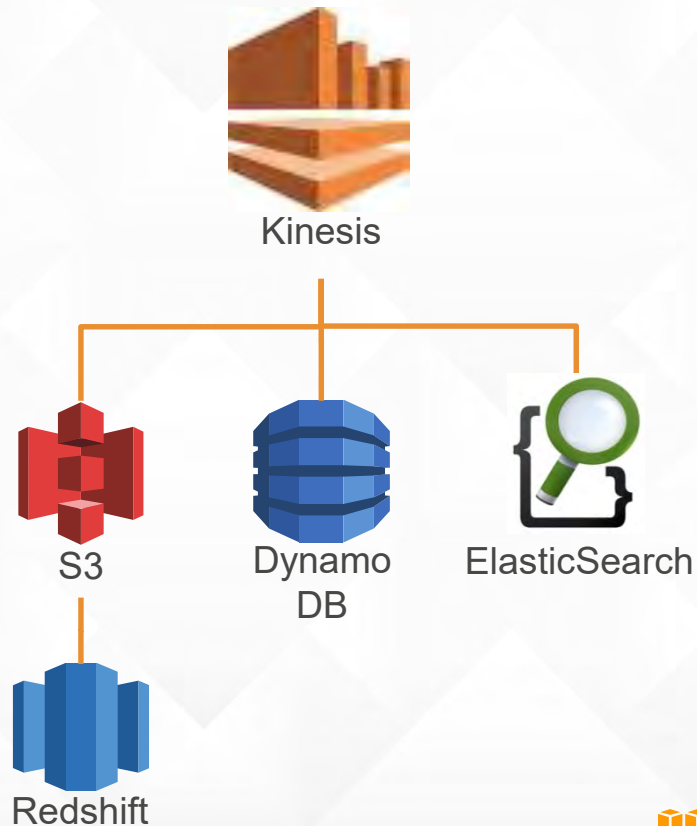
DynamoDB

BatchPut Append to Table

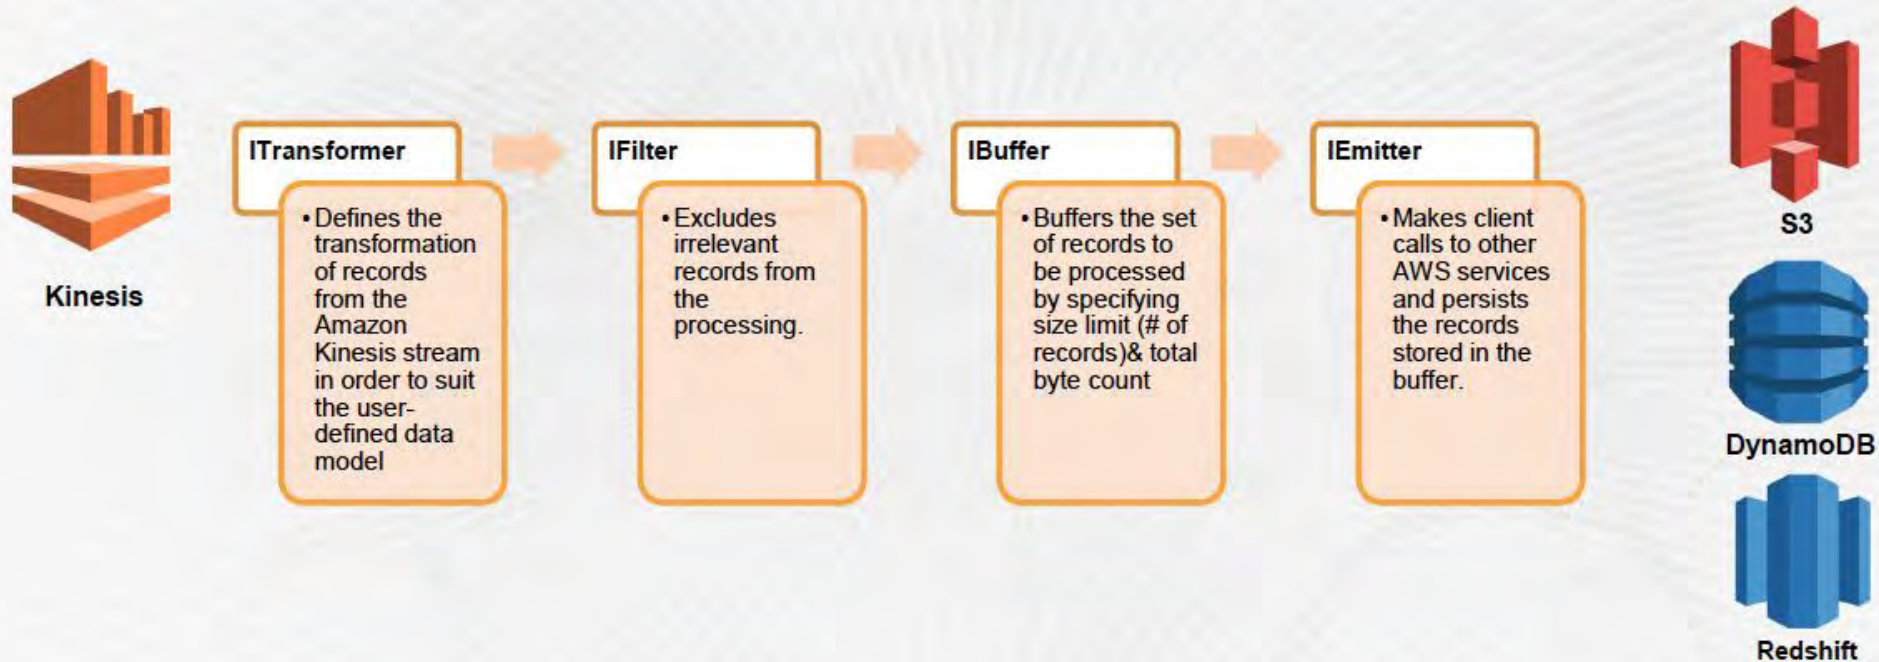
User Defined Transformers

ElasticSearch

Automatically index Stream



Connectors Architecture





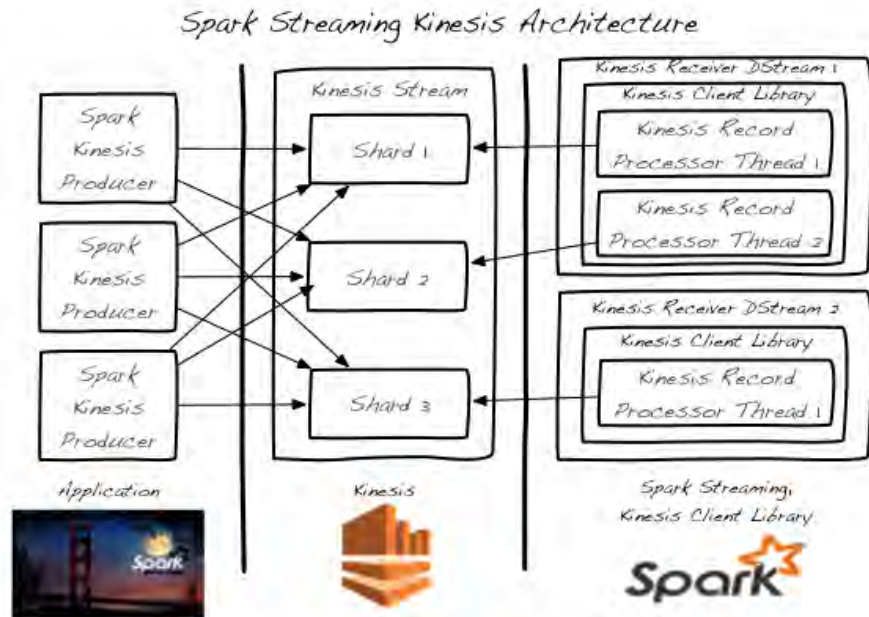
Apache Storm

Kinesis Spout

**Automatic Checkpointing with
Zookeeper**

<https://github.com/awslabs/kinesis-storm-spout>

Kinesis Ecosystem - Spark



Apache Spark

DStream Receiver runs KCL

One DStream per Shard

Checkpointed via KCL

Spark Natively Available on EMR

EMRFS overlay on HDFS

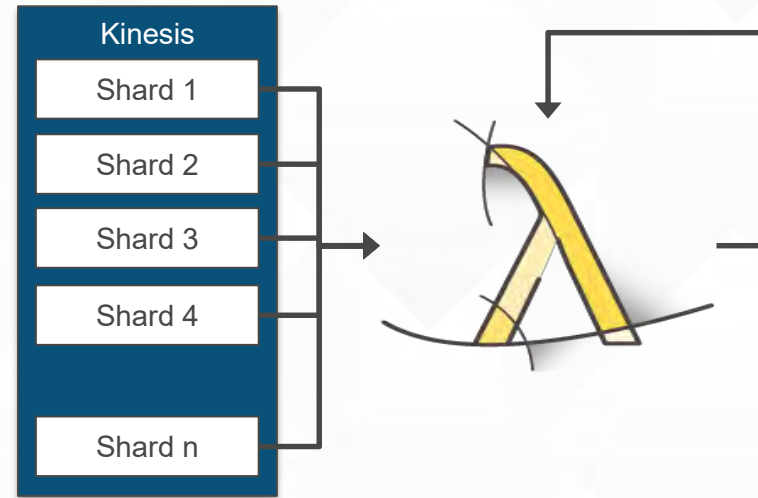
AMI 3.8.0

<https://aws.amazon.com/elasticmapreduce/details/spark>

Consuming Data - AWS Lambda

Distributed Event Processing Platform

- ❏ **Stateless JavaScript & Java functions run against an Event Stream**
- ❏ **AWS SDK Built In**
- ❏ **Configure RAM and Execution Timeout**
- ❏ **Functions automatically invoked against a Shard**
- ❏ **Community libraries for Python & Go**
- ❏ **Access to underlying filesystem for read/write**
- ❏ **Call other Lambda Functions**



Why Kinesis? Durability



Regional Service
Synchronous Writes to Multiple
AZ's
Extremely High Durability



May be in-memory for
Performance
Requirement to understand Disk
Sync Semantics
User Managed Replication
Replication Lag -> RPO

Why Kinesis? Performance



Perform continual processing on streaming big data. Processing latencies fall to a < 1 second, compared with the minutes or hours associated with batch processing



Processing latencies < 1 second
Based on CPU & Disk Performance
Cluster Interruption -> Processing Outage

Why Kinesis? Availability



Regional Service
Synchronous Writes to Multiple
AZ's
Extremely High Durability
AZ, Networking, & Chain Server
Issues Transparent to Producers
& Consumers



Many Depend on a CP Database
Lost Quorum can result in
failure/inconsistency of the cluster
Highest Availability is determined
by Availability of Cross-AZ Links
or Availability of an AZ

Why Kinesis? Operations



Managed service for real-time streaming data collection, processing and analysis. Simply create a new stream, set the desired level of capacity, and let the service handle the rest



Build Instances
Install Software
Operate Cluster
Manage Disk Space
Manage Replication
Migrate to new Stream on Scale Up

Why Kinesis? Elasticity



Seamlessly scale to match your data throughput rate and volume.

You can easily scale up to gigabytes per second. The service will scale up or down based on your operational or business needs



Fixed Partition Count up Front

Maximum Performance ~ 1

Partition/Core | Machine


Convert from 1 Stream to

Another to Scale

Application Reconfiguration

Scaling Streams

Kinesis Stream



Why Kinesis? Cost



Cost-efficient for workloads of any scale. You can get started by provisioning a small stream, and pay low hourly rates only for what you use.

Scale Up/Down Dynamically

\$.015/Hour/1MB



Run your Own EC2 Instances
Multi-AZ Configuration for
increased Durability
Utilise Instance AutoScaling on
Worker Lag from HEAD with
Custom Metrics

Why Kinesis? Cost



- Price Dropped on 2nd June 2015, Restructured to support KPL
- Old Pricing: \$.028 / 1M Records PUT
- New Pricing: \$.014/1M 25KB “Payload Units”

Scenario: 50,000 Events / Second, 512B / Event = 24.4 MB/Second

		Old Pricing		New Pricing + KPL	
	Units	Cost	Units	Cost	
Shards	50	\$558	25	\$279	
PutRecords	4,320M Records	\$120.96	2,648M Payload Units	\$37.50	
		\$678.96			\$316.50

A middle-aged man with a beard and mustache, wearing a dark suit, light blue shirt, and dark tie, stands in the center of a busy city street. He is smiling and holding a large white rectangular sign in front of his chest. The sign contains the text "Everything Fails, All the Time" in a large, dark, sans-serif font. The background is a blurred city street with buildings, pedestrians, and a traffic light. The lighting is bright, suggesting daytime.


Everything
Fails,
All the Time

Kinesis – Consumer Application Best Practices

- ❏ Tolerate Failure of: Threads – Consider Data Serialisation issues and Lease Stealing; Hardware – AutoScaling may add nodes as needed
- ❏ Scale Consumers up and down as the number of Shards increase or decrease
- ❏ Don't store data in memory in the workers. Use an elastic data store such as Dynamo DB for State

- ❏ Elastic Beanstalk provides all Best Practices in a simple to deploy, multi-version Application Container
- ❏ KCL will automatically redistribute Workers to use new Instances

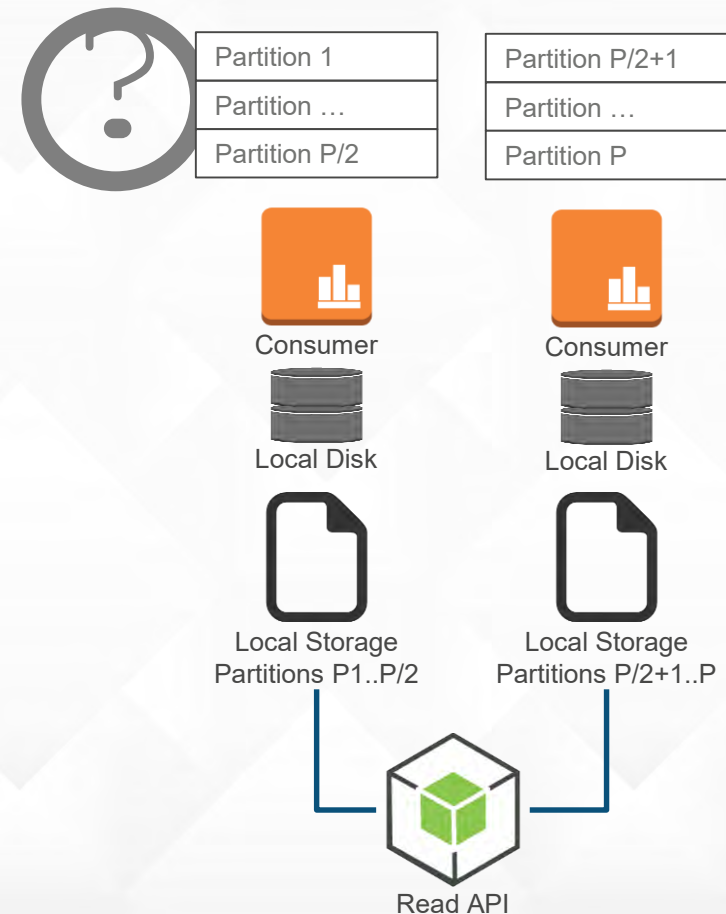
- ❏ Logic implemented in Lambda doesn't require any Servers at all!





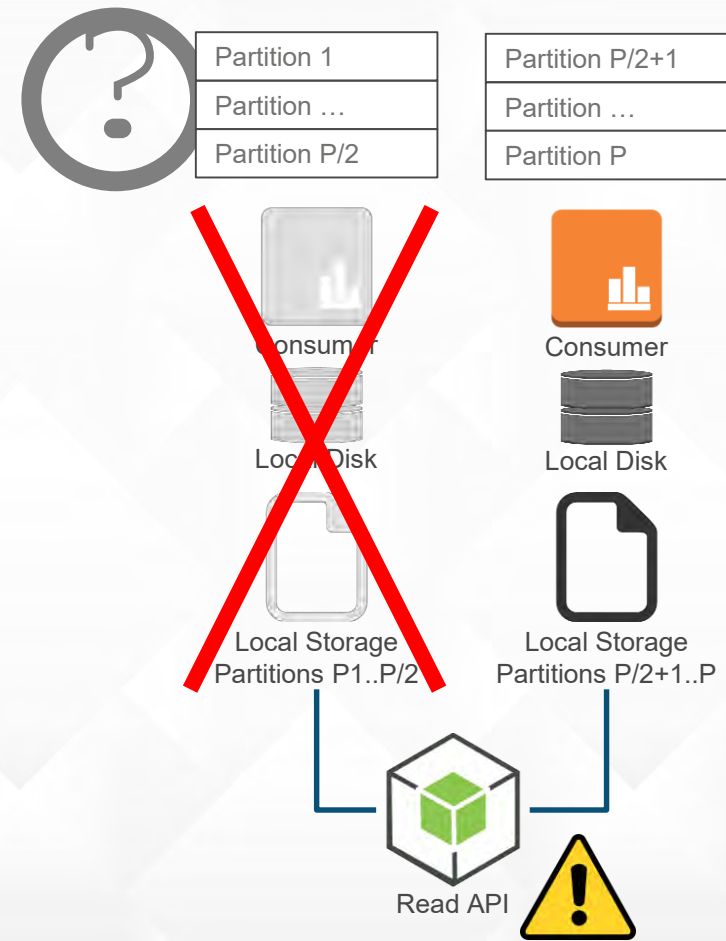
Managing Application State

Consumer Local State Anti-Pattern



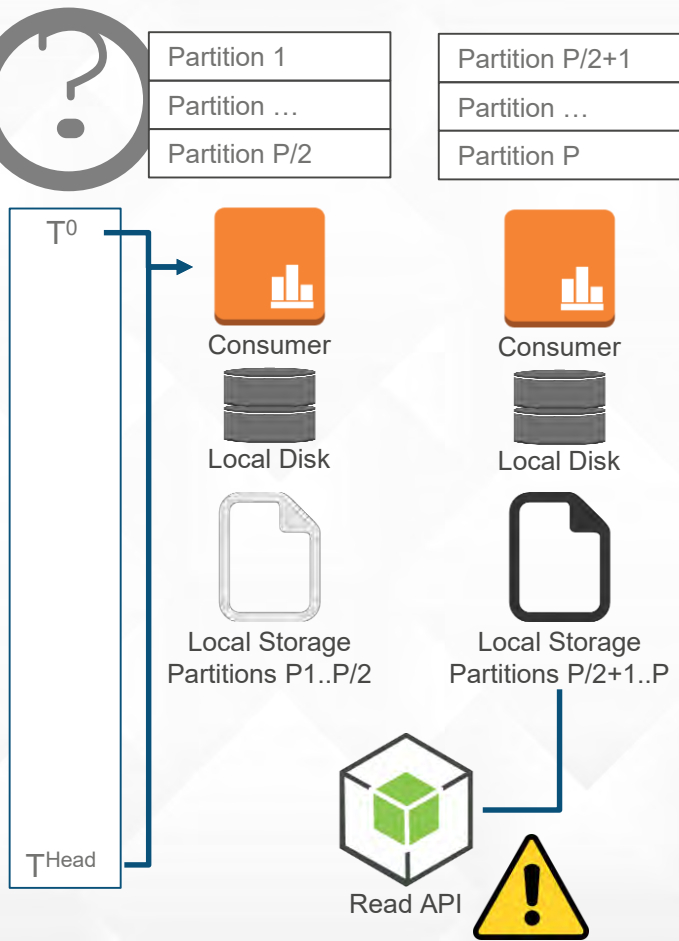
- Consumer binds to a configured number of Partitions
- Consumer stores the 'state' of a data structure, as defined by the event stream, on local storage
- Read API can access that local storage as a 'shard' of the overall database

Consumer Local State Anti-Pattern



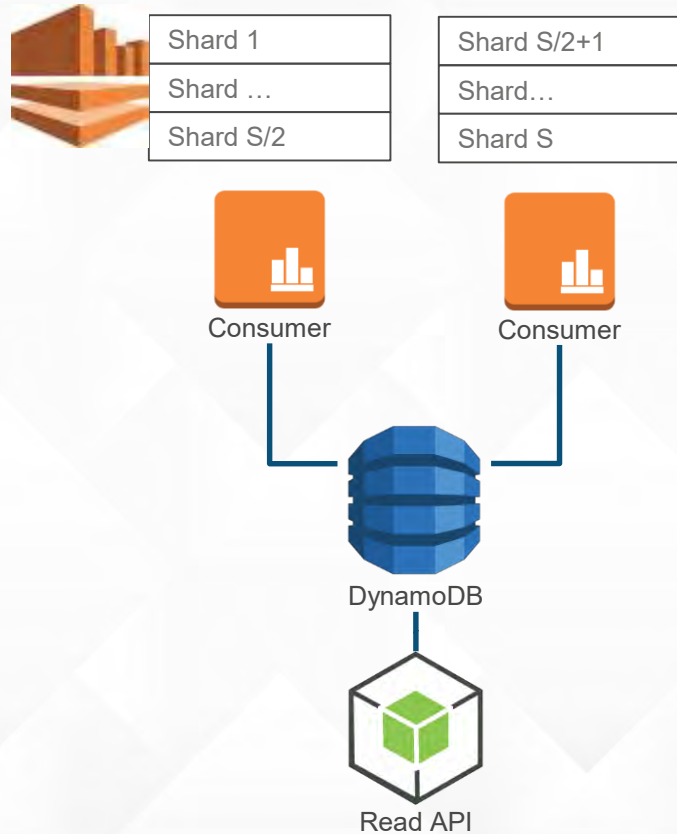
But what happens when an instance fails?

Consumer Local State Anti-Pattern



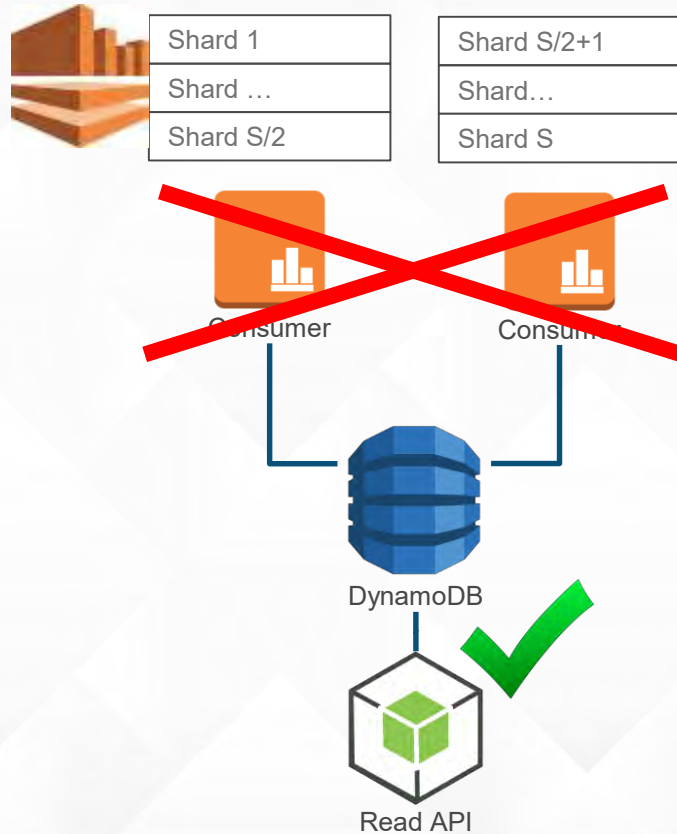
- ❏ A new consumer process starts up for the required Partitions
- ❏ Consumer must read from the beginning of the Stream to rebuild local storage
- ❏ Complex, error prone, user constructed software
- ❏ Long Startup Time

External Highly Available State – Best Practice



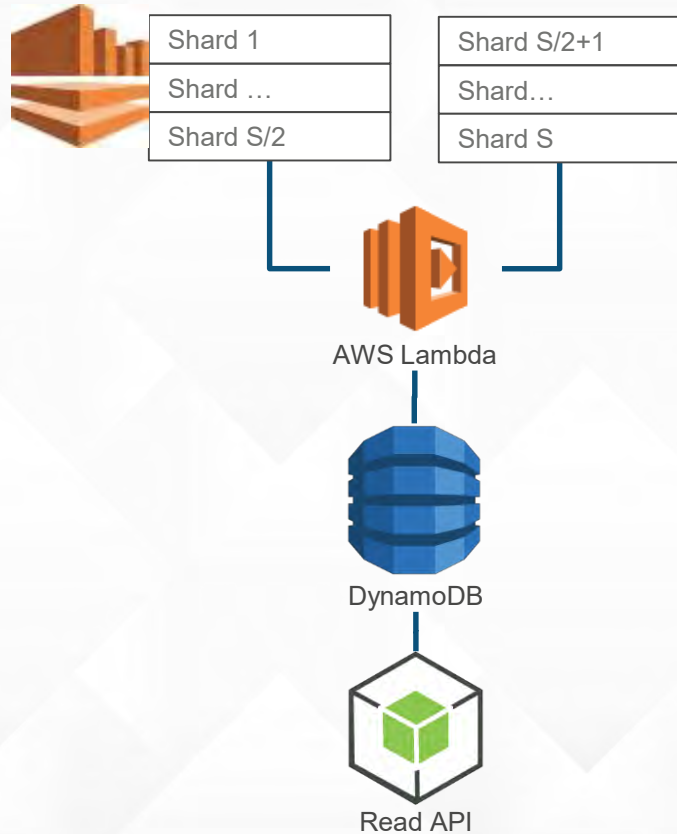
- Consumer binds to an even number of Shards based on number of Consumers
- Consumer stores the 'state' in Dynamo DB
- Dynamo DB is Highly Available, Elastic & Durable
- Read API can access Dynamo DB

External Highly Available State – Best Practice



- Consumer binds to an even number of Shards based on number of Consumers
- Consumer stores the 'state' in Dynamo DB
- Dynamo DB is Highly Available, Elastic & Durable
- Read API can access Dynamo DB

External Highly Available State – Best Practice



- Consumer binds to an even number of Shards based on number of Consumers
- Consumer stores the 'state' in Dynamo DB
- Dynamo DB is Highly Available, Elastic & Durable
- Read API can access Dynamo DB




Idempotency

Property of a system whereby the repeated application of a function on a single input results in the same end state of the system

...

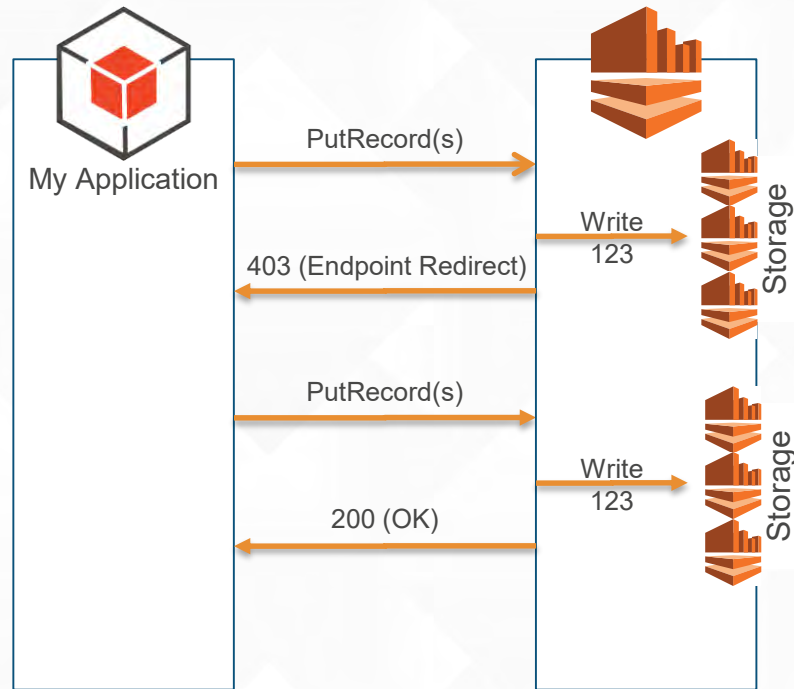
Exactly Once Processing

Idempotency – Writing Data



- ❏ The Kinesis SDK & KPL may retry PUT in certain circumstances
- ❏ Kinesis Record acknowledged with a Sequence Number **is** durable to Multiple Availability Zones...
- ❏ But there could be a duplicate entry

Idempotency – Writing Data




- ❏ The Kinesis SDK & KPL may retry PUT in certain circumstances
- ❏ Kinesis Record acknowledged with a Sequence Number **is** durable to Multiple Availability Zones...
- ❏ But there could be a duplicate entry



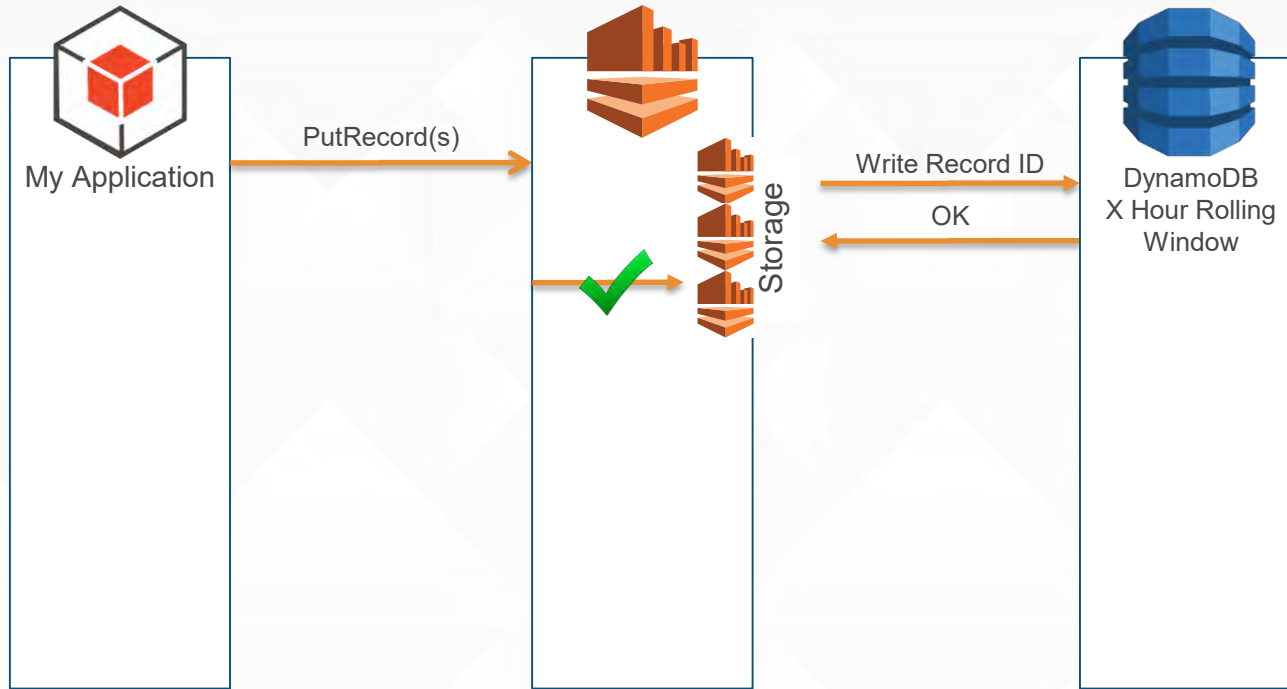
Coming Soon...

Idempotency – Rolling Idempotency Check



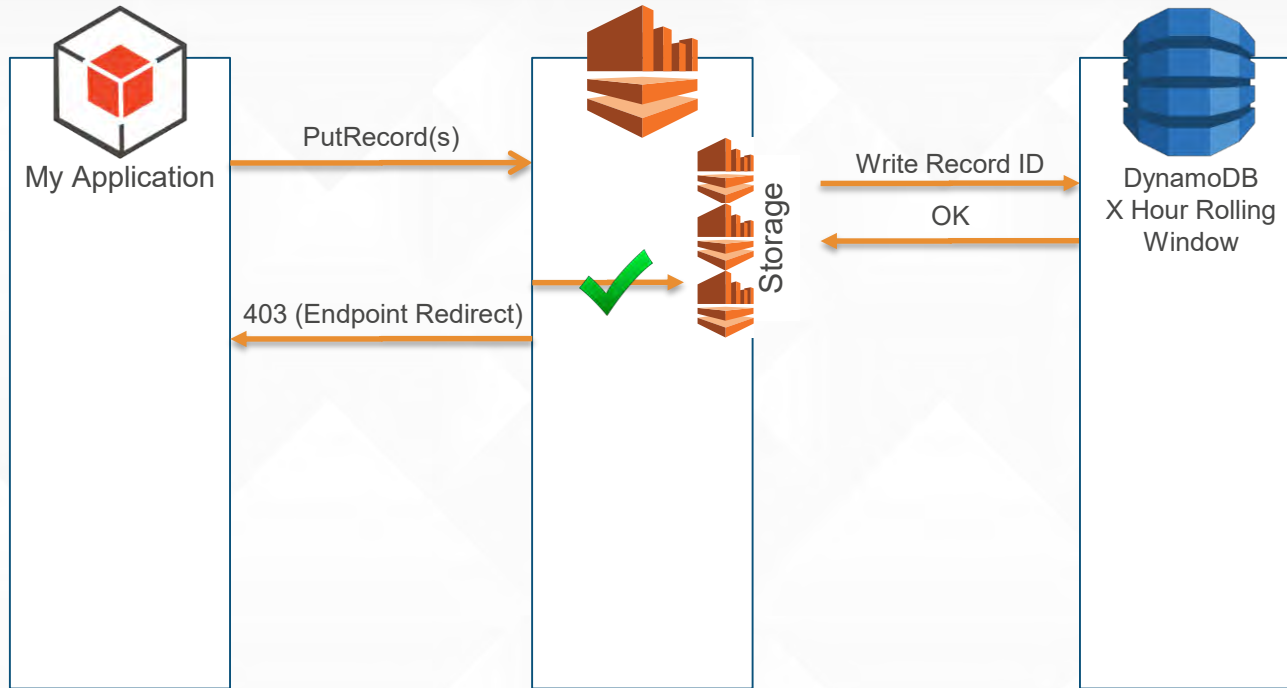
- 📦 Kinesis will manage a rolling time window of Record ID's in Dynamo DB
- 📦 Record ID's are User Based
- 📦 Duplicates in storage tier will be acknowledged as Successful

Idempotency – Rolling Idempotency Check



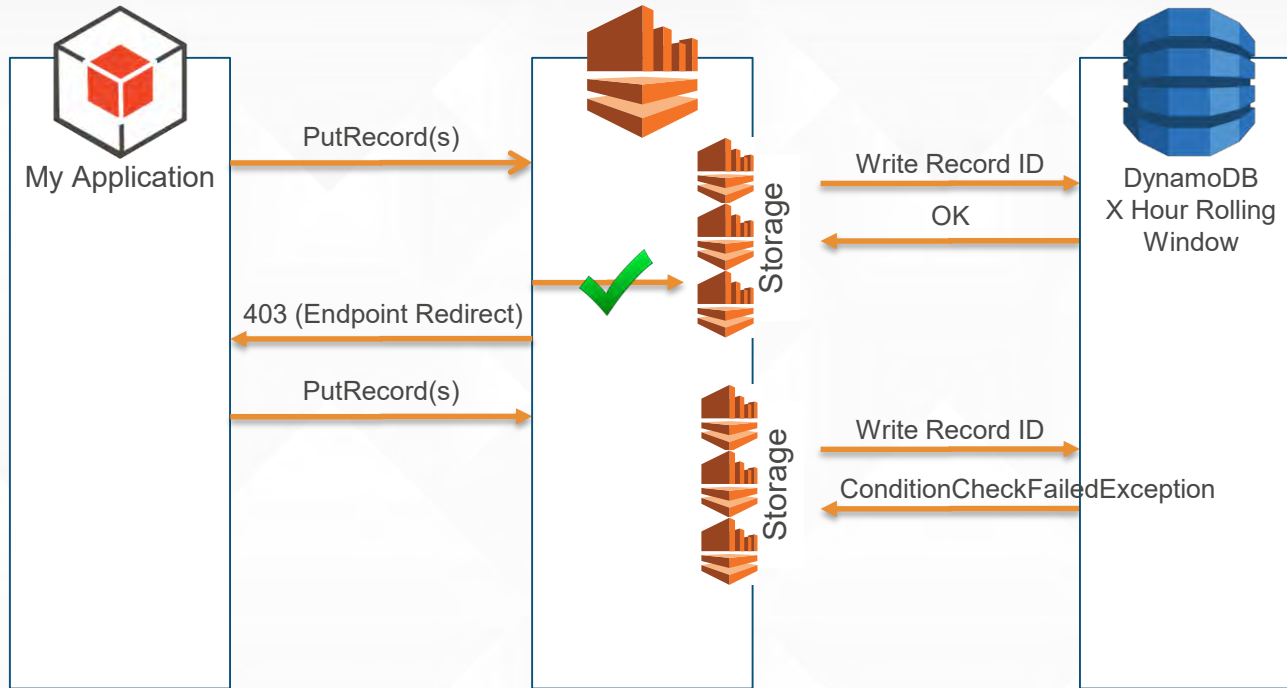
- 📦 Kinesis will manage a rolling time window of Record ID's in Dynamo DB
- 📦 Record ID's are User Based
- 📦 Duplicates in storage tier will be acknowledged as Successful

Idempotency – Rolling Idempotency Check



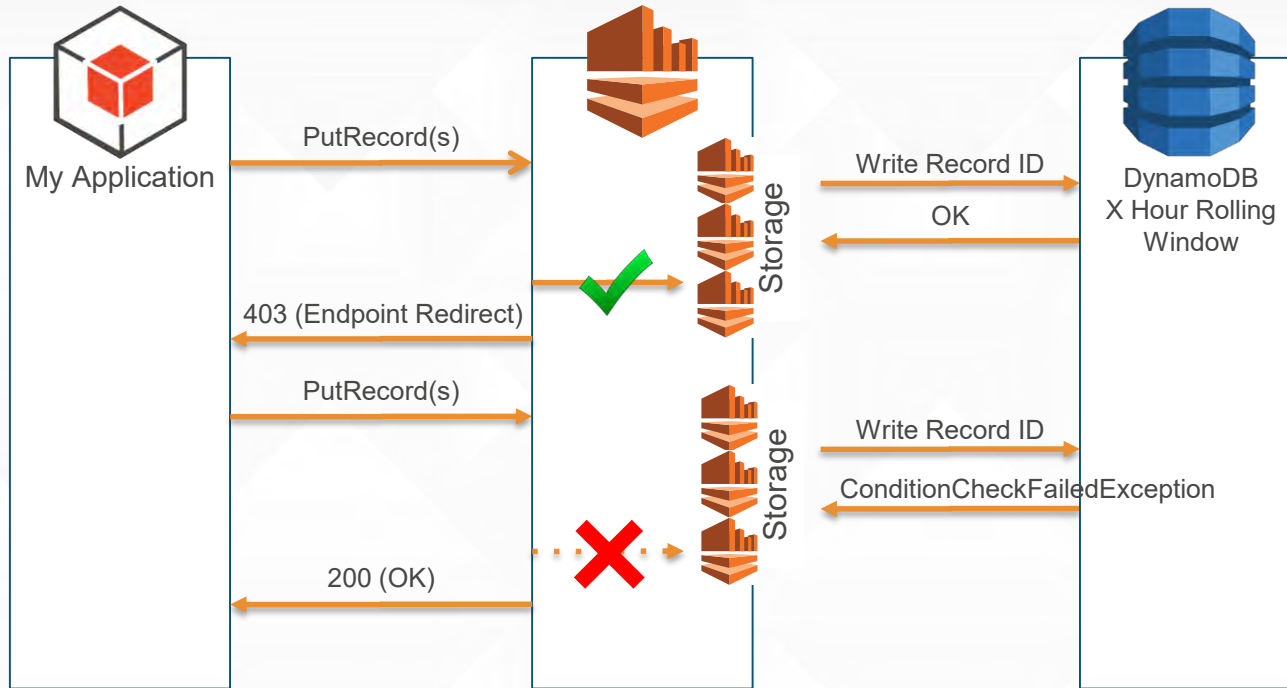
- 📦 Kinesis will manage a rolling time window of Record ID's in Dynamo DB
- 📦 Record ID's are User Based
- 📦 Duplicates in storage tier will be acknowledged as Successful

Idempotency – Rolling Idempotency Check



- ❏ Kinesis will manage a rolling time window of Record ID's in Dynamo DB
- ❏ Record ID's are User Based
- ❏ Duplicates in storage tier will be acknowledged as Successful

Idempotency – Rolling Idempotency Check



- ❏ Kinesis will manage a rolling time window of Record ID's in Dynamo DB
- ❏ Record ID's are User Based
- ❏ Duplicates in storage tier will be acknowledged as Successful

In Short...



Easy Administration



Real-time
Performance.
High Durability



High Throughput.
Elastic



S3, Redshift, &
DynamoDB
Integration



Large Ecosystem



Low Cost





Thank You