



# SPARK STREAMING PROGRAMMING TECHNIQUES YOU SHOULD KNOW

Gerard Maas  Lightbend

#EUstr2

# Gerard Maas


Sr SW Engineer

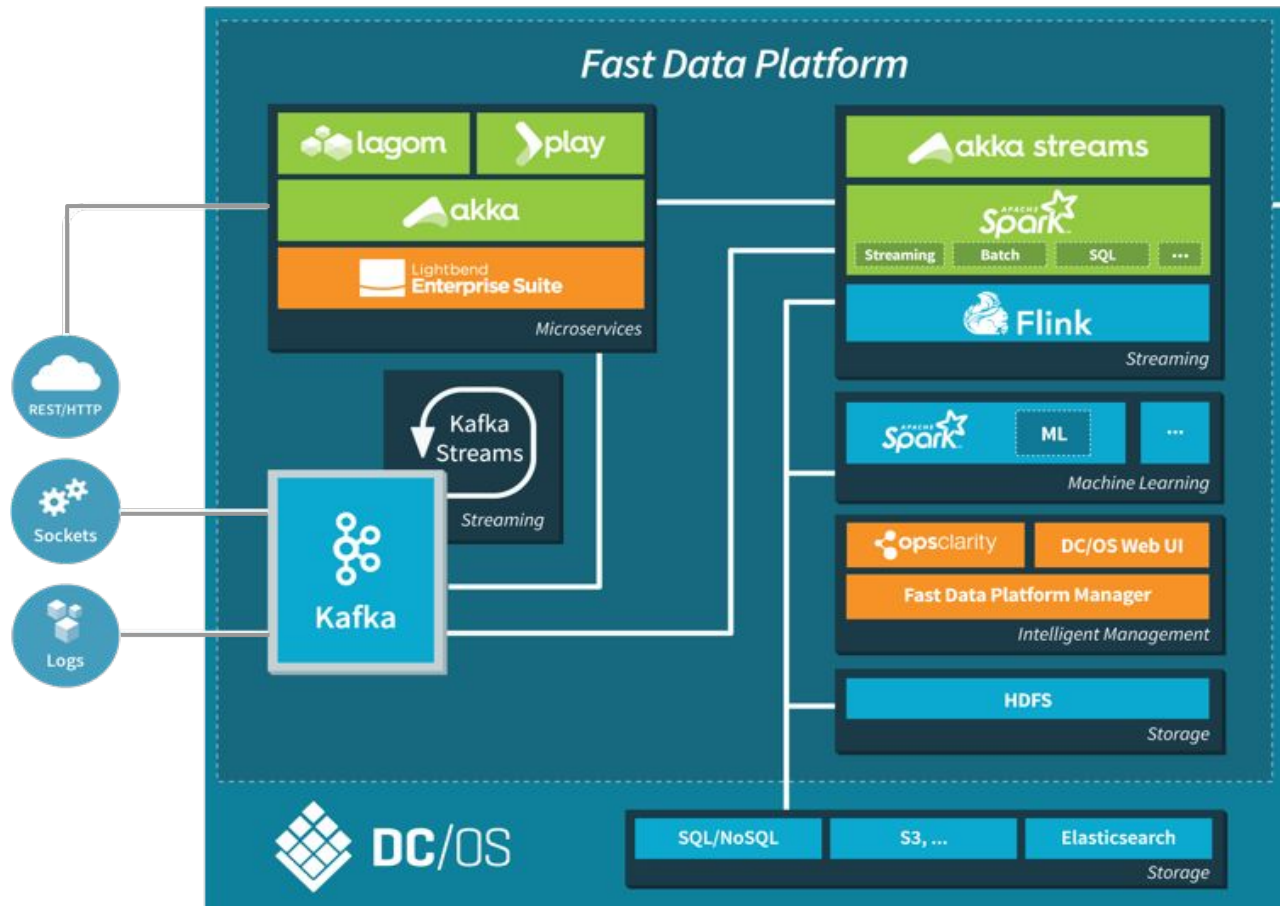


Computer Engineer  
Scala Programmer  
Early Spark Adopter (v0.9)  
Spark Notebook Contributor

Cassandra MVP (2015, 2016)  
Stack Overflow Top Contributor  
(Spark, Spark Streaming, Scala)

Wannabe {  
IoT Maker  
Drone crasher/tinkerer  
}





[lightbend.com/fast-data-platform](http://lightbend.com/fast-data-platform)

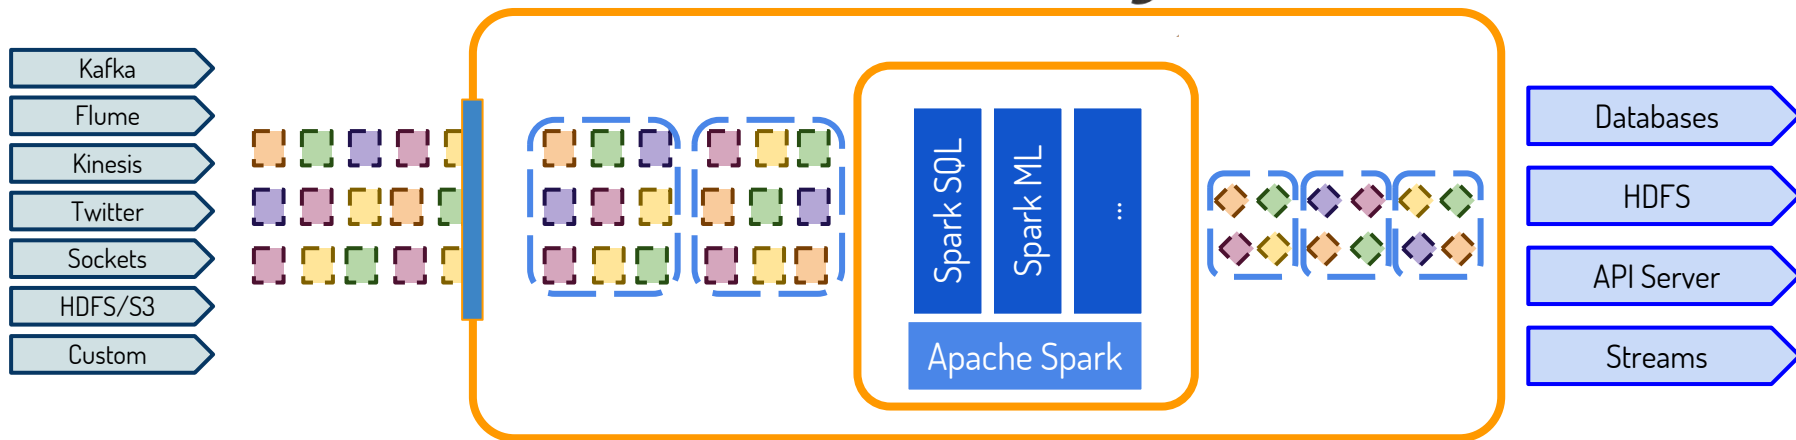
# Agenda

- Spark Streaming Refresher
  - Model
  - Operations
- Techniques
  - Self-contained stream generation
  - Refreshing external data
  - Structured Streaming compatibility
  - Keeping arbitrary state
  - Probabilistic accumulators

# Spark Streaming Refresher



# Spark Streaming






# API




# Transformations


map,  
flatMap,  
filter



count,  
reduce,  
countByValue,  
reduceByKey




union,  
join  
cogroup





# Transformations

mapWithState



# Transformations



```
val iotDstream = MQTTUtils.createStream(...)
val devicePriority = sparkContext.cassandraTable(...)
val prioritizedDStream = iotDstream.transform{rdd =>
  rdd.map(d => (d.id, d)).join(devicePriority)
}
```

# Actions

print




Time: 1459875469000 ms

data1

data2

saveAsTextFiles,  
saveAsObjectFiles,  
saveAsHadoopFiles



foreachRDD



# Actions

```
def print(num: Int): Unit = ssc.withScope {  
  def foreachFunc: (RDD[T], Time) => Unit = {  
    (rdd: RDD[T], time: Time) => {  
      val firstNum = rdd.take(num + 1)  
      // scalastyle:off println  
      println("-----")  
      println(s"Time: $time")  
      println("-----")  
      firstNum.take(num).foreach(println)  
      if (firstNum.length > num) println("...")  
      println()  
      // scalastyle:on println  
    }  
  }  
  foreachRDD(context.sparkContext.clean(foreachFunc), displayInnerRDDOps = false)  
}
```

# Actions – foreachRDD

```
dstream.foreachRDD{rdd =>
```

```
  rdd.cache()
```

```
  val alternatives = restServer.get("/v1/alternatives").toSet
```

```
  alternatives.foreach{alternative =>
```

```
    val byAlternative = rdd.filter(element => element.kind == alternative)
```

```
    val asRecords = byAlternative.map(element => asRecord(element))
```

```
    asRecords.foreachPartition{partition =>
```

```
      val conn = DB.connect(server)
```

```
      partition.foreach(element => conn.insert(element))
```

```
    }
```

```
  }
```

```
  rdd.unpersist(true)
```


Executes **local** on the Driver

Executes **distributed** on the Workers

# Actions – foreachRDD

```
dstream.foreachRDD{rdd =>
  rdd.cache()
  val alternatives = restServer.get("/v1/alternatives").toSet
  alternatives.foreach{alternative =>
    val byAlternative = rdd.filter(element => element.kind == alternative)
    val asRecords = byAlternative.map(element => asRecord(element))
    asRecords.foreachPartition{partition =>
      val conn = DB.connect(server)
      partition.foreach(element => conn.insert(element))
    }
  }
  rdd.unpersist(true)
}
```

Spark Cluster








Ready to Dive in?





## Self Contained Stream Generation

# ConstantInputDStream

```
/**  
 * An input stream that always returns the same RDD on each time step. Useful for testing.  
 */  
class ConstantInputDStream[T: ClassTag](_ssc: StreamingContext, rdd: RDD[T])  
  
// Usage  
val constantDStream = new ConstantInputDStream(streamingContext, rdd)
```

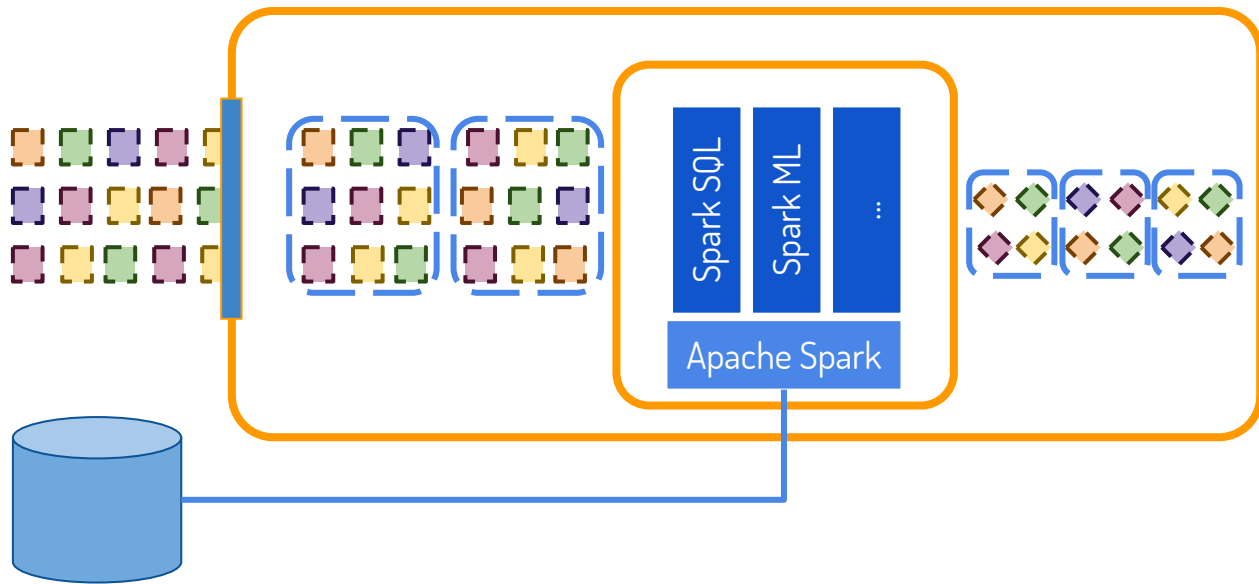
# ConstantInputDStream: Generate Data

```
import scala.util.Random
val sensorId: () => Int = () => Random.nextInt(sensorCount)
val data: () => Double = () => Random.nextDouble
val timestamp: () => Long = () => System.currentTimeMillis
// Generates records with Random data
val recordFunction: () => String = { () =>
  if (Random.nextDouble < 0.9) {
    Seq(sensorId().toString, timestamp(), data()).mkString(",")
  } else {
    // simulate 10% crap data as well... real world streams are seldom clean
    "!!~corrupt~^&##$"
  }
}
val sensorDataGenerator = sparkContext.parallelize(1 to n).map(_ => recordFunction)
```

RDD[() => Record]

```
val sensorData = sensorDataGenerator.map(recordFun => recordFun())
```

```
val rawDStream = new ConstantInputDStream(streamingContext, sensorData)
```



## Stream Enrichment with External Data

# ConstantInputDStream + foreachRDD= Reload External Data Periodically

```
var sensorReference = sparkSession.read.parquet(s"$referenceFile")
```

```
sensorRef.cache()
```

```
val refreshDStream = new ConstantInputDStream(streamingContext, sparkContext.emptyRDD[Int])
```

```
// Refresh data every 5 minutes
```

```
val refreshIntervalDStream = refreshDStream.window(Seconds(300), Seconds(300))
```

```
refreshIntervalDStream.foreachRDD{ _ =>
```

```
  sensorRef.unpersist(false)
```

```
  sensorRef = sparkSession.read.parquet(s"$referenceFile")
```

```
  sensorRef.cache()
```

```
}
```

# DStream + foreachRDD= Reload External Data with a Trigger

```
var sensorReference = sparkSession.read.parquet(s"$referenceFile")  
sensorRef.cache()
```

```
val triggerRefreshDStream: DStream = // create a DStream from a source. e.g. Kafka
```

```
val referenceStream = triggerRefreshDStream.transform { rdd =>  
  if (rdd.take(1) == "refreshNow") {  
    sensorRef.unpersist(false)  
    sensorRef = sparkSession.read.parquet(s"$referenceFile")  
    sensorRef.cache()  
  }  
  sensorRef.rdd  
}  
incomingStream.join(referenceStream) ...
```

  
**Spark**  
*Streaming*

  
**Spark**  
Structured Streaming



# ForeachRDD + Datasets + Functional = Structured Streaming Portability

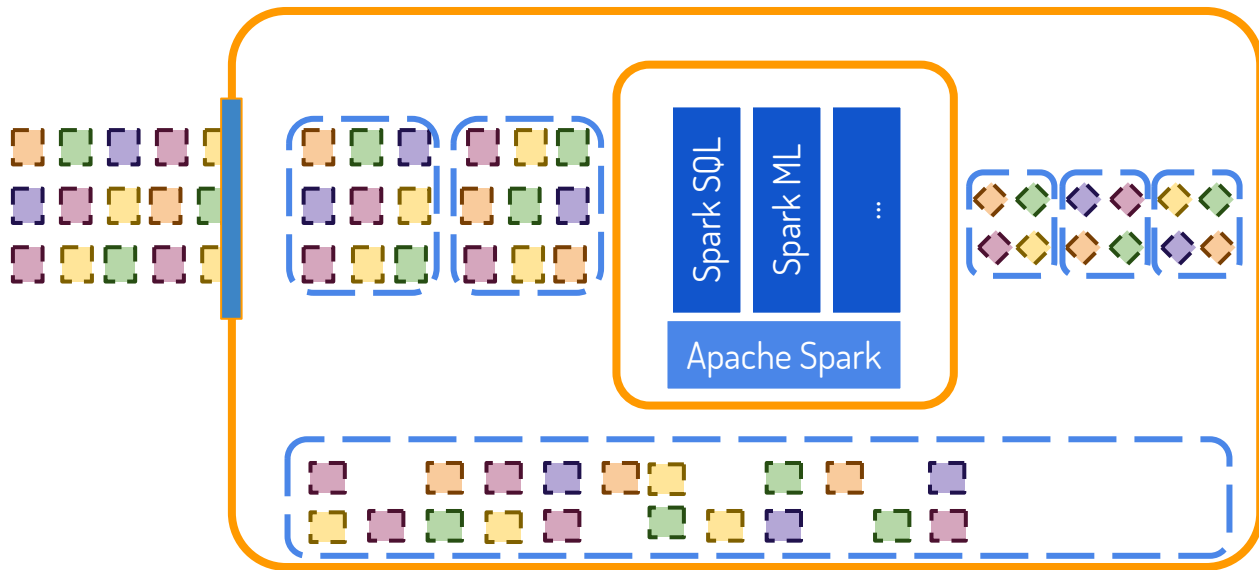
```
val parse: Dataset[String] => Dataset[Record] = ???  
val process: Dataset[Record] => Dataset[Result] = ???  
val serialize: Dataset[Result] => Dataset[String] = ???
```

## Spark Streaming

```
val dstream = KafkaUtils.createDirectStream(...)  
dstream.map{rdd =>  
  val ds = sparkSession.createDataset(rdd)  
  val f = parse andThen process andThen serialize  
  val result = f(ds)  
  result.write.format("kafka")  
    .option("kafka.bootstrap.servers", bootstrapServers)  
    .option("topic", writeTopic)  
    .option("checkpointLocation", checkpointLocation)  
    .save()  
}
```

## Structured Streaming

```
val kafkaStream = spark.readStream...  
val f = parse andThen process andThen serialize  
val result = f(kafkaStream)  
result.writeStream  
  .format("kafka")  
  .option("kafka.bootstrap.servers", bootstrapServers)  
  .option("topic", writeTopic)  
  .option("checkpointLocation", checkpointLocation)  
  .start()
```



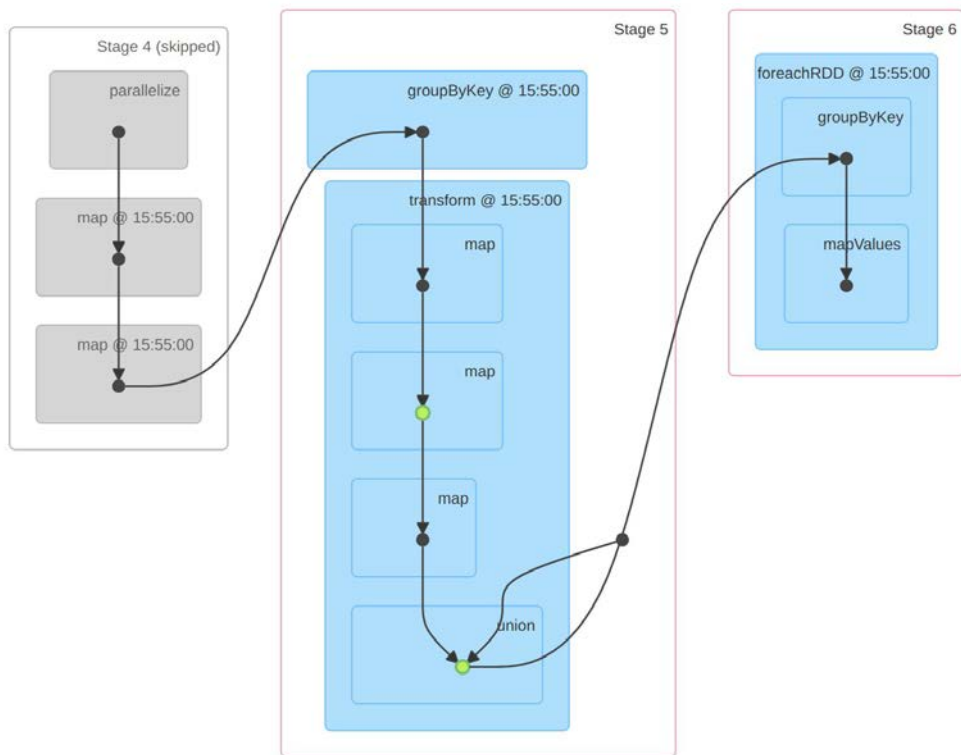
Keep Arbitrary State

# Keeping Arbitrary State


```
var baseline: Dataset[Features] = sparkSession.read.parquet(targetFile).as[Features]
...
stream.foreachRDD{ rdd =>
  val incomingData = sparkSession.createDataset(rdd)
  val incomingFeatures = rawToFeatures(incomingData)
  val analyzed = compare(incomingFeatures, baseline)
  // store analyzed data
  baseline = (baseline union incomingFeatures).filter(isExpired)
}
```

<https://gist.github.com/maasg/9d51a2a42fc831e385cf744b84e80479>

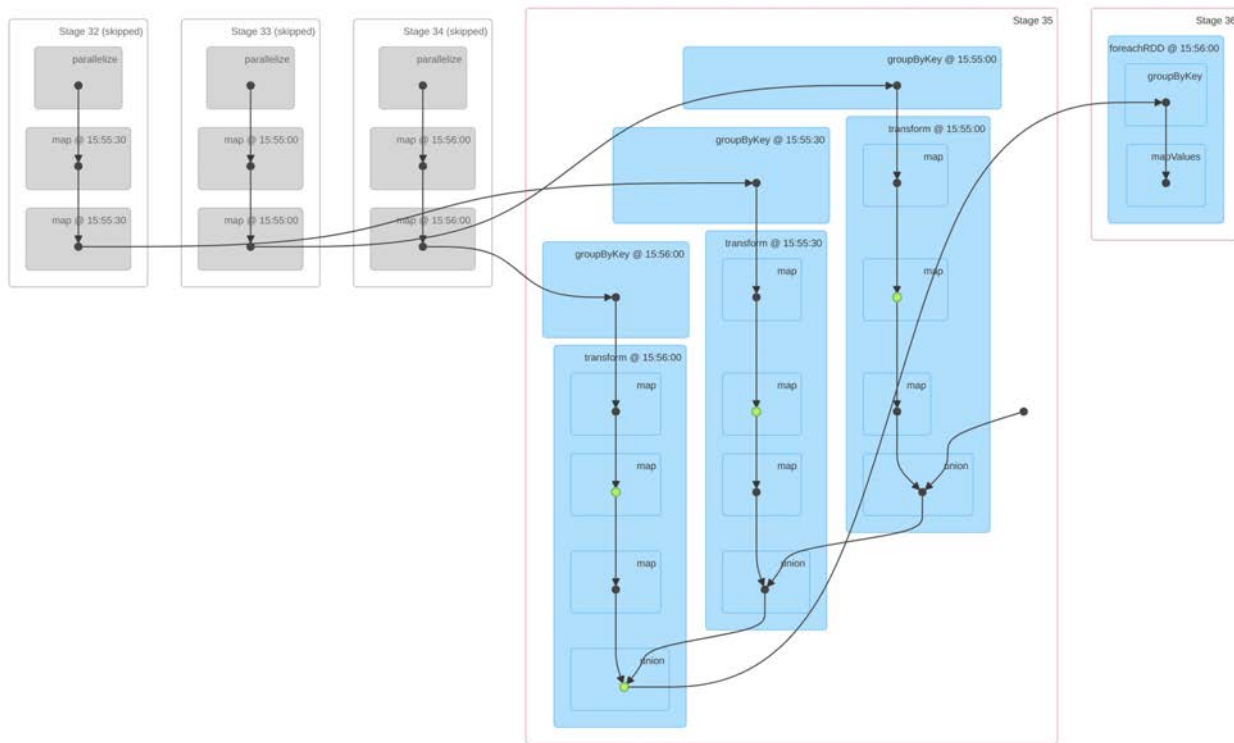
# Keeping Arbitrary State



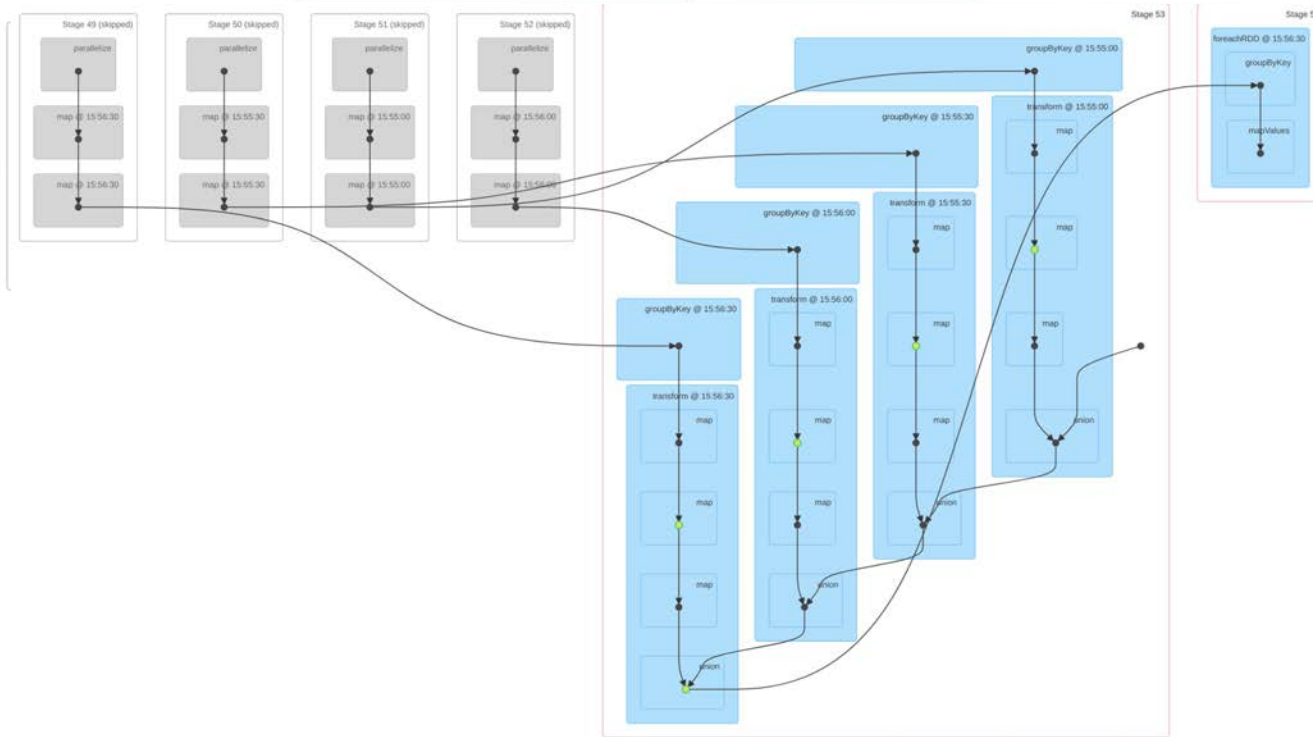
# Keeping Arbitrary State



# Keeping Arbitrary State

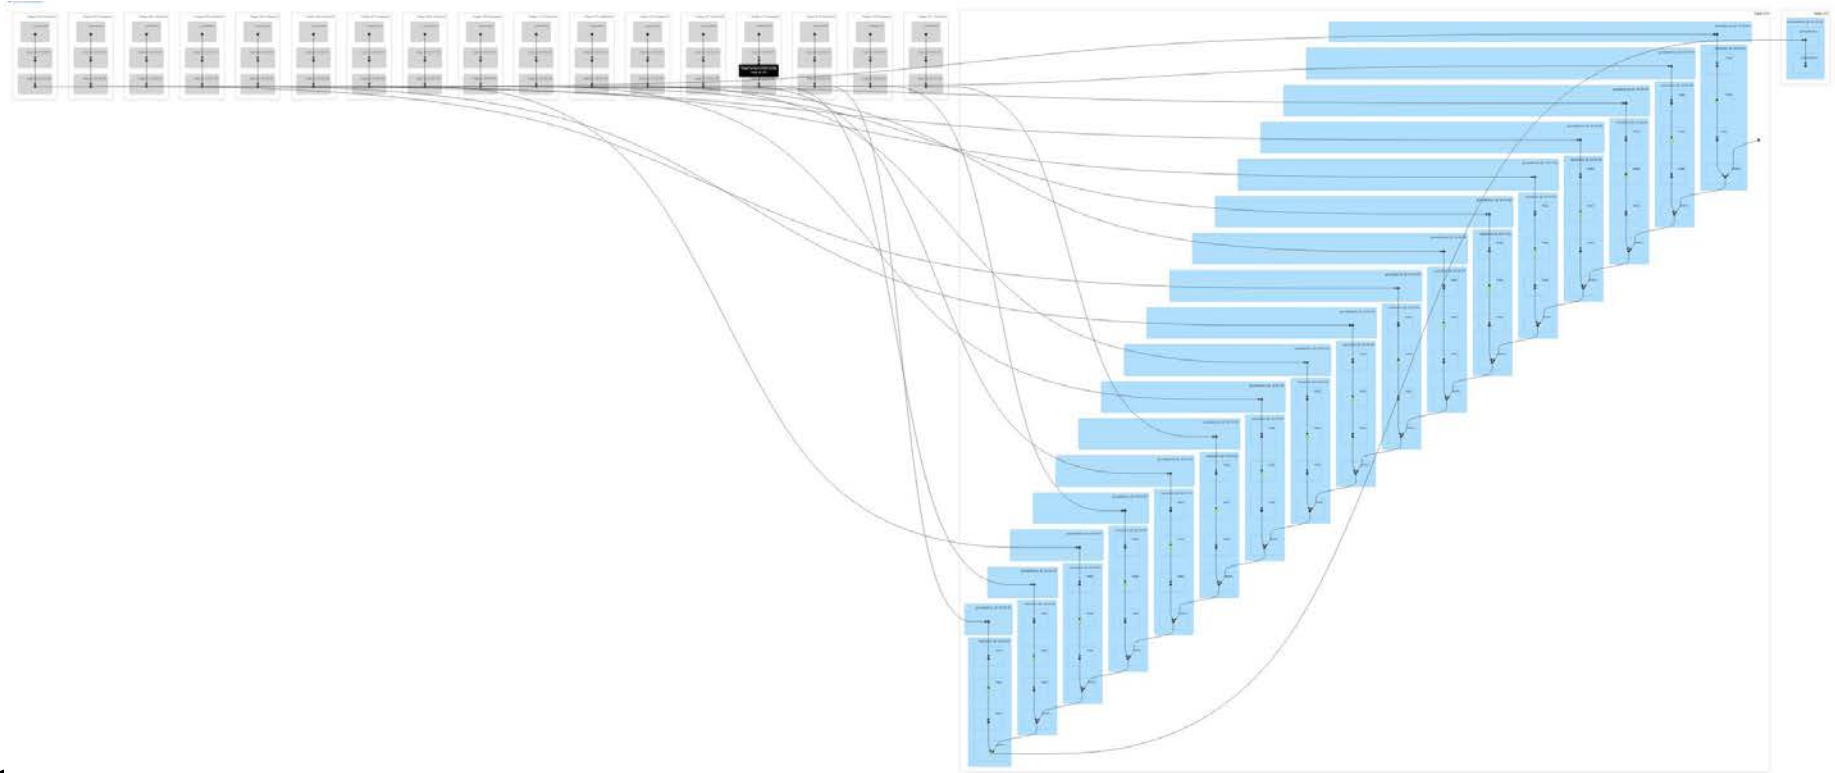


# Keeping Arbitrary State





# Keeping Arbitrary State




# Keeping Arbitrary State: Roll **your own** checkpoints !

```
var baseline: Dataset[Features] = sparkSession.read.parquet(targetFile).as[Features]


var cycle = 1

var checkpointFile = 0


stream.foreachRDD{ rdd =>
  val incomingData = sparkSession.createDataset(rdd)
  val incomingFeatures = rawToFeatures(incomingData)
  val analyzed = compare(incomingFeatures, baseline)
  // store analyzed data
  baseline = (baseline union incomingFeatures).filter(isOldFeature)
  cycle = (cycle + 1) % checkpointInterval
  if (cycle == 0) {
    checkpointFile = (checkpointFile + 1) % 2
    baseline.write.mode("overwrite").parquet(s"$targetFile_$checkpointFile")
    baseline = baseline.read(s"$targetFile_$checkpointFile")
  }
}
```




## Probabilistic Accumulators



# HyperLogLog: Cardinality Estimation



accuracy =  $1.054 / \sqrt{2^p}$



# HLL Accumulator

```
class HLLAccumulator[T](precisionValue: Int = 12) extends AccumulatorV2[T, Long] {  
  private def instance(): HyperLogLogPlus = new HyperLogLogPlus(precisionValue, 0)  
  
  override def add(v: T): Unit = hll.offer(v)  
  
  override def merge(other: AccumulatorV2[T, Long]): Unit = other match {  
    case otherHllAcc: HLLAccumulator[T] => hll.addAll(otherHllAcc.hll)  
    case _ => throw new UnsupportedOperationException(  
      s"Cannot merge ${this.getClass.getName} with ${other.getClass.getName}")  
  }  
}
```

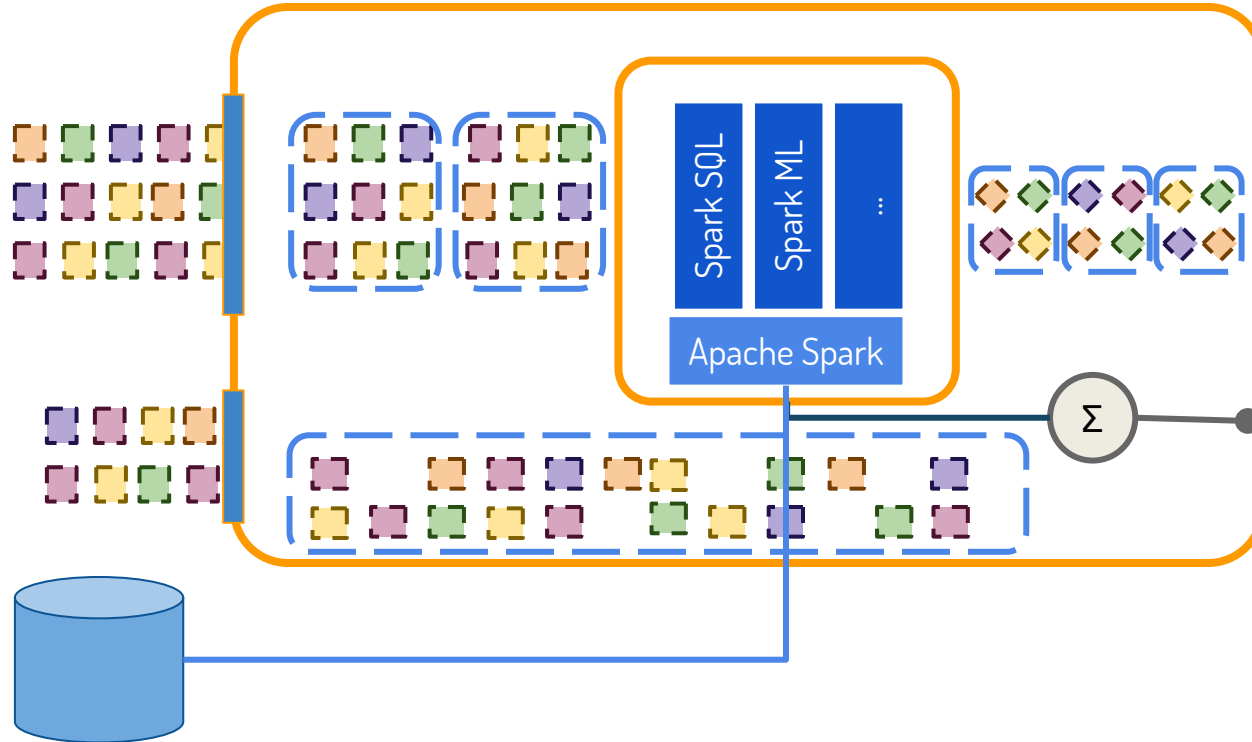
<https://github.com/LearningSparkStreaming/HLLAccumulator>

# Using Probabilistic Accumulators

```
import learning.spark.streaming.HLLAccumulator
val uniqueVisitorsAccumulator= new HLLAccumulator[String](precisionValue = 12)
sc.register(uniqueVisitorsAccumulator, "unique-visitors")
...
clickStream.foreachRDD{rdd =>
  rdd.foreach{
    case BlogHit(ts, user, url) => uniqueVisitorsAccumulator.add(user)
  }
  ...
  val currentUniqueVisitors = uniqueVisitorsAccumulator.value
  ...
}
```



# Putting it all Together



# Questions?

# Thank You

 @maasg

 Lightbend