



Spark Structured Streaming Helps Smart Manufacturing

Xiaochang Wu

xiaochang.wu@intel.com

Hao Cheng

hao.cheng@intel.com

Intel Big Data Engineering Team

#EUstr1

Who am I?

- Currently working at Intel Big Data Engineering team based in Shanghai, China
- Joined Intel from 2006, focus on performance optimization for CPU / GPU / System / App for ~10 years
- Delivering the best Spark performance on Intel platforms
- Building reference big data platform for Intel customers

Agenda


- Motivation
- Architecture and Features Overview
- Adopting Spark Structured Streaming
- Improvement Proposals

Motivation

- Building Spark-centric IOT Data Platform for PCB Manufacturing
 - Data collection, processing, analytics, visualization and alerting
 - Short-term goal: Enhance predictive fault repair and material tracking efficiency
- Take advantage of latest Structured Streaming
 - Streaming ETL
 - Stateless and Stateful processing
- Learn from real-world problems

Overview

- Transform the PCB surface mount lines

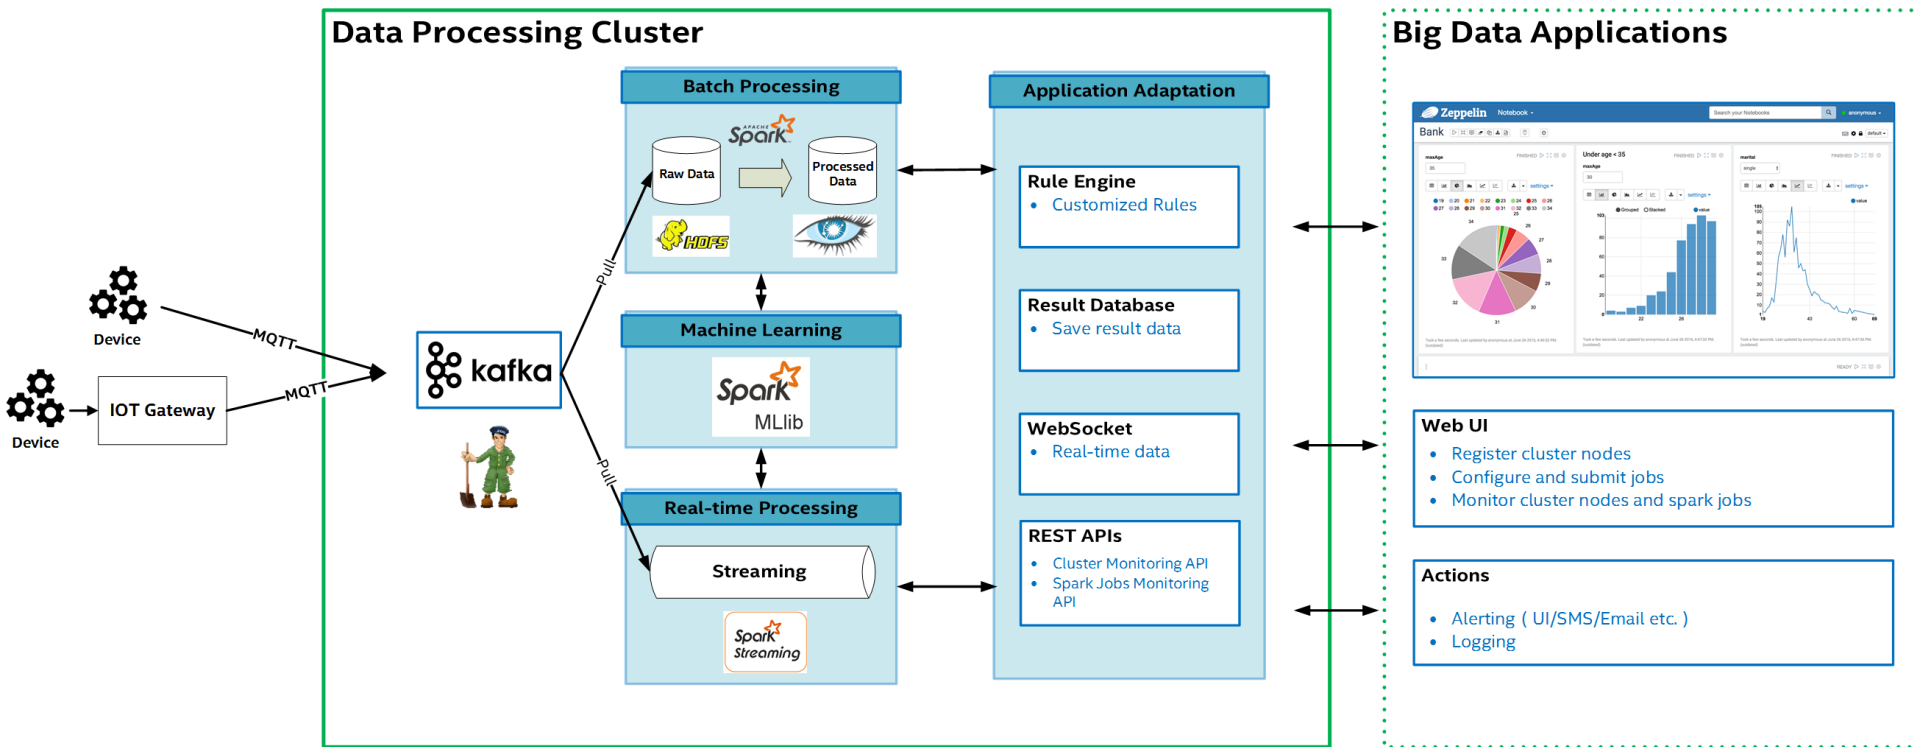


- SPI: Solder Paste Inspection
- AOI: Automated Optical Inspection

Data Platform Features

- Device Management
 - Register and manage devices
 - RPC commands to access device attributes
- Data Collection
 - Collect and store device data in reliable way
- Data Processing & Analytics
 - Domain-specific scripts for both batch and real-time processing
 - Support using Machine Learning algorithms for advanced analytics
- Rule Engine & Actions
 - Process incoming data with flexible user-defined rules.
 - Trigger alarms or send commands to devices using rules.
- Data Visualization
 - Customizable user-friendly UI & Dashboards
 - Generate interactive data reports to fulfil business needs
- Multi-tenancy, Scale-out, Fault-tolerance & Security

Data Platform Architecture



*Other names and brands may be claimed as the property of others.

Data Characteristics

- Data Sources
 - Manufacturing Execution System (MES)
 - Sensor, Camera, Log
 - OCR for machine screen output
- Structured, semi-structured and unstructured data
 - Low frequency but large amount
 - E.g. Picture and video
 - High frequency continuous data
 - Not large amount per unit. Total amount is very large
 - E.g. vibration sensor data used to detect the quality of the equipment.

Spark Structured Streaming

- Built on the Spark SQL engine
 - Express streaming the same way as batch
 - Dataset/DataFrame API in Scala, Java, Python or R
- Incrementally and continuously updating the final result
 - Handling Event-time and Late Data
 - Delivering end-to-end exactly-once fault tolerance semantics
 - Output modes: Append, Complete, Update mode
- Support Stateless and Stateful operations

Fast, scalable, fault-tolerant, end-to-end exactly-once

Streaming ETL

- Data are produced from sources as JSON objects, transferred with MQTT and saved to Kafka
- JSON objects are complex and have nested structures
- Use Spark SQL functions to extract
 - *selectExpr, get_json_object, cast, alias*

```
val ds = dataframe.selectExpr("CAST(value AS STRING)").as[String]
val parsedDF = ds.select(
  get_json_object($"value", "$.key").alias("key"),
  get_json_object($"value", "$.idno").cast("int").alias("idno"),
  get_json_object($"value", "$.ts").cast("timestamp").alias("ts"),
  get_json_object($"value", "$.OffsetX").cast("double").alias("OffsetX"),
  get_json_object($"value", "$.OffsetY").cast("double").alias("OffsetY")
)
```

Stateless and Stateful Operations

- Stateless
 - Projection, Filter
- Stateful (Implicit)
 - Aggregation
 - Time window
 - Join
- Advanced arbitrary stateful operations
 - `mapGroupsWithState` / `flatMapGroupsWithState`

Structured Streaming Considerations

- Output mode: **append, update, complete**
- Handling late data
 - Method 1: use time window and automatic watermark
 - Method 2: use manual watermark
- Post-processing
 - Debug: console sink / memory sink
 - File sink / Foreach sink
 - Collect to driver side and process

Case 1: Find out offset AVG and STDDEV in Solder Paste Inspection (SPI)

```
val resultDS = dataframe
  .groupBy("idno")
  .agg(mean("OffsetX").as("meanX"),
        mean("OffsetY").as("meanY"),
        stddev("OffsetX").as("stddevX"),
        stddev("OffsetY").as("stddevY"),
        count("idno").as("recordsCount"))
  .filter($"idno".isNotNull)
  .orderBy("idno")
  .as[ResultRecord]
```

Case 2: Find out machine off time from event logs (1)

- Track sessions from streams of events and find out machine off time
 - mapGroupsWithState / flatMapGroupsWithState
 - Off time = Duration from status “off” to next “on”

```
{"device_id":1, "timestamp":"2017-01-07 15:05:00", "status":"off"}
{"device_id":2, "timestamp":"2017-01-07 15:02:00", "status":"on" }
{"device_id":3, "timestamp":"2017-01-07 15:04:00", "status":"off"}
{"device_id":1, "timestamp":"2017-01-07 15:02:00", "status":"on" }
{"device_id":1, "timestamp":"2017-01-07 15:01:00", "status":"on" }
{"device_id":3, "timestamp":"2017-01-07 15:03:00", "status":"off"}
{"device_id":2, "timestamp":"2017-01-07 15:03:00", "status":"on" }
{"device_id":2, "timestamp":"2017-01-07 15:04:00", "status":"off"}
{"device_id":2, "timestamp":"2017-01-07 15:08:00", "status":"on" }
{"device_id":3, "timestamp":"2017-01-07 15:10:00", "status":"off"}
```

Case 2: Find out machine off time from event logs (2)

- Event: define data types for events
- SessionInfo: Session information as state
- SessionUpdate: Update information returned as result

```
case class Event(device_id: Long, timestamp: Timestamp, machine_status: String)

case class SessionInfo(offTime: Long,           // Total machine off time so far
                       currentTime: Timestamp,  // Latest event timestamp from last update
                       currentStatus: String,    // Latest machine status from last update: "on" or "off"
                       eventList: List[Event])   // Current event list

case class SessionUpdate(deviceId: Long,        // Device ID
                          offTime: Long,        // Total machine off time
                          expired: Boolean)      // If session is expired
```

```

val sessionUpdate = IoTEvents
    .groupByKey(event => event.device_id)
    .mapGroupsWithState[SessionInfo, SessionUpdate](GroupStateTimeout.ProcessingTimeTimeout) {
        case (deviceId: Long, events: Iterator[Event], state: GroupState[SessionInfo]) =>
            // If timed out, then remove session and send final update
            if (state.hasTimedOut) {
                val finalUpdate =
                    SessionUpdate(deviceId, state.get.offTime, expired = true)
                state.remove()
                finalUpdate
            } else {
                val eventList = events.toList

                val updatedSessionInfo = if (state.exists) { // update existing state
                    // Get existing state
                    val oldSession = state.get
                    // Update and sort the event list based on timestamp
                    val eventList_new = (oldSession.eventList ::: eventList).sortWith(_.event_time.getTime < _.event_time.getTime)

                    //
                    // Iterate thru event list to calculate off time, then update to updatedSessionInfo
                    //

                } else { // setup initial state
                    // Sort the event list based on timestamp
                    val eventList_new = eventList.sortWith(_.event_time.getTime < _.event_time.getTime)

                    //
                    // Iterate thru event list to calculate off time, then set initial value to updatedSessionInfo
                    //

                } // end if (state.exists)
                state.update(updatedSession)
                // Set timeout such that the session will be expired if no data received for 5 minutes
                state.setTimeoutDuration("5 minutes")
                SessionUpdate(deviceId, state.get.offTime, expired = false)
            } // end if (state.hasTimedOut)
    } // end mapGroupsWithState

```


Case 2: Find out machine off time from event logs (3)

- Event time ordering
 - Every subsequent event depends on previous event
 - Required to sort events based on timestamp due to logical dependence
- Handling late events
 - Late data will break previous calculation result, forced to maintain all events and recalculate
 - Solution: manual watermarking for late data, remove old events if timeout, incremental calculation

Improvement Proposals

- Flexible window and trigger operation
- Complex Event Processing (CEP) APIs
 - Stream Join
 - Session Tracking
 - Pattern
 - Alert
- Streaming Domain Specific Language (DSL)
 - SQL-like language for Streams
- We are working on them!

See Also from Intel team

- Today 14:00 @ AUDITORIUM

An Adaptive Execution Engine For Apache Spark SQL
by Carson Wang

- Tomorrow 11:00 @ LIFFEY HALL

FPGA-Based Acceleration Architecture for Spark SQL
by Qi Xie

Legal Notices and Disclaimers

- You may not use or facilitate the use of this document in connection with any infringement or other legal analysis concerning Intel products described herein. You agree to grant Intel a non-exclusive, royalty-free license to any patent claim thereafter drafted which includes subject matter disclosed herein.
- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document.
- Intel disclaims all express and implied warranties, including without limitation, the implied warranties of merchantability, fitness for a particular purpose, and non-infringement, as well as any warranty arising from course of performance, course of dealing, or usage in trade.

References

- Surface mount process: <http://www.surfacemountprocess.com/>
- Structured Streaming Programming Guide:
<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>
- Five Spark SQL Utility Functions to Extract and Explore Complex Data Types:
<https://databricks.com/blog/2017/06/13/five-spark-sql-utility-functions-extract-explore-complex-data-types.html>
- Spark Structured Sessionization example:
<https://github.com/apache/spark/blob/v2.2.0/examples/src/main/scala/org/apache/spark/examples/sql/streaming/StructuredSessionization.scala>