

Introduction to DynamoDB

<http://bit.ly/NoSQLDesignPatterns>

Rick Houlihan, Principal Solutions Architect

February 2016

What to expect from the session

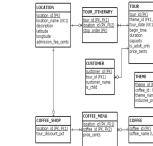
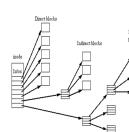
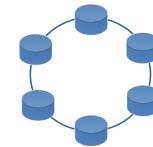
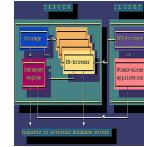
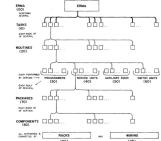
- Brief history of data processing
- DynamoDB Internals
 - Tables, API, data types, indexes
 - Scaling and data modeling
- Design patterns and best practices
- Event driven applications and DDB Streams

What is a Database?

“A place to put stuff my app needs.”

– Average Developer

Timeline of Database Technology



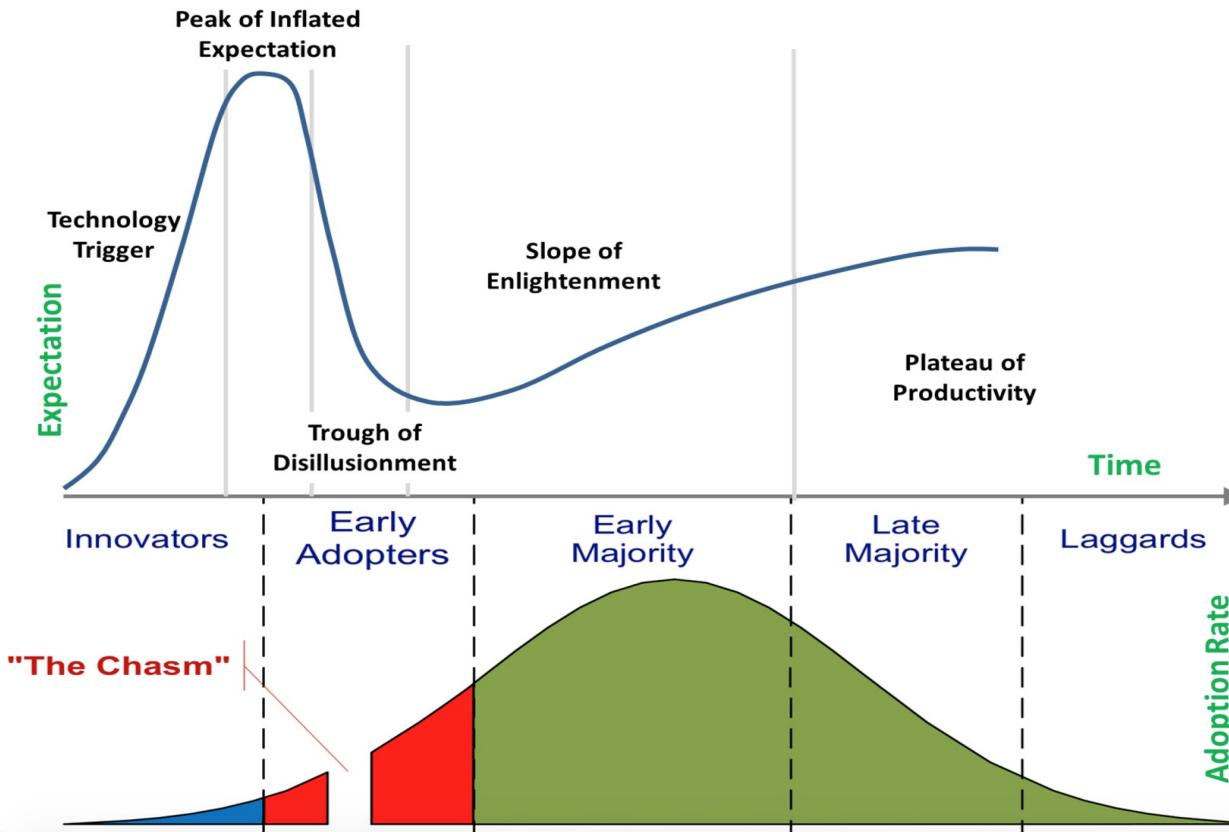
Data Volume Since 2010

- 90% of stored data generated in last 2 years
- 1 Terabyte of data in 2010 equals 6.5 Petabytes today
- Linear correlation between data pressure and technical innovation
- No reason these trends will not continue over time

5

Historical Current

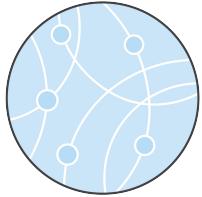
Technology Adoption and the Hype Curve



Why NoSQL?

SQL	NoSQL
Optimized for storage	Optimized for compute
Normalized/relational	Denormalized/hierarchical
Ad hoc queries	Instantiated views
Scale vertically	Scale horizontally
Good for OLAP	Built for OLTP at scale

It's all about aggregations...



Social Network



Document Management



Process Control

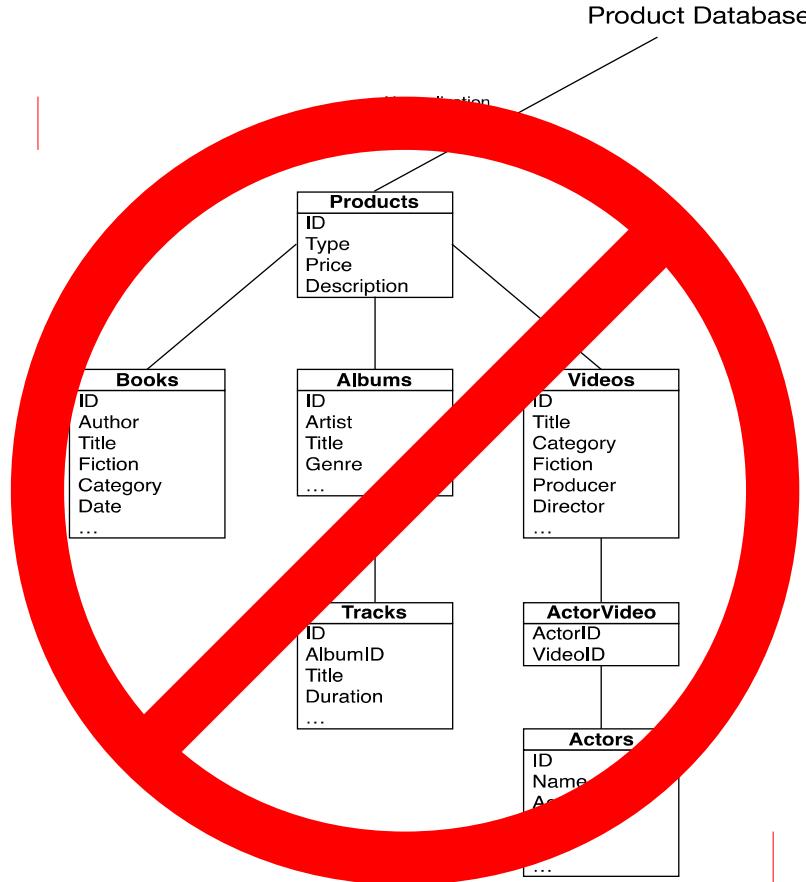


IT Monitoring

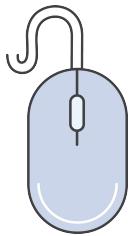


Data Trees

SQL vs. NoSQL Access Pattern



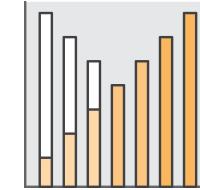
Amazon DynamoDB



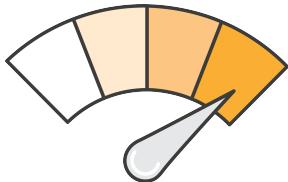
Fully Managed NoSQL



Document or Key-Value



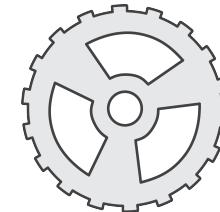
Scales to Any Workload



Fast and Consistent

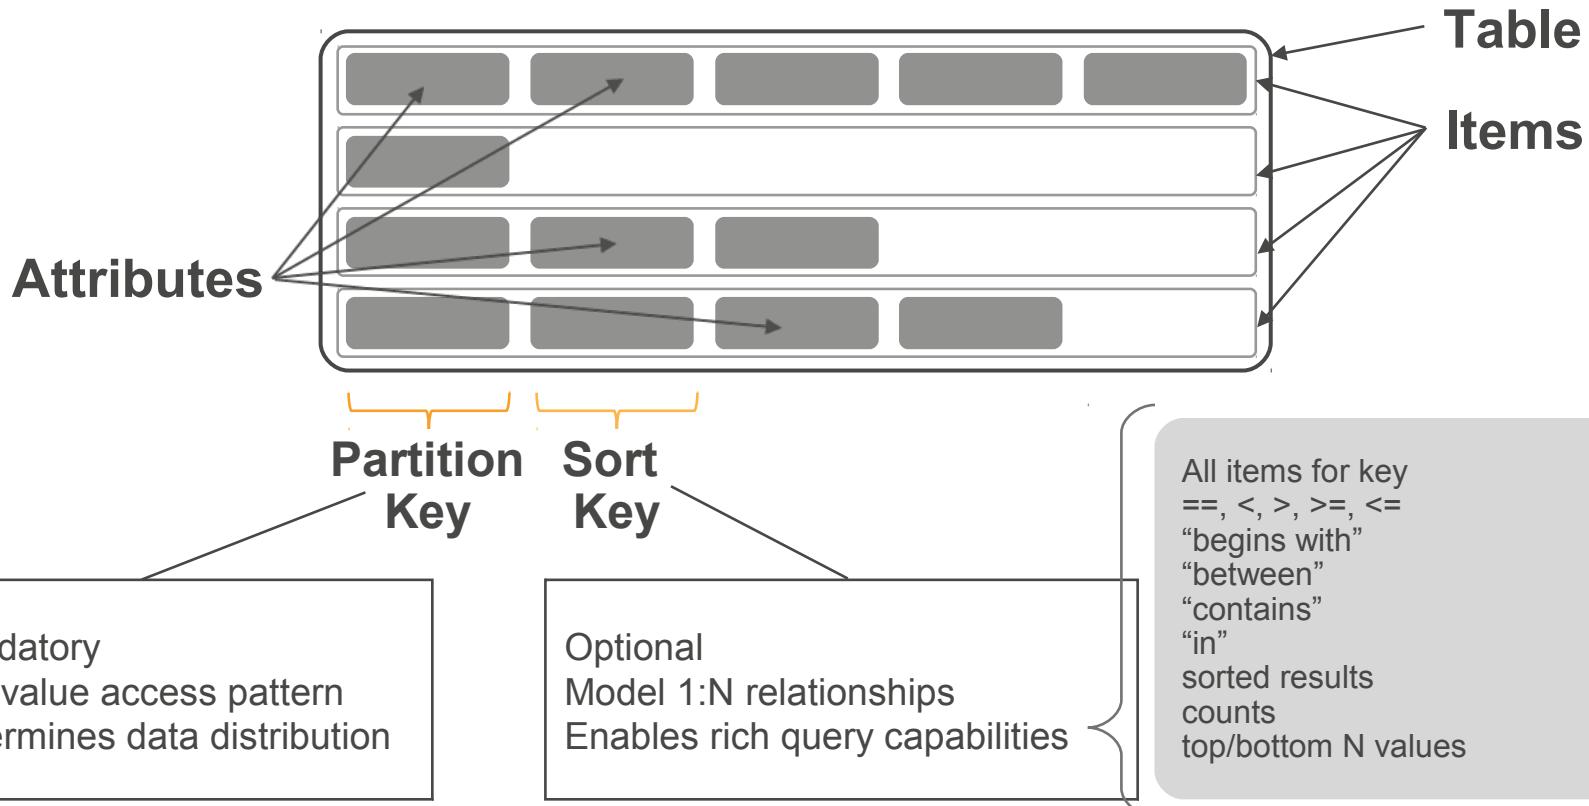


Access Control



Event Driven Programming

Table

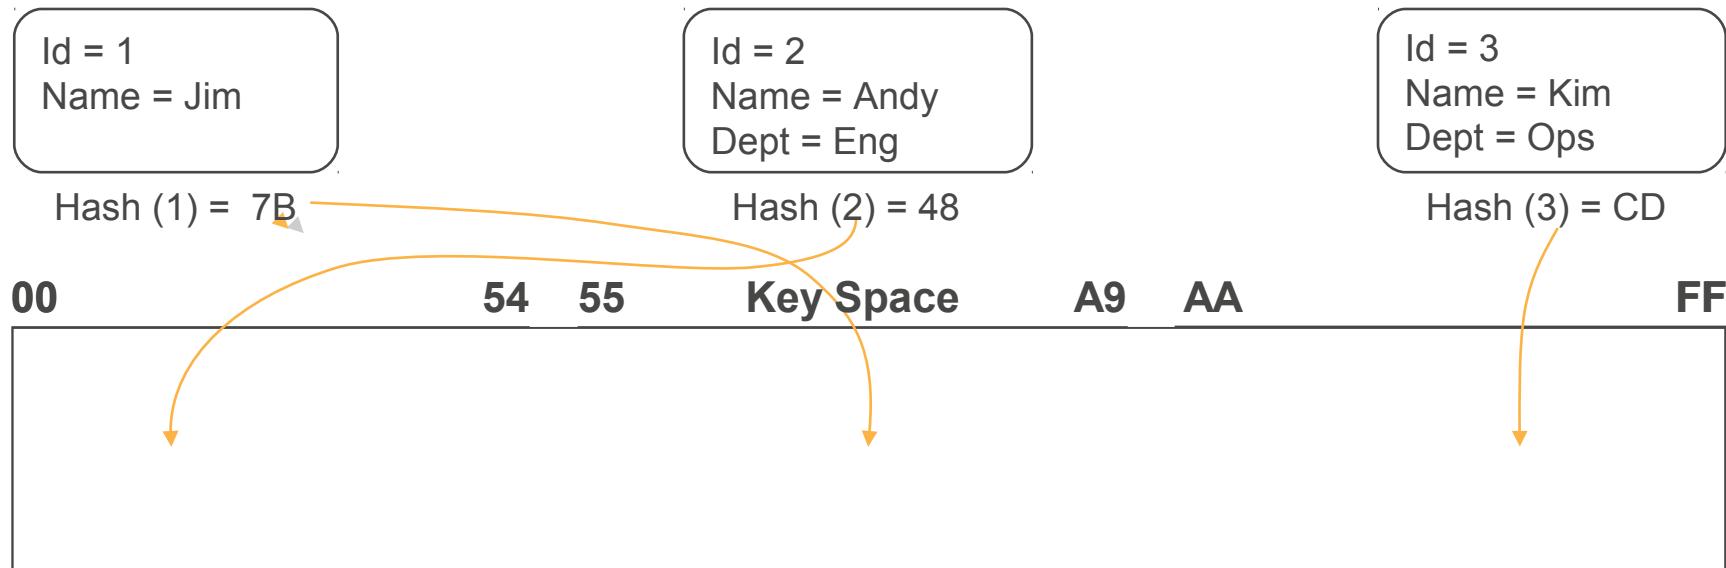


Partition Keys

Partition Key uniquely identifies an item

Partition Key is used for building an unordered hash index

Allows table to be partitioned for scale



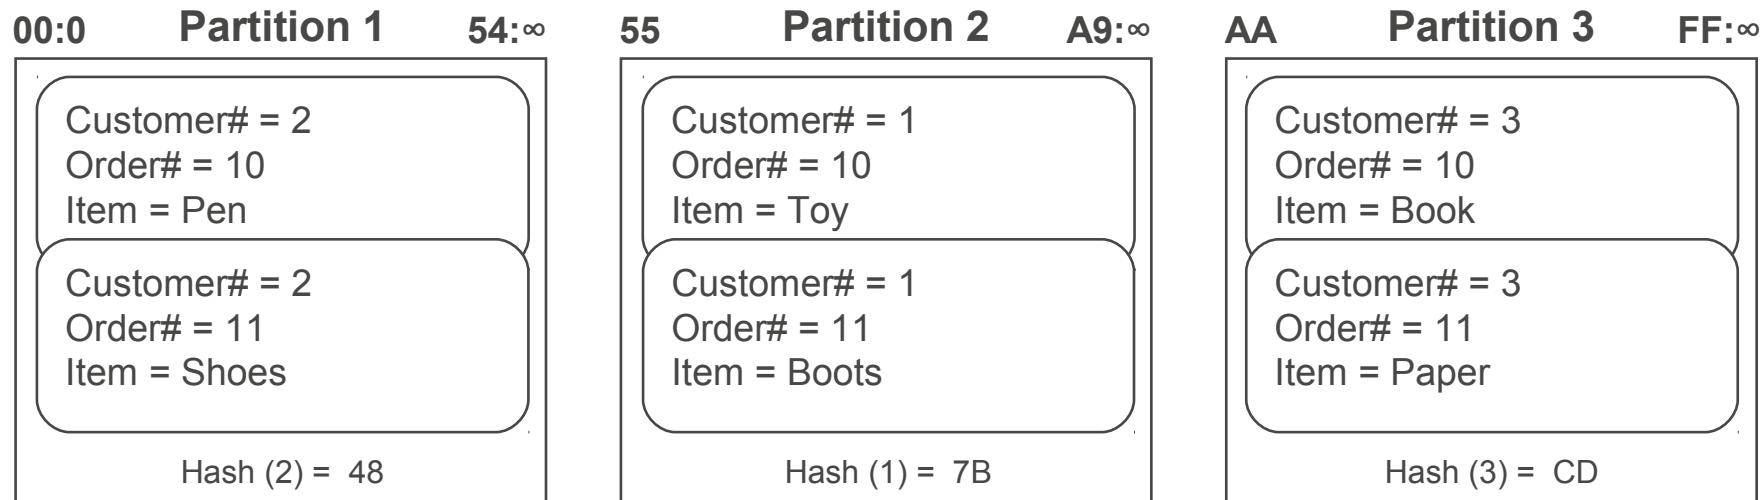
Partition:Sort Key

Partition:Sort Key uses two attributes together to uniquely identify an Item

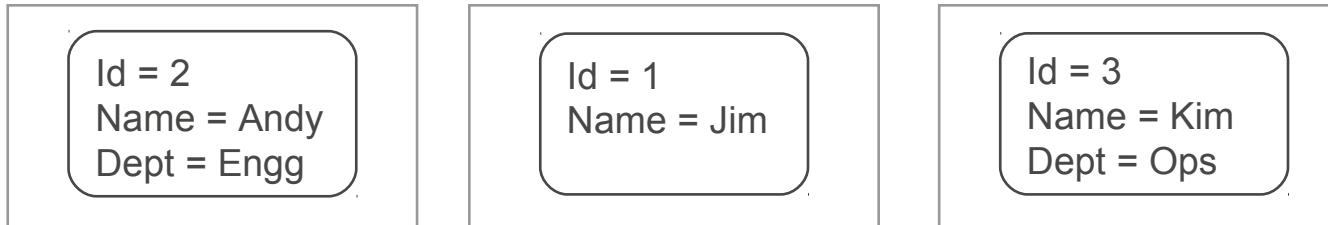
Within unordered hash index, data is arranged by the sort key

No limit on the number of items (∞) per partition key

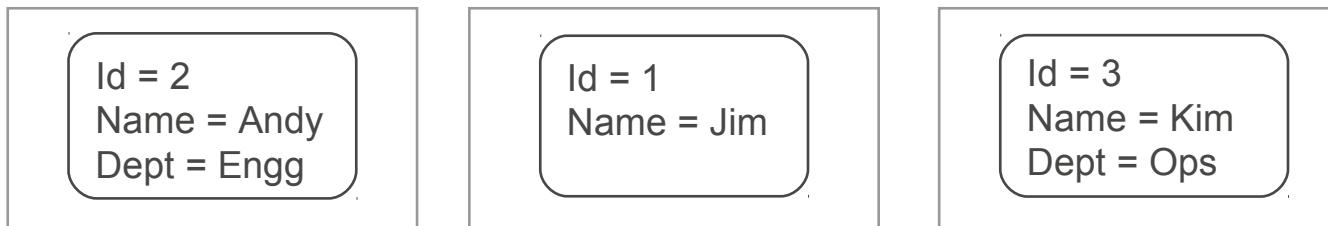
- Except if you have local secondary indexes



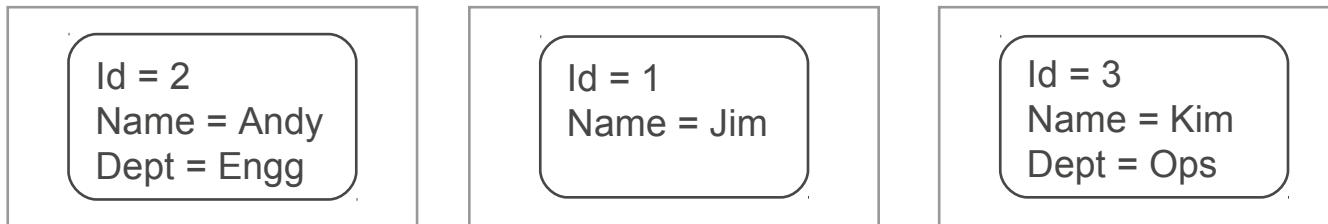
Partitions are three-way replicated



Replica 1



Replica 2



Replica 3

Partition 1

Partition 2

Partition N

Indexes

Local secondary index (LSI)

Alternate sort key attribute

Index is local to a partition key

Table

A1 (partition)	A2 (sort)	A3	A4	A5
-------------------	--------------	----	----	----

10 GB max per partition key, i.e.
LSIs limit the # of range keys!

LSIs

A1 (partition)	A3 (sort)	A2 (item key)
-------------------	--------------	------------------

KEYS_ONLY

A1 (partition)	A4 (sort)	A2 (item key)	A3 (projected)
-------------------	--------------	------------------	-------------------

INCLUDE A3

A1 (partition)	A5 (sort)	A2 (item key)	A3 (projected)	A4 (projected)
-------------------	--------------	------------------	-------------------	-------------------

ALL

Global secondary index (GSI)

Online indexing

Alternate partition and/or sort key

Index is across all partition keys

Use composite sort keys for compound indexes

Table

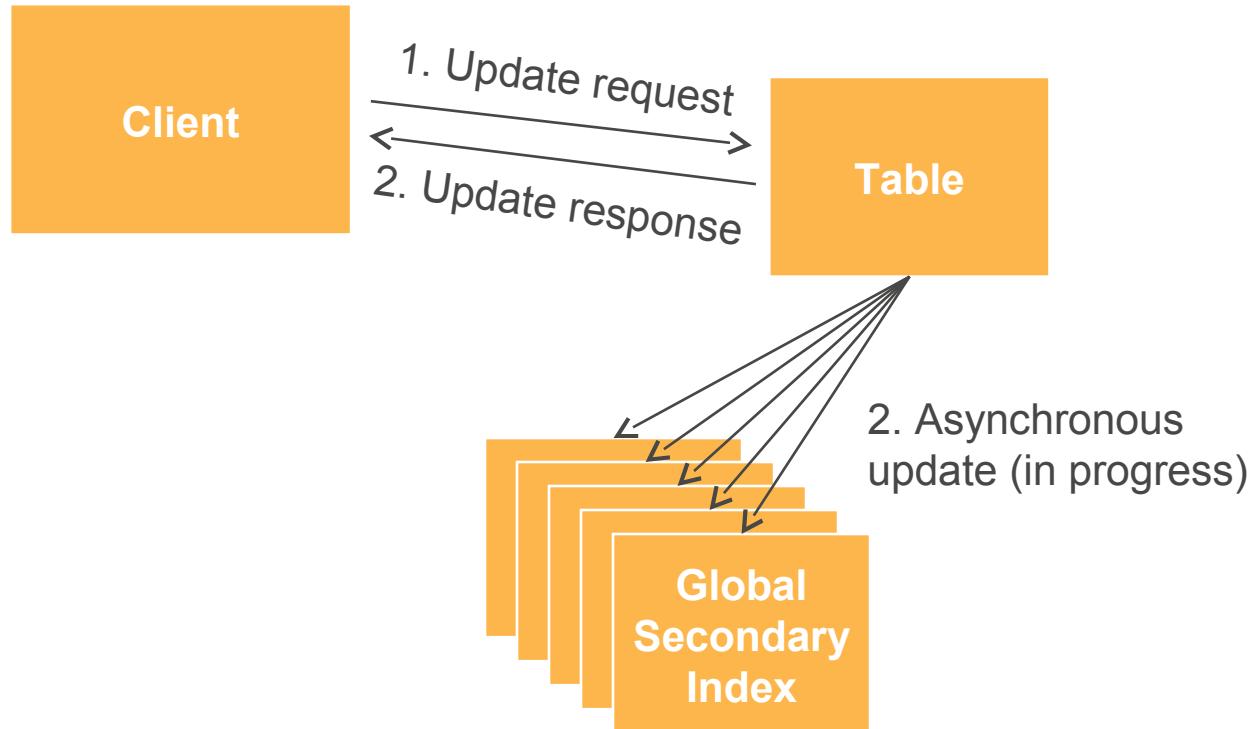
A1 (partition)	A2	A3	A4	A5
-------------------	----	----	----	----

RCUs/WCUs provisioned
separately for GSIs

GSI

A2 (partition)	A1 (itemkey)			KEYS_ONLY
A5 (partition)	A4 (sort)	A1 (item key)	A3 (projected)	INCLUDE A3
A4 (partition)	A5 (sort)	A1 (item key)	A2 (projected)	A3 (projected) ALL

How do GSI updates work?



If GSIs don't have enough write capacity, table writes will be throttled!

LSI or GSI?

LSI can be modeled as a GSI

If data size in an item collection > 10 GB, use GSI

**If eventual consistency is okay for your scenario,
use GSI!**

Scaling

Scaling

Throughput

- Provision any amount of throughput to a table

Size

- Add any number of items to a table
 - Max item size is 400 KB
 - LSIs limit the number of range keys due to 10 GB limit

Scaling is achieved through partitioning

Throughput

Provisioned at the table level

- Write capacity units (WCUs) are measured in 1 KB per second
- Read capacity units (RCUs) are measured in 4 KB per second
 - RCUs measure strictly consistent reads
 - Eventually consistent reads cost 1/2 of consistent reads

Read and write throughput limits are independent



RCU



WCU

Partitioning Math

Number of Partitions	
By Capacity	$(\text{Total RCU} / 3000) + (\text{Total WCU} / 1000)$
By Size	$\text{Total Size} / 10 \text{ GB}$
Total Partitions	$\text{CEILING}(\text{MAX}(\text{Capacity, Size}))$

In the future, these details might change...

Partitioning Example

Table size = 8 GB, RCUs = 5000, WCUs = 500

Number of Partitions	
By Capacity	$(5000 / 3000) + (500 / 1000) = 2.17$
By Size	$8 / 10 = 0.8$
Total Partitions	$\text{CEILING}(\text{MAX} (2.17, 0.8)) = 3$

RCUs and WCUs are uniformly spread across partitions

RCUs per partition = $5000/3 = 1666.67$
WCUs per partition = $500/3 = 166.67$
Data/partition = $10/3 = 3.33 \text{ GB}$

What causes throttling?

If sustained throughput goes beyond provisioned throughput per partition

Non-uniform workloads

- Hot keys/hot partitions
- Very large bursts

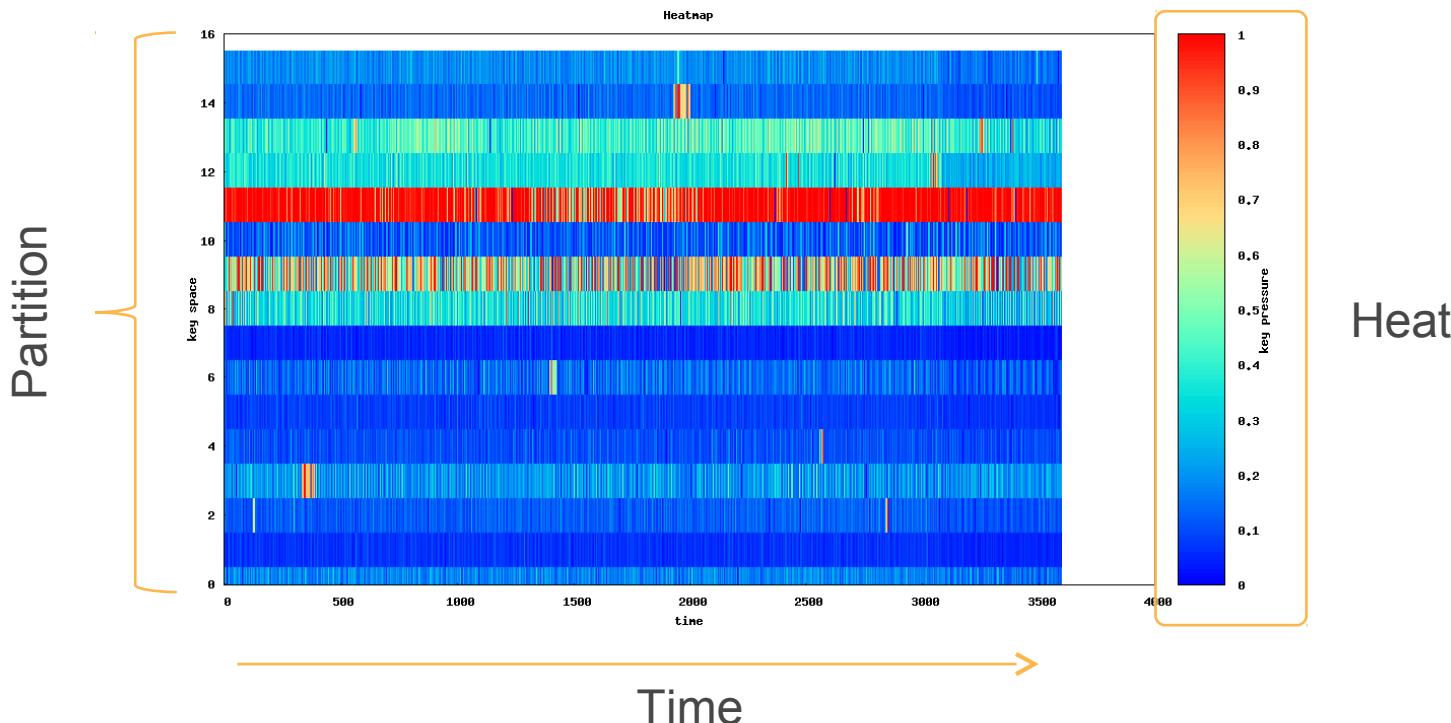
Mixing hot data with cold data

- Use a table per time period

From the example before:

- Table created with 5000 RCUs, 500 WCUs
- RCUs per partition = 1666.67
- WCUs per partition = 166.67
- If sustained throughput > (1666 RCUs or 166 WCUs) per key or partition, DynamoDB may throttle requests
 - Solution: Increase provisioned throughput

What bad NoSQL looks like...



Getting the most out of DynamoDB throughput

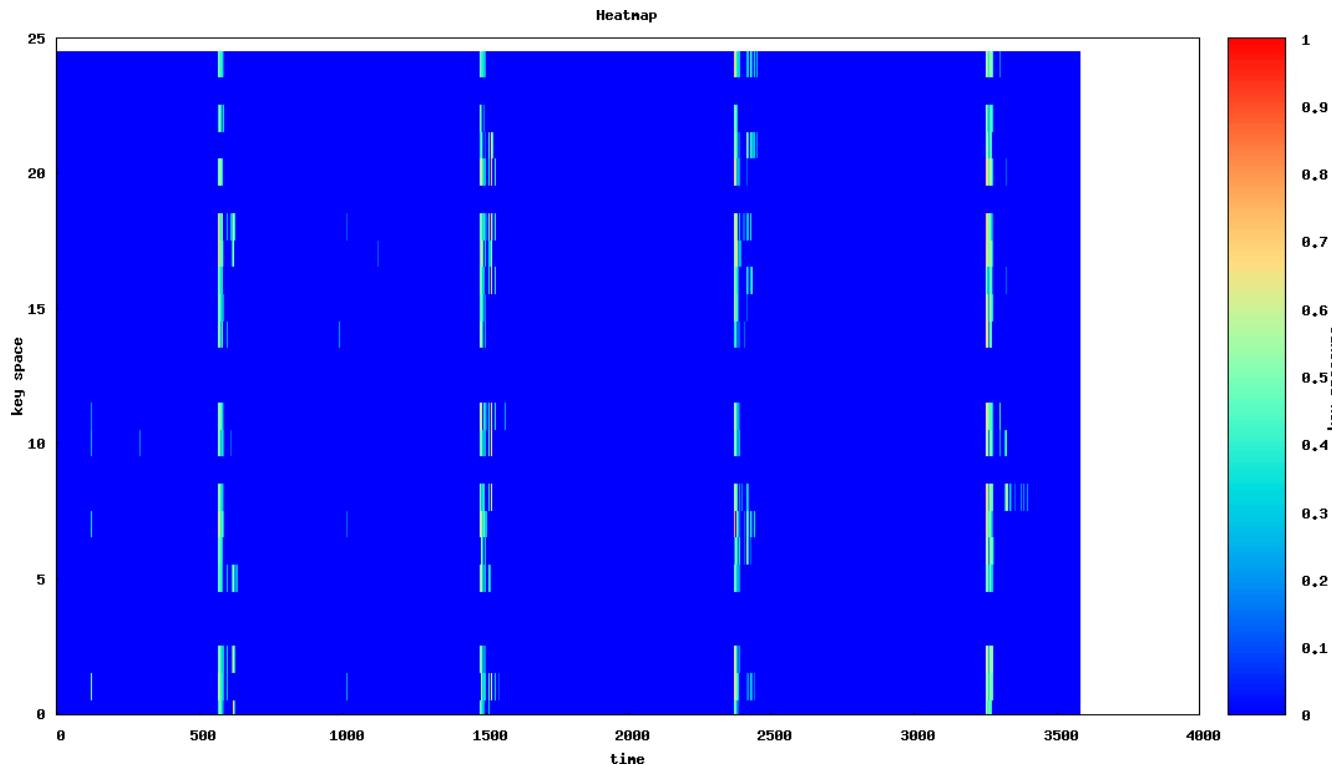
“To get the most out of DynamoDB throughput, create tables where the hash key element has a large number of distinct values, and values are requested fairly uniformly, as randomly as possible.”

—*DynamoDB Developer Guide*

Space: access is evenly spread over the key-space

Time: requests arrive evenly spaced in time

Much better picture...



Data Modeling

1:1 relationships or key-values

Use a table or GSI with an alternate partition key

Use GetItem or BatchGetItem API

Example: Given an SSN or license number, get attributes

Users Table	
Partition key	Attributes
SSN = 123-45-6789	Email = johndoe@nowhere.com, License = TDL25478134
SSN = 987-65-4321	Email = maryfowler@somewhere.com, License = TDL78309234

Users-License-GSI	
Partition key	Attributes
License = TDL78309234	Email = maryfowler@somewhere.com, SSN = 987-65-4321
License = TDL25478134	Email = johndoe@nowhere.com, SSN = 123-45-6789

1:N relationships or parent-children

Use a table or GSI with partition and sort key

Use Query API

Example: Given a device, find all readings between epoch X, Y

Device-measurements		
Partition Key	Sort key	Attributes
DeviceId = 1	epoch = 5513A97C	Temperature = 30, pressure = 90
DeviceId = 1	epoch = 5513A9DB	Temperature = 30, pressure = 90

N:M relationships

Use a table and GSI with partition and sort key elements switched

Use Query API

Example: Given a user, find all games. Or given a game, find all users.

User-Games-Table	
Partition Key	Sort key
UserId = bob	GameId = Game1
UserId = fred	GameId = Game2
UserId = bob	GameId = Game3

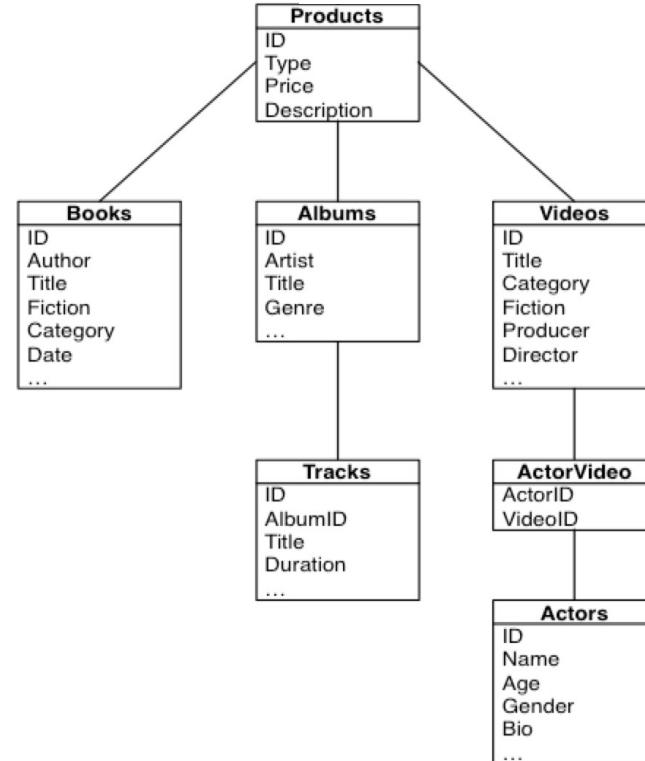
Game-Users-GSI	
Partition Key	Sort key
GameId = Game1	UserId = bob
GameId = Game2	UserId = fred
GameId = Game3	UserId = bob

Hierarchical Data

Tiered relational data structures

How OLTP Apps Use Data

- Mostly Hierarchical Structures
- Entity Driven Workflows
- Data Spread Across Tables
- Requires Complex Queries



Hierarchical Structures as Items...

Use composite sort key to define a Hierarchy

Highly selective result sets with sort queries

Index anything, scales to any size

Items	Primary Key		Attributes						
	ProductID	type	title	author	genre	publisher	datePublished	ISBN	
1	bookID	Ringworld	Larry Niven	Science Fiction	Ballantine	Oct-70	0-345-02046-4		
		Dark Side of the Moon	Pink Floyd	Progressive Rock	Harvest	Abbey Road	3/1/73	Pink Floyd	
	albumID	Speak to Me	1:30	Mason	Instrumental				
		Breathe	2:43	Waters, Gilmour, Wright	Gilmour				
	albumID:trackID	On the Run	3:30	Gilmour, Waters	Instrumental				
		Idiocracy	Scifi Comedy	Mike Judge	20th Century Fox				
	movieID:actorID	Luke Wilson	Joe Bowers	image					
		Maya Rudolph	Rita	img3.jpg					
	movieID:actorID	Dax Shepard	Frito Pendejo	img1.jpg					

... or as Documents (JSON)

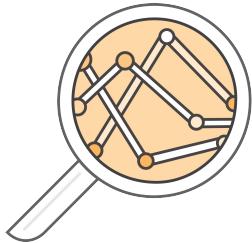
JSON data types (M, L, BOOL, NULL)

Indexing via Streams/Lambda

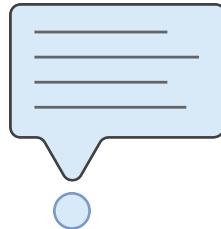
Fully Atomic Updates

Items	Primary Key	Attributes						
	ProductID	id	title	author	genre	publisher	datePublished	ISBN
1	bookID	1	Ringworld	Larry Niven	Science Fiction	Ballantine	Oct-70	0-345-02046-4
		id	title	artist	genre	Attributes		
	2	albumID	Dark Side of the Moon	Pink Floyd	Progressive Rock	{ label: "Harvest", studio: "Abbey Road", published: "3/1/73", producer: "Pink Floyd", tracks: [{title: "Speak to Me", length: "1:30", music: "Mason", vocals: "Instrumental"}, {title: "Breathe", length: "2:43", music: "Waters, Gilmour, Wright", vocals: "Gilmour"}, {title: "On the Run", length: "3:30", music: "Gilmour, Waters", vocals: "Instrumental"}]}		
2	movieID	id	title	genre	writer	Attributes		
		3	Idiocracy	Scifi Comedy	Mike Judge	{ producer: "20th Century Fox", actors: [{name: "Luke Wilson", dob: "9/21/71", character: "Joe Bowers", image: "img2.jpg"}, {name: "Maya Rudolph", dob: "7/27/72", character: "Rita", image: "img1.jpg"}, {name: "Dax Shepard", dob: "1/2/75", character: "Frito Pendejo", image: "img3.jpg"}]}		

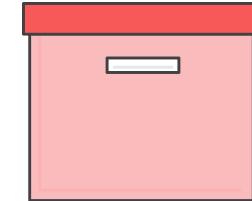
Scenarios and Best Practices



Time Series Data



Messaging



Online Voting



Product Catalog



Gaming

Event Logging

Storing time series data

Time Series Tables

Current
table

Events_table_2015_April					
Event_id (Partition)	Timestamp (Sort)	Attribute1	Attribute N	

RCUs = 10000
WCUs = 10000

Hot data

Older
tables

Events_table_2015_March					
Event_id (Partition)	Timestamp (Sort)	Attribute1	Attribute N	

RCUs = 1000
WCUs = 1

Events_table_2015_February					
Event_id (Partition)	Timestamp (Sort)	Attribute1	Attribute N	

RCUs = 100
WCUs = 1

Older
tables

Events_table_2015_January					
Event_id (Partition)	Timestamp (Sort)	Attribute1	Attribute N	

RCUs = 10
WCUs = 1

Don't mix hot and cold data; archive cold data to Amazon S3

Use a table per time period

Pre-create daily, weekly, monthly tables

Provision required throughput for current table

Writes go to the current table

Turn off (or reduce) throughput for older tables

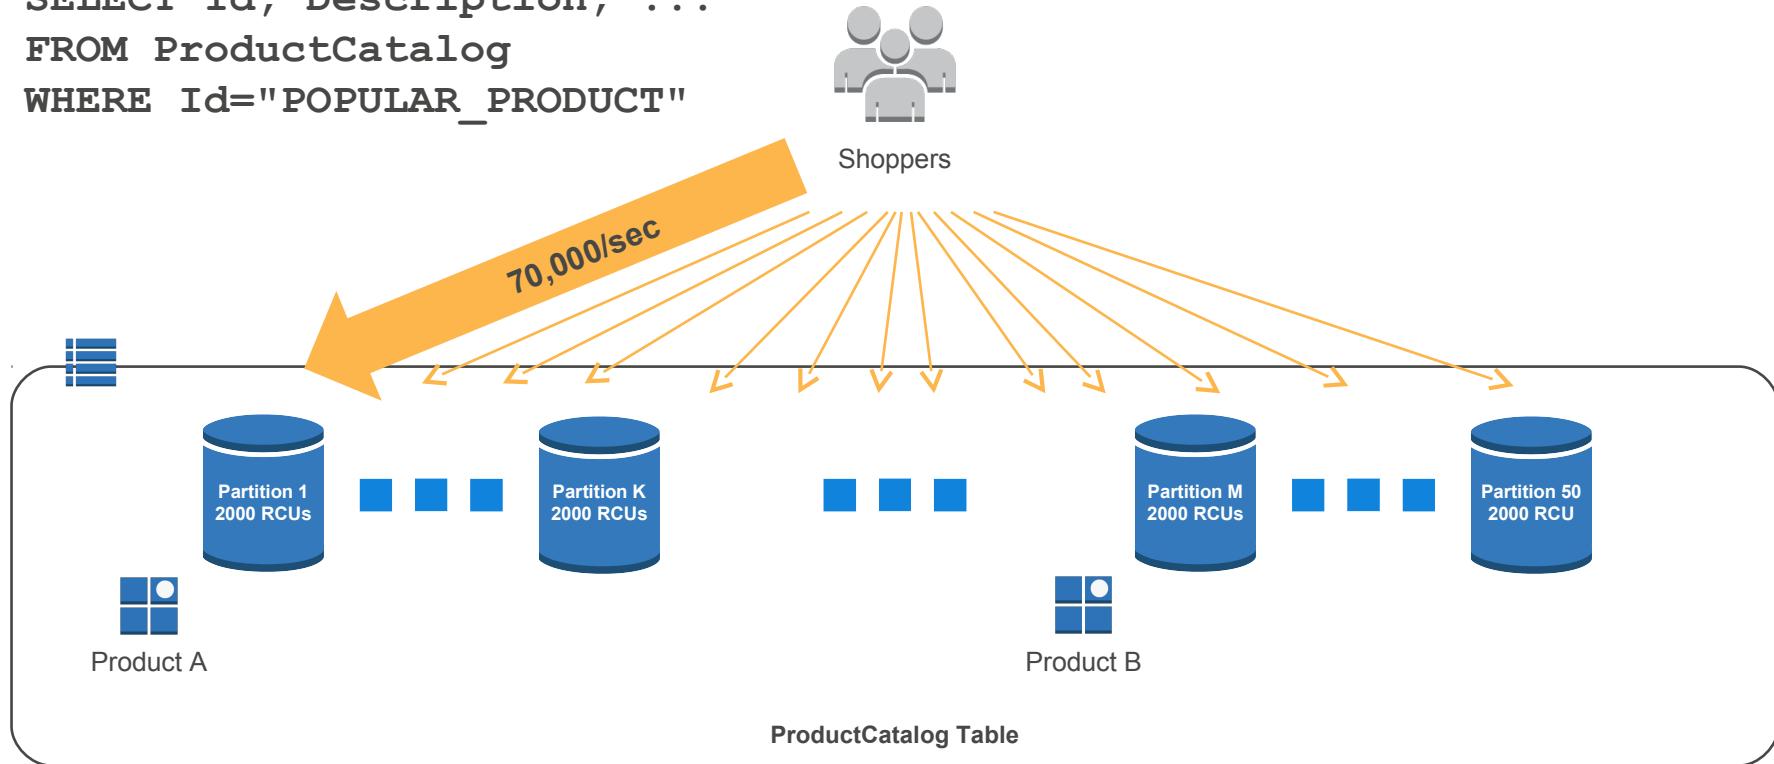
Dealing with time series data

Product Catalog

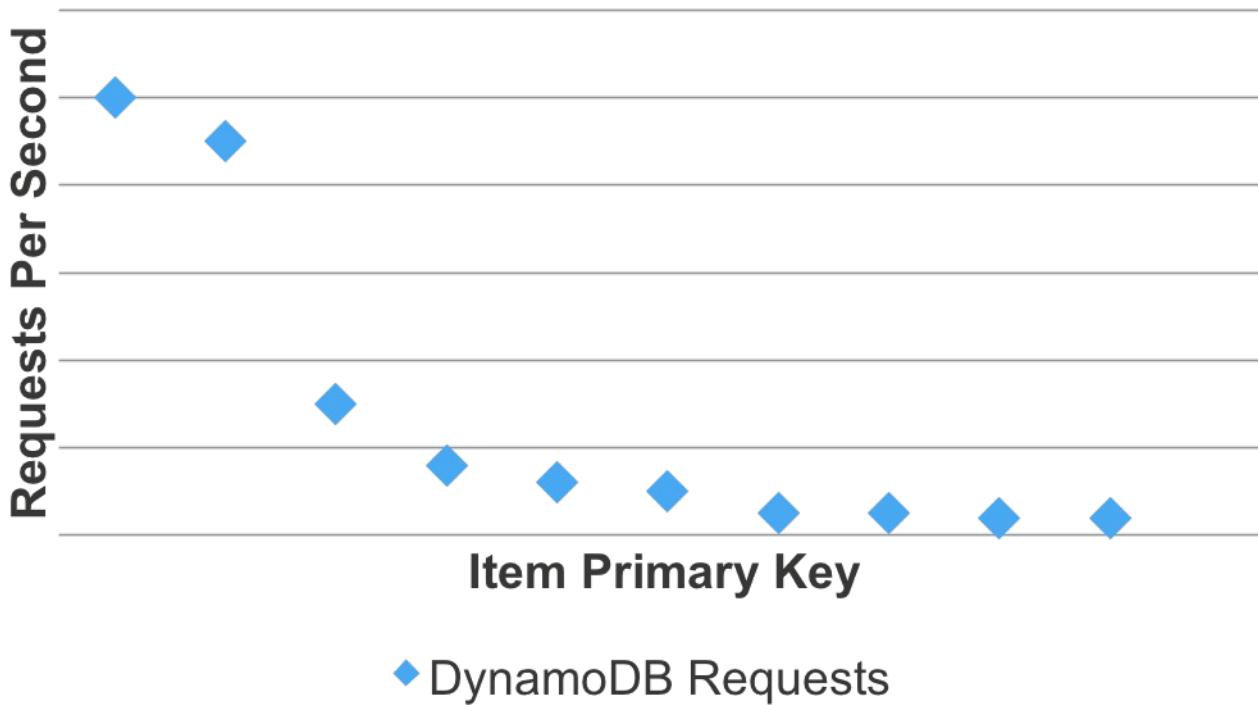
Popular items (read)

Scaling Bottlenecks

```
SELECT Id, Description, ...
FROM ProductCatalog
WHERE Id="POPULAR_PRODUCT"
```

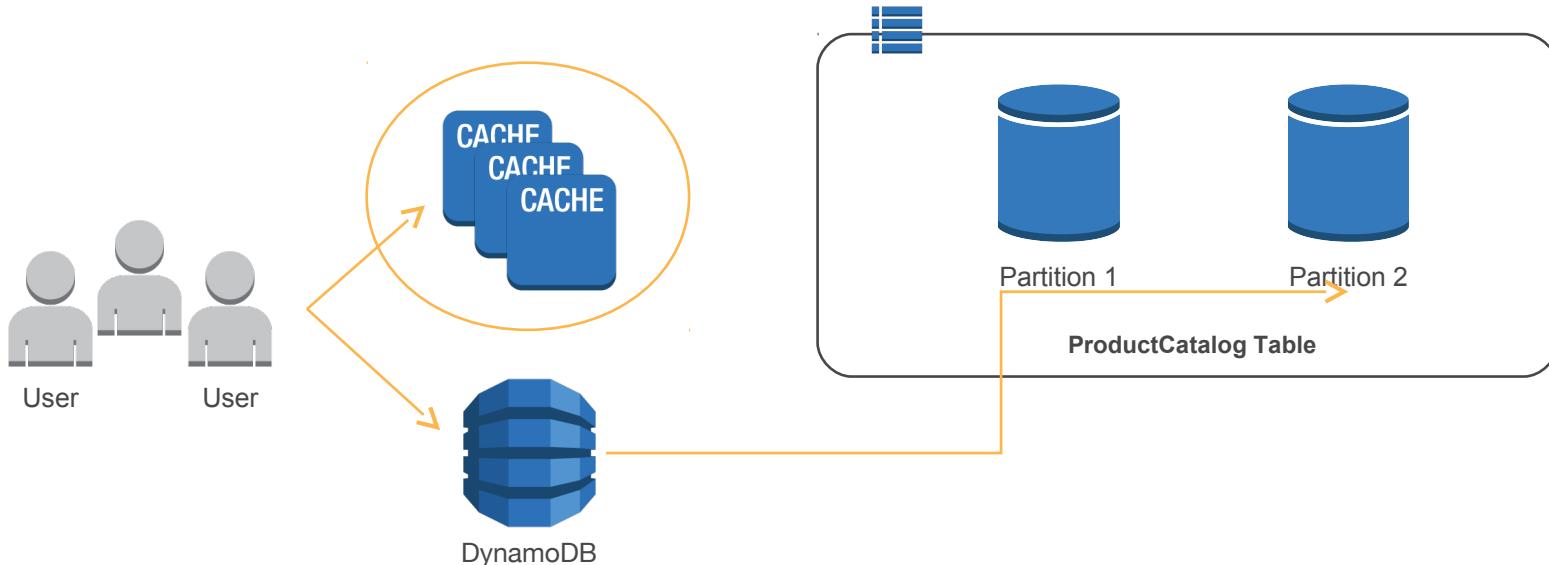


Request Distribution Per Partition Key

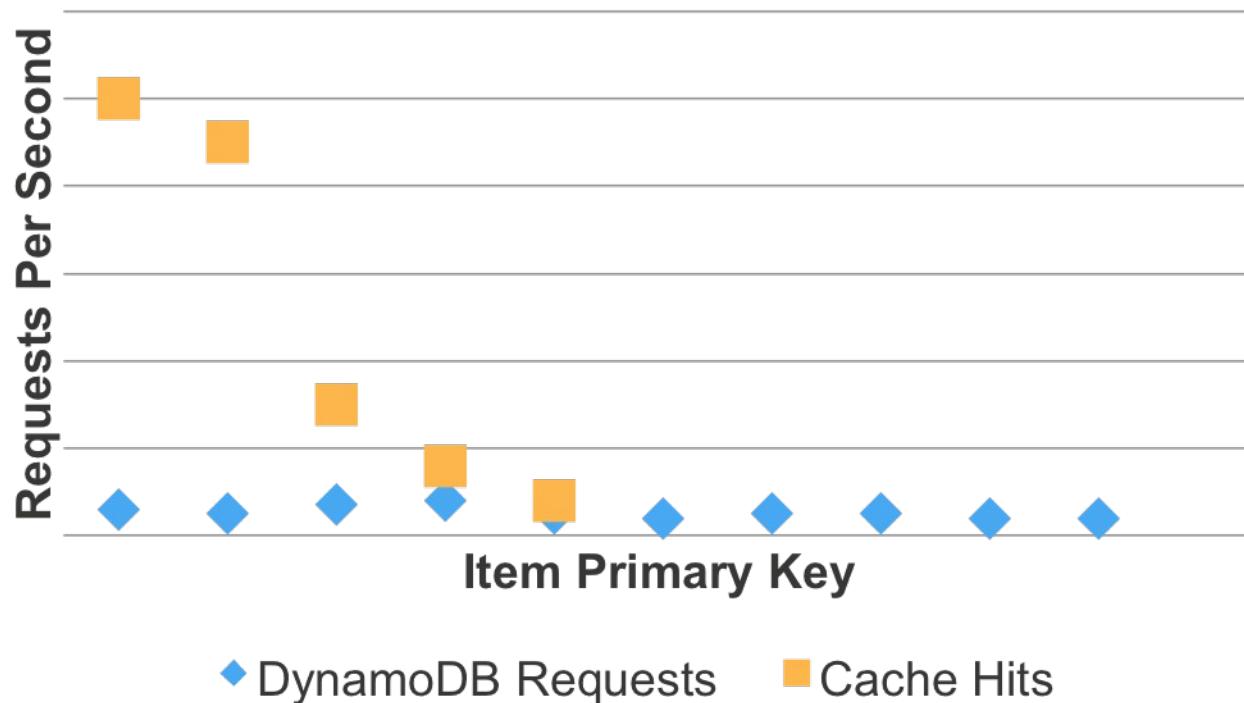


Cache Popular Items

```
SELECT Id, Description, ...  
FROM ProductCatalog  
WHERE Id="POPULAR_PRODUCT"
```



Request Distribution Per Partition Key



Messaging App

Large items

Filters vs. indexes

M:N Modeling—inbox and outbox



David



Messages App



Messages
Table

Inbox

```
SELECT *  
FROM Messages  
WHERE Recipient='David'  
LIMIT 50  
ORDER BY Date DESC
```

Outbox

```
SELECT *  
FROM Messages  
WHERE Sender = 'David'  
LIMIT 50  
ORDER BY Date DESC
```

Large and small attributes mixed

Messages Table

Recipient	Date	Sender	Message
David	2014-10-02	Bob	...
... 48 more messages for David ...			
David	2014-10-03	Alice	...
Alice	2014-09-28	Bob	...
Alice	2014-10-01	Carol	...

(Many more messages)

Inbox



```
SELECT *
FROM Messages
WHERE Recipient='David'
LIMIT 50
ORDER BY Date DESC
```

50 items × 256 KB each

Large message bodies
Attachments

Computing inbox query cost

$$50 * 256\text{KB} * (1 \text{ RCU} / 4\text{KB}) * (1 / 2) = 1600 \text{ RCU}$$

Items evaluated by query

Average item size

Conversion ratio

Eventually consistent reads

Separate the bulk data

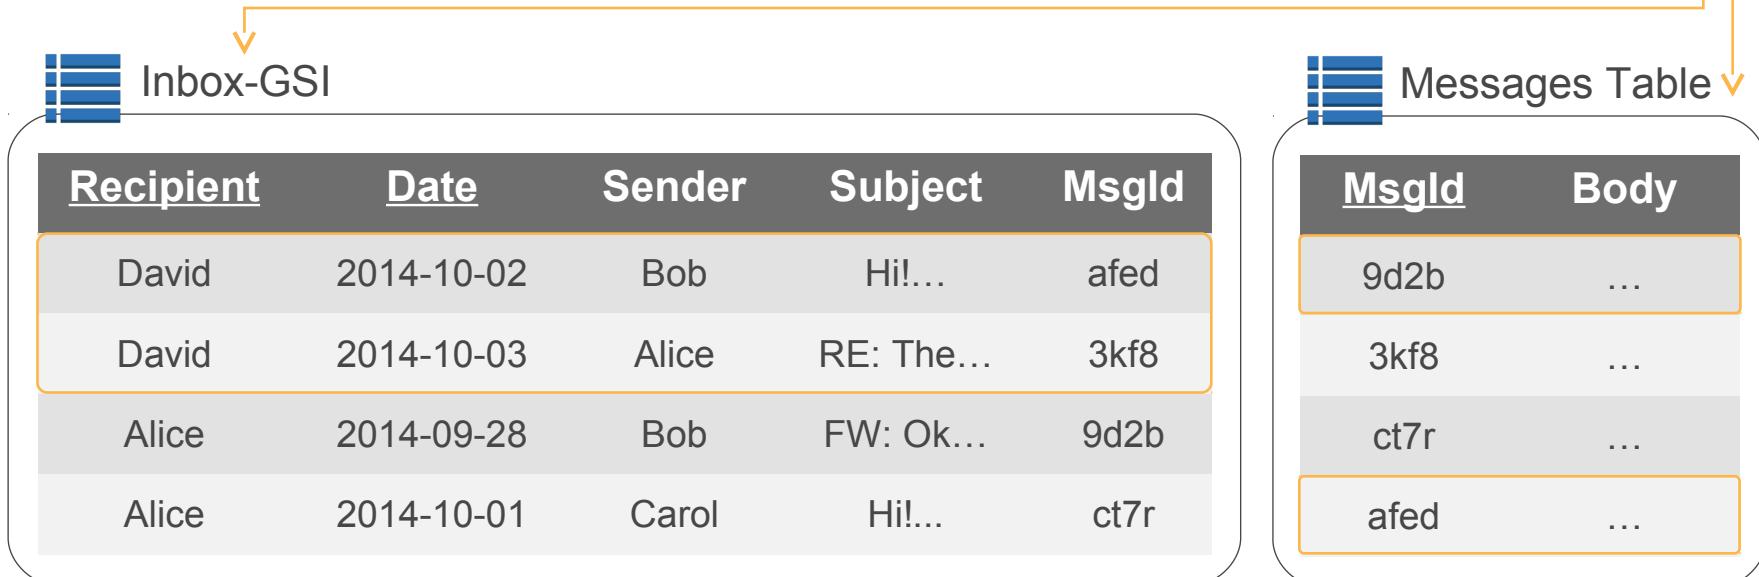


David

Uniformly distributes large item reads

Query Inbox-GSI: 1 RCU (50 sequential items at 128 bytes)

BatchGetItem Messages: 1600 RCU (50 separate items at 256 KB)



Inbox GSI

Define which attributes to copy into the index

The screenshot shows the AWS Lambda console interface. At the top, there are four tabs: 'Details', 'Indexes' (which is highlighted with an orange border), 'Monitoring', and 'Alarm Setup'. Below the tabs, the heading 'Global Secondary Indexes' is displayed. A table lists the configuration for a single index named 'Inbox'. The columns are 'Index Name', 'Hash Key', and 'Range Key'. The 'Index Name' row contains 'Inbox'. The 'Hash Key' row contains 'Recipient (String)'. The 'Range Key' row contains 'Date (String)'. To the right of the table, a callout box with an orange border and arrow points to the 'Projected Attributes' section. This section lists the attributes to be copied: 'MsgId, Recipient, Date, Subject, Sender'.

Index Name	Hash Key	Range Key	Projected Attributes
Inbox	Recipient (String)	Date (String)	MsgId, Recipient, Date, Subject, Sender

Outbox GSI

Details

Indexes

Monitoring

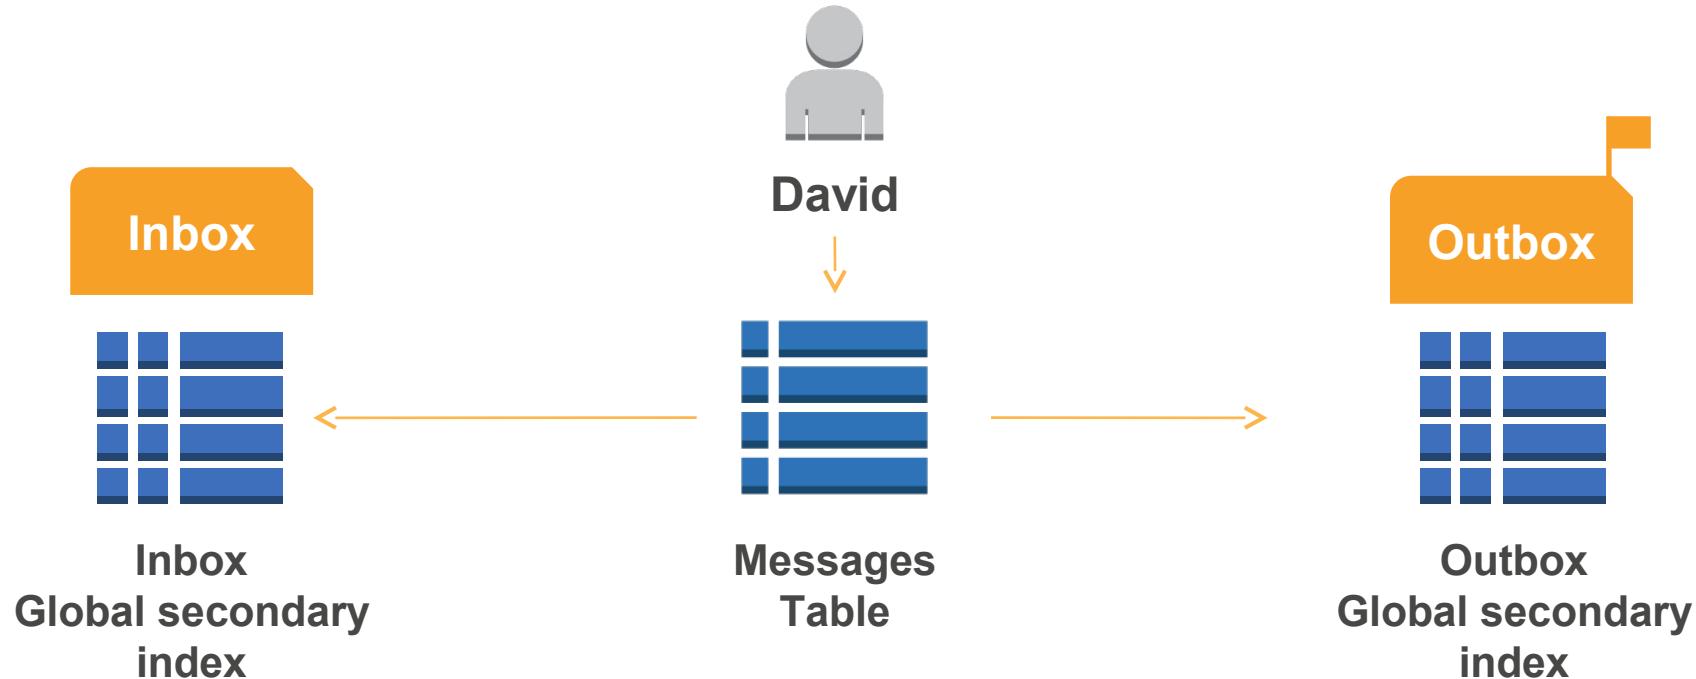
Alarm Setup

Global Secondary Indexes

Index Name	Hash Key	Range Key	Projected Attributes
Outbox	Sender (String)	Date (String)	MsgId, Recipient, Date, Subject, Sender

```
SELECT *
FROM Messages
WHERE Sender = 'David'
LIMIT 50
ORDER BY Date DESC
```

Messaging app



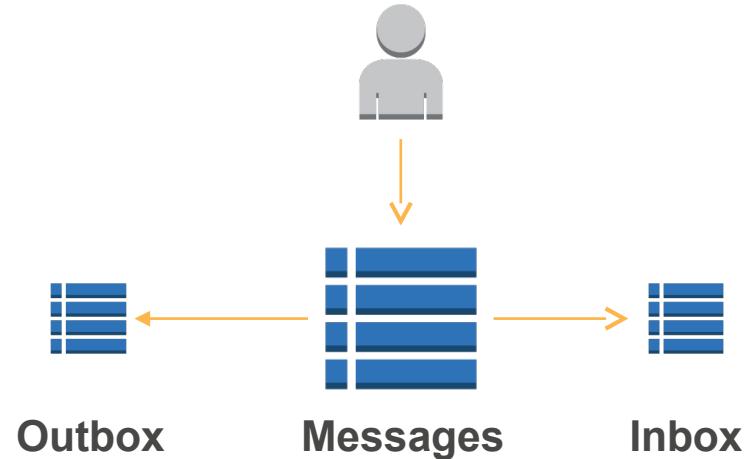
Distribute large items

Reduce one-to-many item sizes

Configure secondary index projections

Use GSIs to model M:N relationship
between sender and recipient

Important when: Querying many
large items at once



Multiplayer Online Gaming

Query filters vs.
composite key indexes

Hierarchical Data Structures

Partition Key



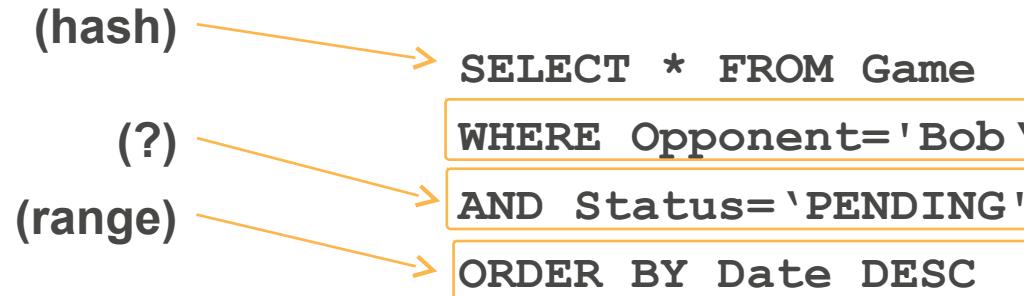
Games Table

Gameld	Date	Host	Opponent	Status
d9bl3	2014-10-02	David	Alice	DONE
72f49	2014-09-30	Alice	Bob	PENDING
o2pnb	2014-10-08	Bob	Carol	IN_PROGRESS
b932s	2014-10-03	Carol	Bob	PENDING
ef9ca	2014-10-03	David	Bob	IN_PROGRESS

Query for incoming game requests

DynamoDB indexes provide partition and sort

What about queries for two equalities and a sort?



Approach 1: Query filter



Bob

Partition Key Sort Key

Secondary Index

Opponent	Date	Gameld	Status	Host
Alice	2014-10-02	d9bl3	DONE	David
Carol	2014-10-08	o2pnb	IN_PROGRESS	Bob
Bob	2014-09-30	72f49	PENDING	Alice
Bob	2014-10-03	b932s	PENDING	Carol
Bob	2014-10-03	ef9ca	IN_PROGRESS	David

Approach 1: Query filter



Bob

```
SELECT * FROM Game  
WHERE Opponent='Bob'  
ORDER BY Date DESC  
FILTER ON Status='PENDING'
```



Secondary Index

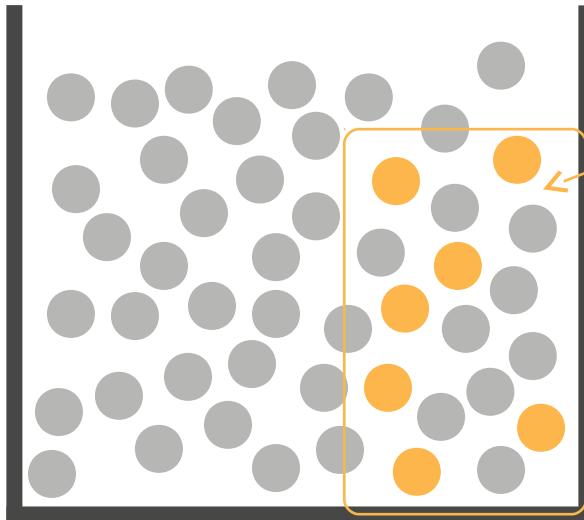
Opponent	Date	GameId	Status	Host
Alice	2014-10-02	d9bl3	DONE	David
Carol	2014-10-08	o2pnb	IN_PROGRESS	Bob
Bob	2014-09-30	72f49	PENDING	Alice
Bob	2014-10-03	b932s	PENDING	Carol
Bob	2014-10-03	ef9ca	IN_PROGRESS	David

(filtered out)

Needle in a haystack



Bob



Use query filter

Send back less data “on the wire”

Simplify application code

Simple SQL-like expressions

- AND, OR, NOT, ()

Important when: Your index isn’t entirely selective

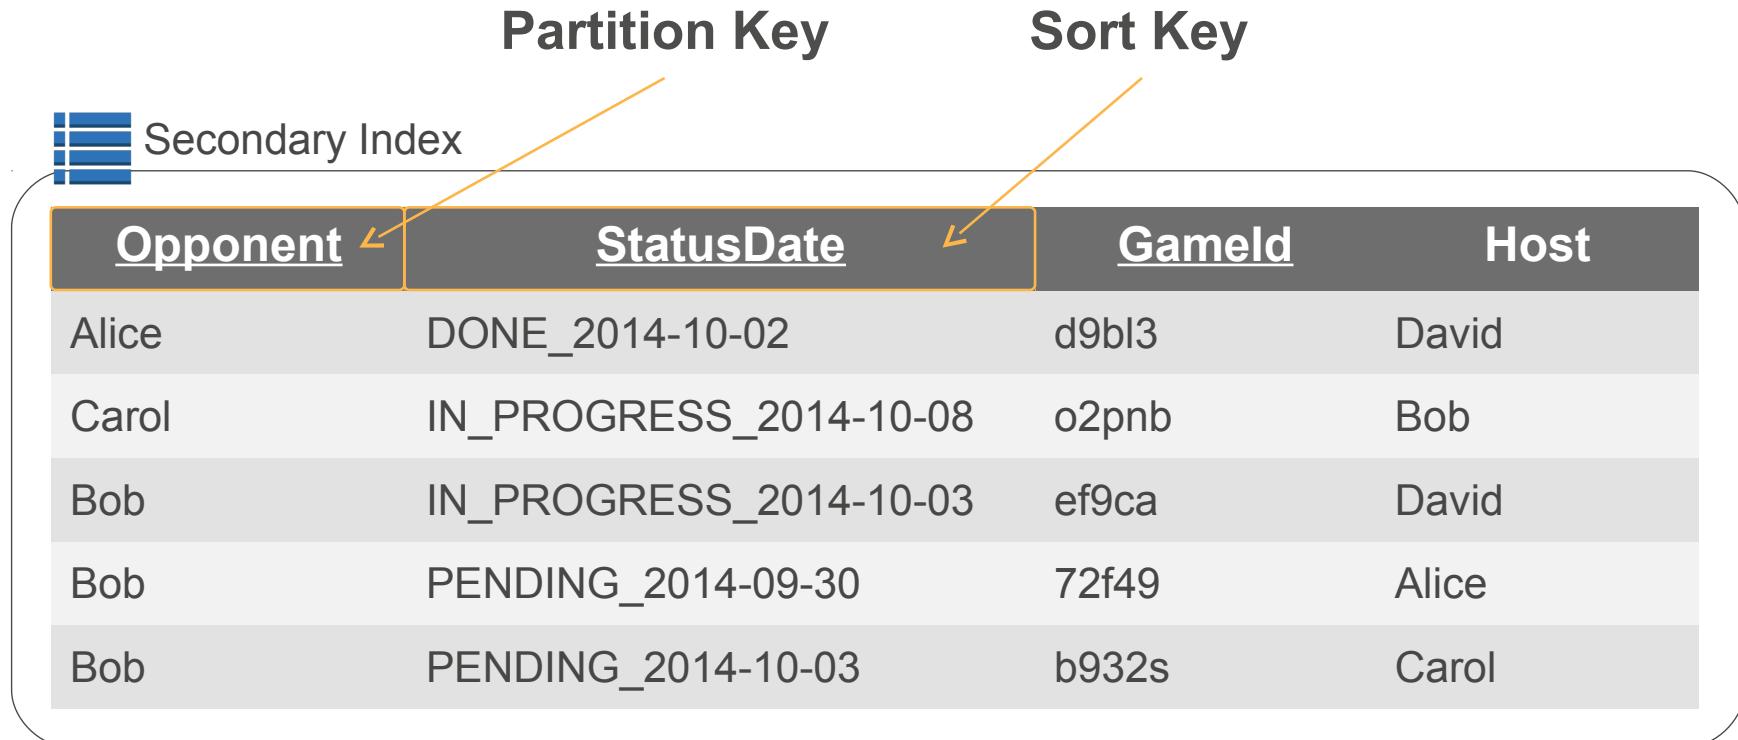
Approach 2: Composite key

Status	Date	StatusDate
DONE	2014-10-02	DONE_2014-10-02
IN_PROGRESS	2014-10-08	IN_PROGRESS_2014-10-08
IN_PROGRESS	2014-10-03	IN_PROGRESS_2014-10-03
PENDING	2014-10-03	PENDING_2014-09-30
PENDING	2014-09-30	PENDING_2014-10-03

+

=

Approach 2: Composite key



The diagram illustrates a composite key structure. A blue icon representing a secondary index is positioned above the table. Two orange arrows point from the text labels "Partition Key" and "Sort Key" to the "Opponent" column header and the "StatusDate" column header respectively.

Partition Key				Sort Key			
Opponent		StatusDate		GameId		Host	
Alice		DONE_2014-10-02		d9bl3		David	
Carol		IN_PROGRESS_2014-10-08		o2pnb		Bob	
Bob		IN_PROGRESS_2014-10-03		ef9ca		David	
Bob		PENDING_2014-09-30		72f49		Alice	
Bob		PENDING_2014-10-03		b932s		Carol	

Approach 2: Composite key



Bob

```
SELECT * FROM Game  
WHERE Opponent='Bob'  
AND StatusDate BEGINS_WITH 'PENDING'
```



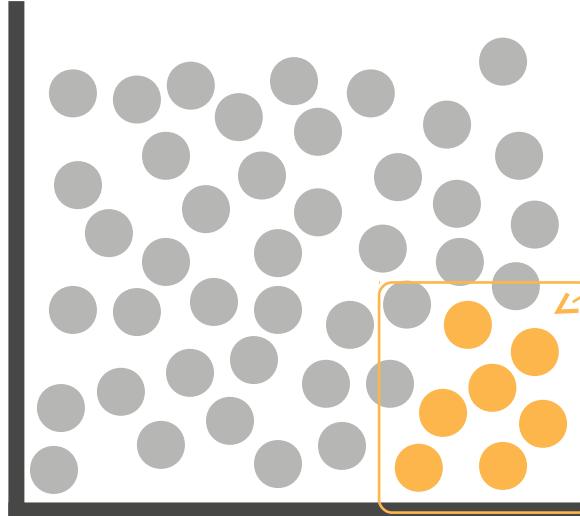
Secondary Index

Opponent	StatusDate	GameId	Host
Alice	DONE_2014-10-02	d9bl3	David
Carol	IN_PROGRESS_2014-10-08	o2pnb	Bob
Bob	IN_PROGRESS_2014-10-03	ef9ca	David
Bob	PENDING_2014-09-30	72f49	Alice
Bob	PENDING_2014-10-03	b932s	Carol

Needle in a sorted haystack



Bob



Sparse indexes

Game-scores-table

Id (Partition)	User	Game	Score	Date	Award
1	Bob	G1	1300	2012-12-23	
2	Bob	G1	1450	2012-12-23	
3	Jay	G1	1600	2012-12-24	
4	Mary	G1	2000	2012-10-24	Champ
5	Ryan	G2	123	2012-03-10	
6	Jones	G2	345	2012-03-20	

Award-GSI

Award (Partition)	ID	User	Score
Champ	4	Mary	2000

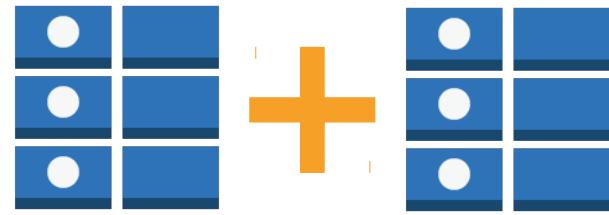
Scan sparse GSIs

Replace filter with indexes

Concatenate attributes to form useful secondary index keys

Take advantage of sparse indexes

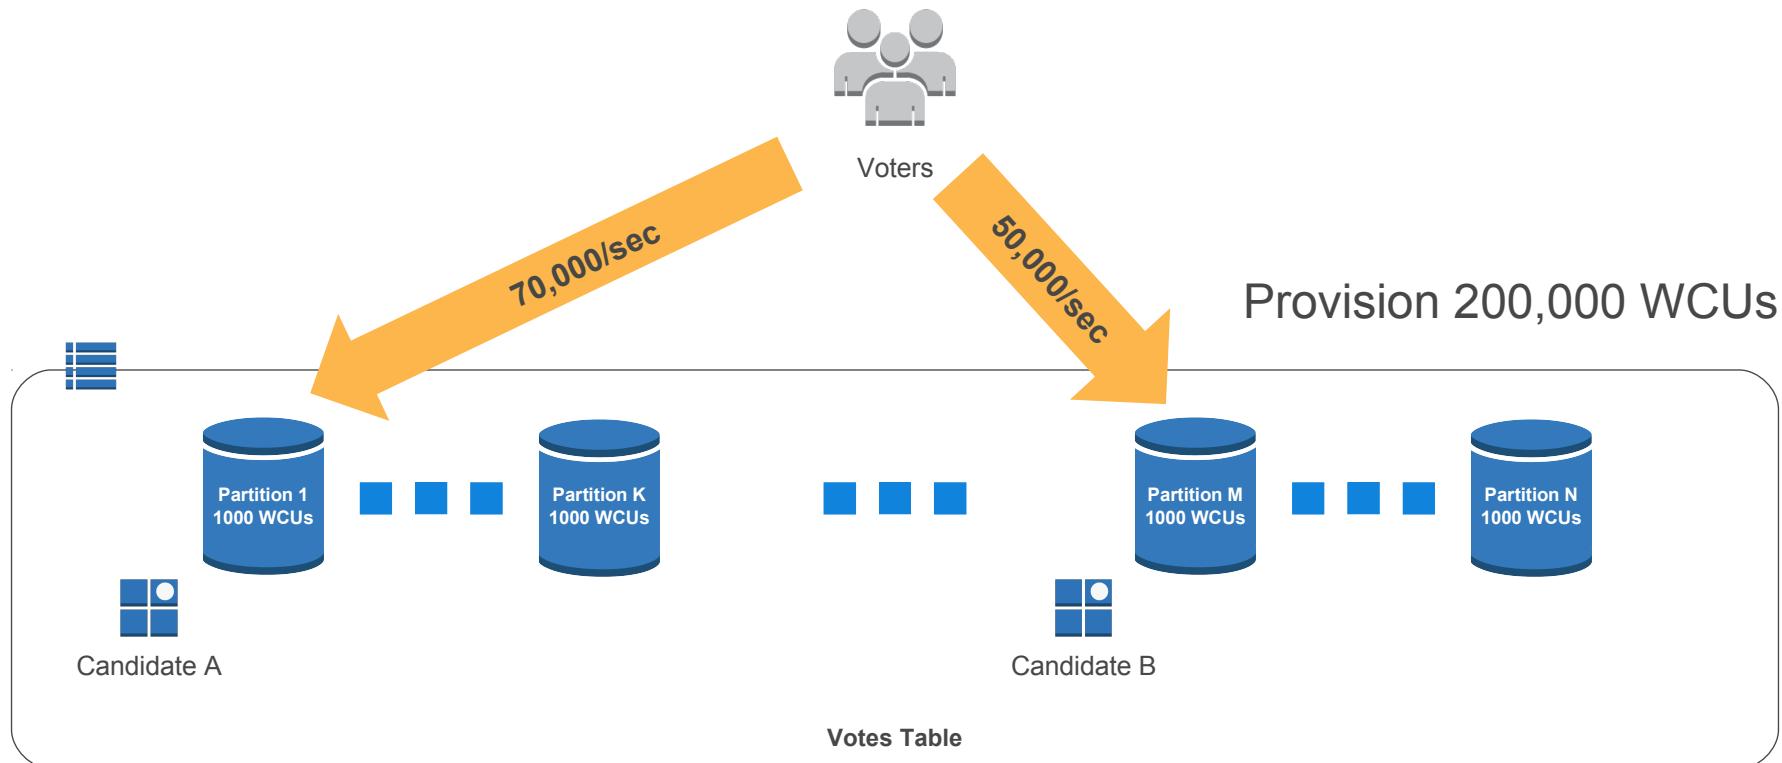
Important when: You want to optimize a query as much as possible



Real-Time Voting

Write-heavy items

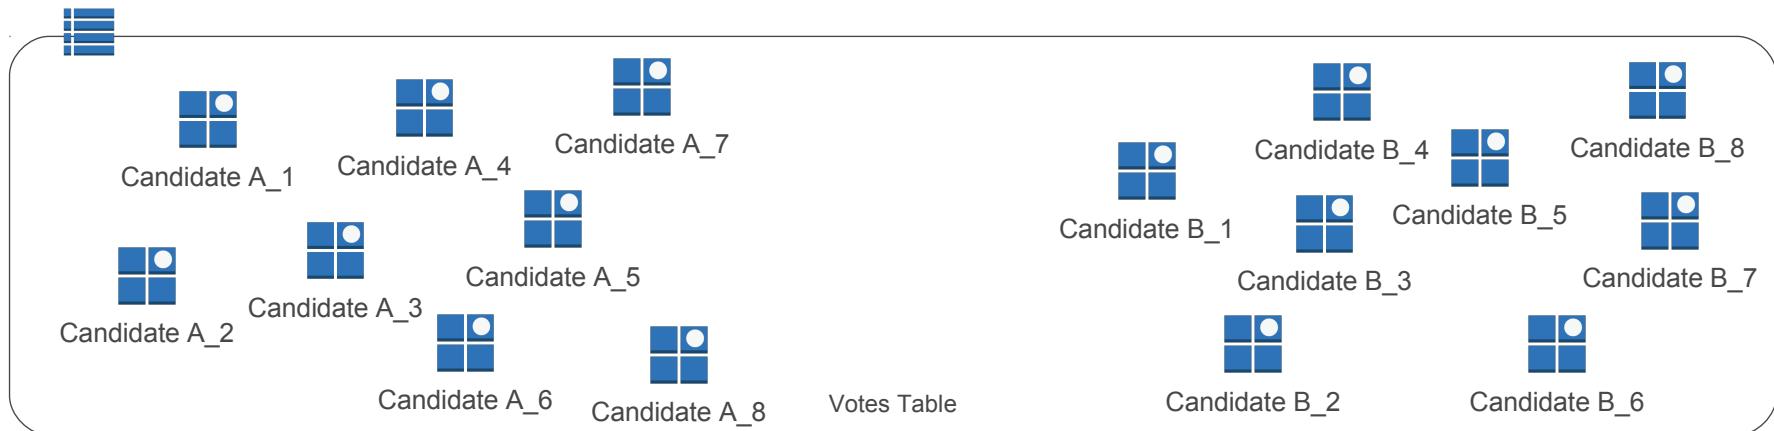
Scaling bottlenecks



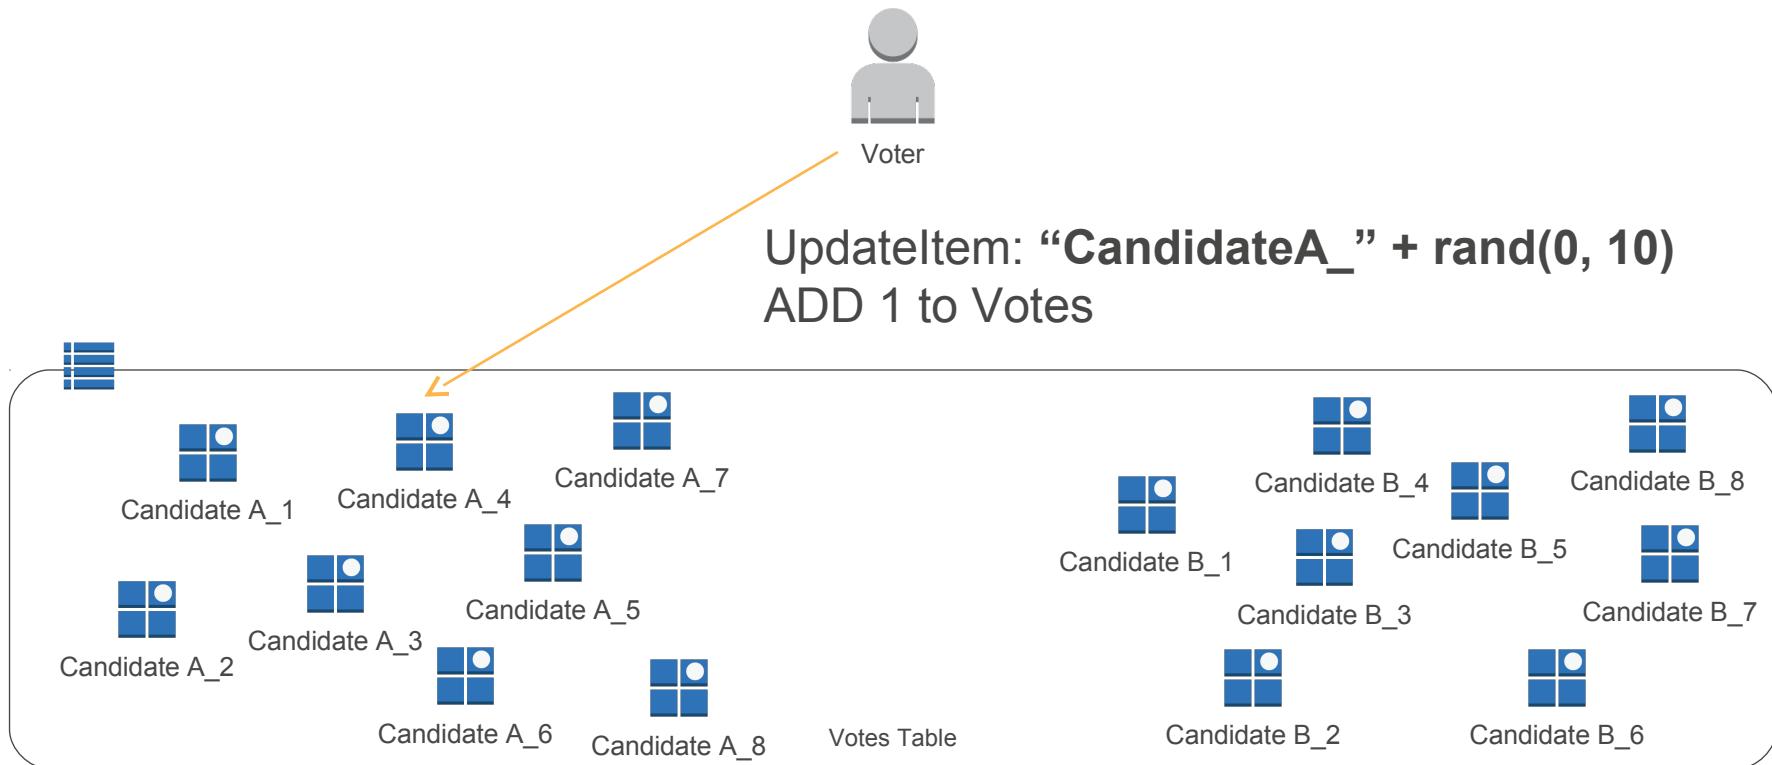
Write sharding



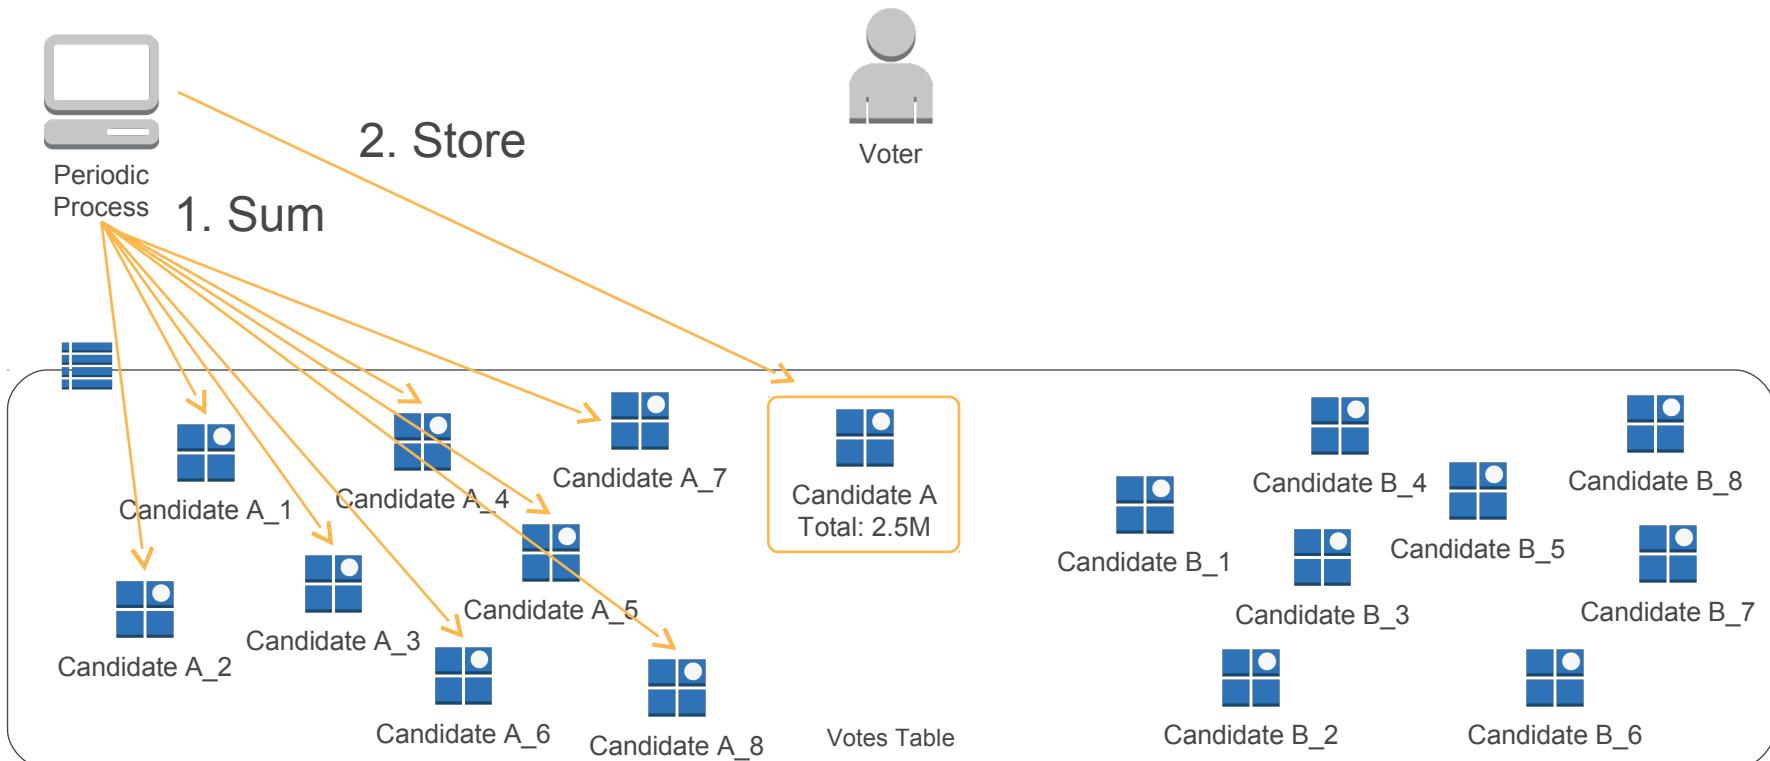
Voter



Write sharding



Shard aggregation

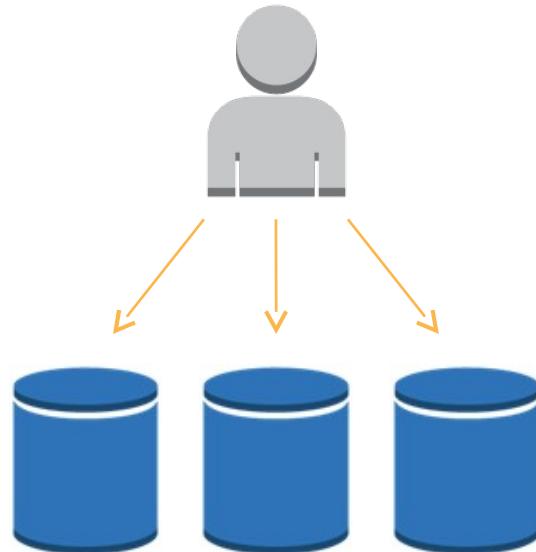


Shard write-heavy partition keys

Trade off read cost for write scalability

Consider throughput per partition key

Important when: Your write workload is not horizontally scalable



DynamoDB Streams

DynamoDB Streams

Stream of updates to a table

Asynchronous

Exactly once

Strictly ordered

- Per item

Highly durable

- Scale with table

24-hour lifetime

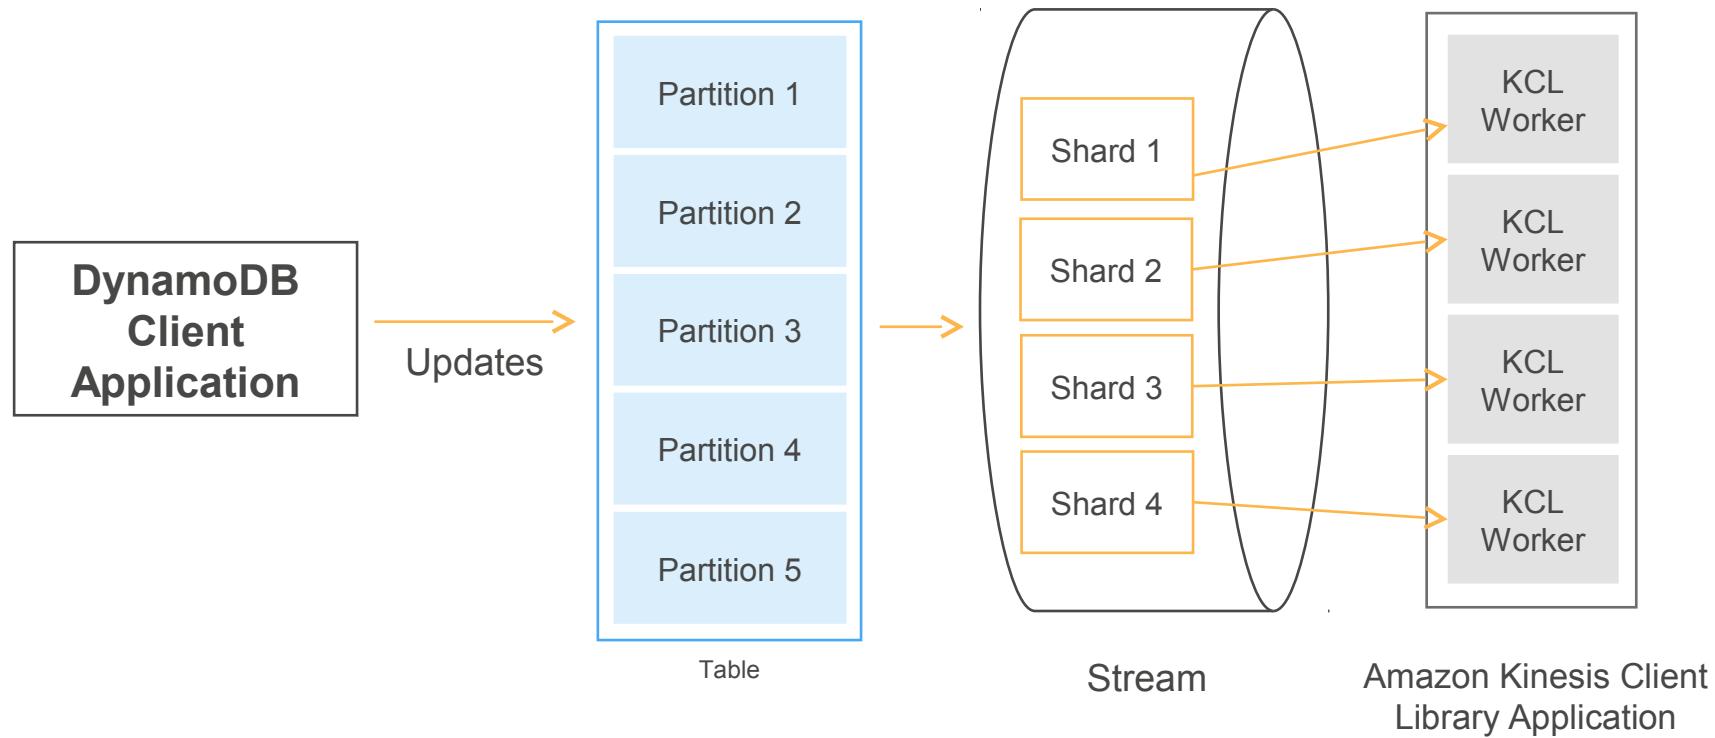
Sub-second latency

View types

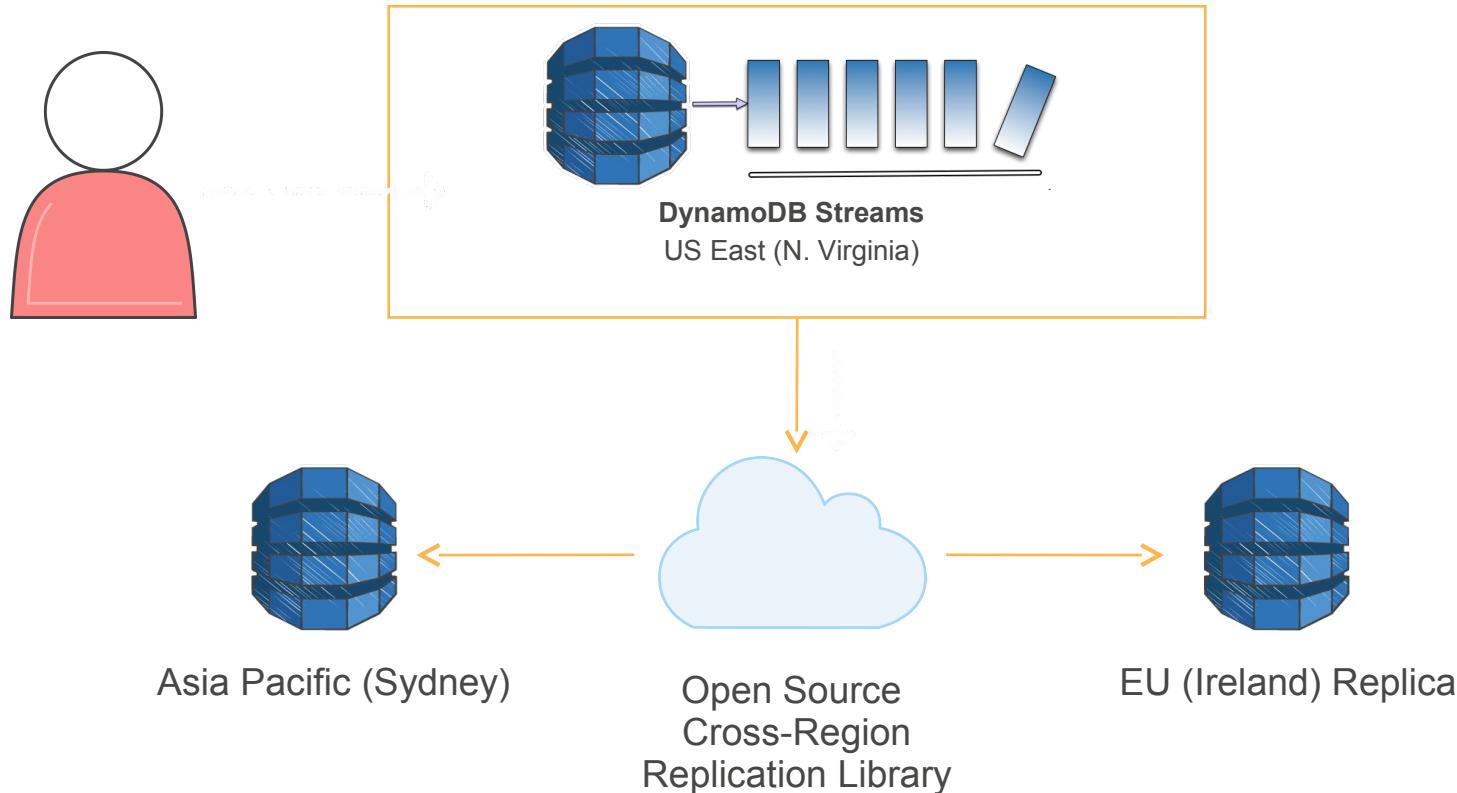
UpdateItem (Name = John, Destination = Pluto)

View Type	Destination
Old image—before update	Name = John, Destination = Mars
New image—after update	Name = John, Destination = Pluto
Old and new images	Name = John, Destination = Mars Name = John, Destination = Pluto
Keys only	Name = John

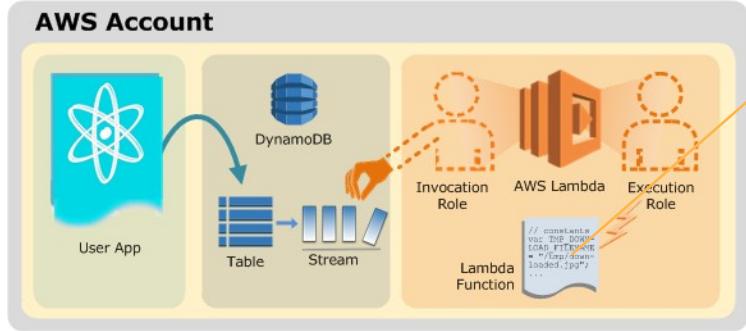
DynamoDB Streams and Amazon Kinesis Client Library



Cross-region replication



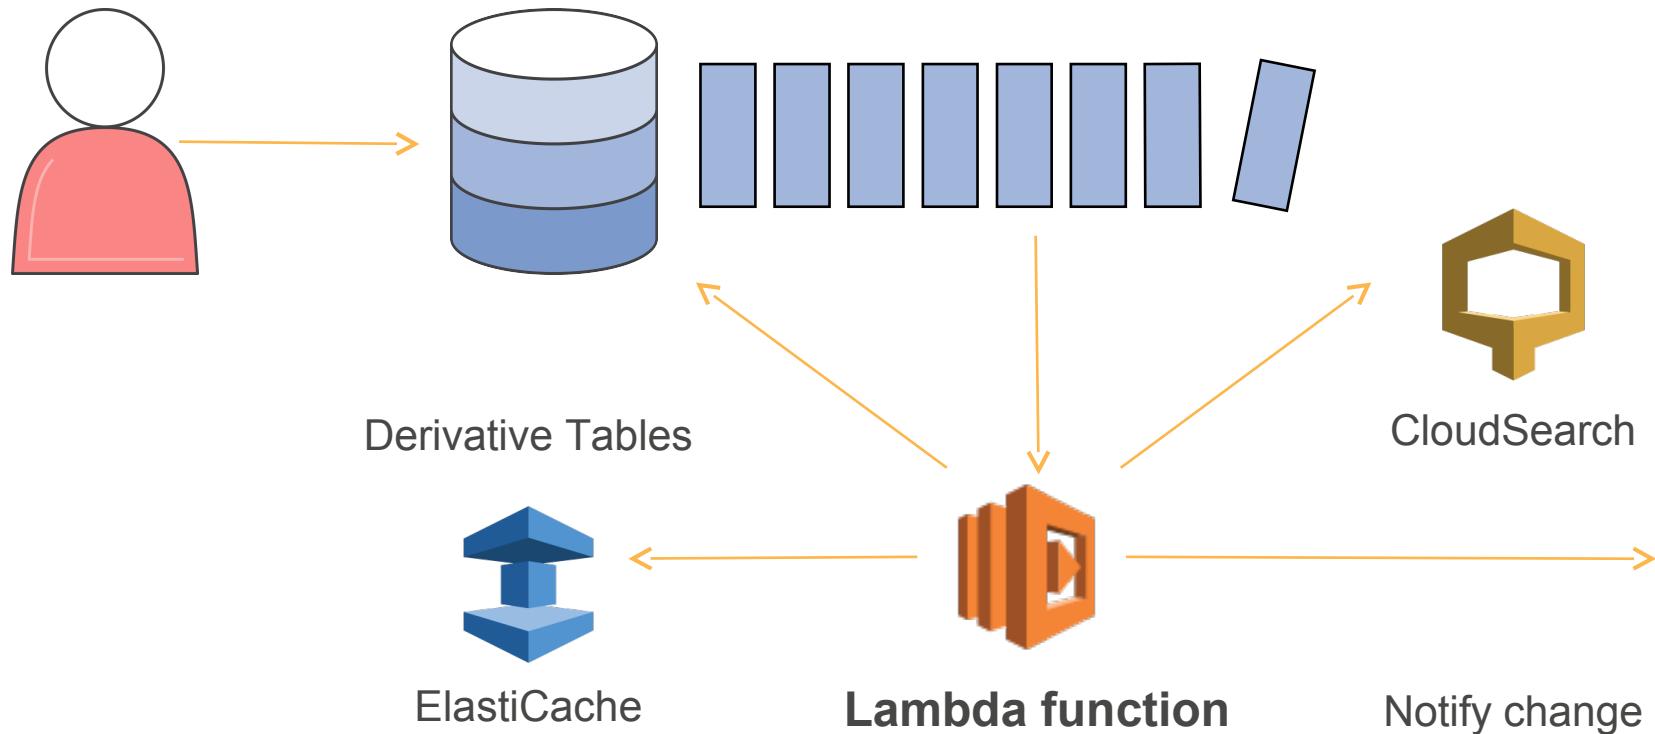
DynamoDB Streams and AWS Lambda



```
1  console.log('Loading event');
2  exports.handler = function(event, context) {
3      console.log("Event: %j", event);
4      for(i = 0; i < event.Records.length; ++i) {
5          record = event.Records[i];
6          console.log(record.EventID);
7          console.log(record.EventName);
8          console.log("DynamoDB Record: %j", record.Dynamodb);
9      }
10     context.done(null, "Hello World"); // SUCCESS with message
11 }
```

```
▶ 2015-03-21T07:44:58.883Z 2ca3769a-cf9e-11e4-b270-ad4d24b312ff INSERT
▶ 2015-03-21T07:44:58.883Z 2ca3769a-cf9e-11e4-b270-ad4d24b312ff DynamoDB Record:{ "NewImage": { "name": { "S": "sivar" }, "hk": { "S": "3" } }, "SizeBytes": 15, "StreamViewType": "NEW_AND_OLD_IMAGES" }
▶ 2015-03-21T07:44:58.883Z 2ca3769a-cf9e-11e4-b270-ad4d24b312ff Message: "Hello World"
```

Triggers

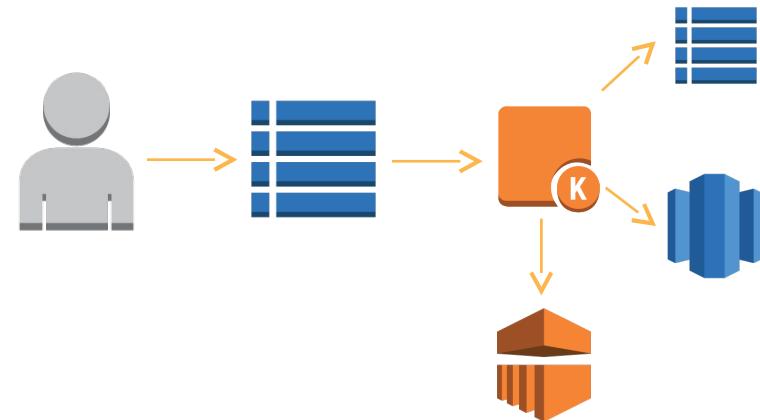


Analytics with DynamoDB Streams

Collect and de-dupe data in
DynamoDB

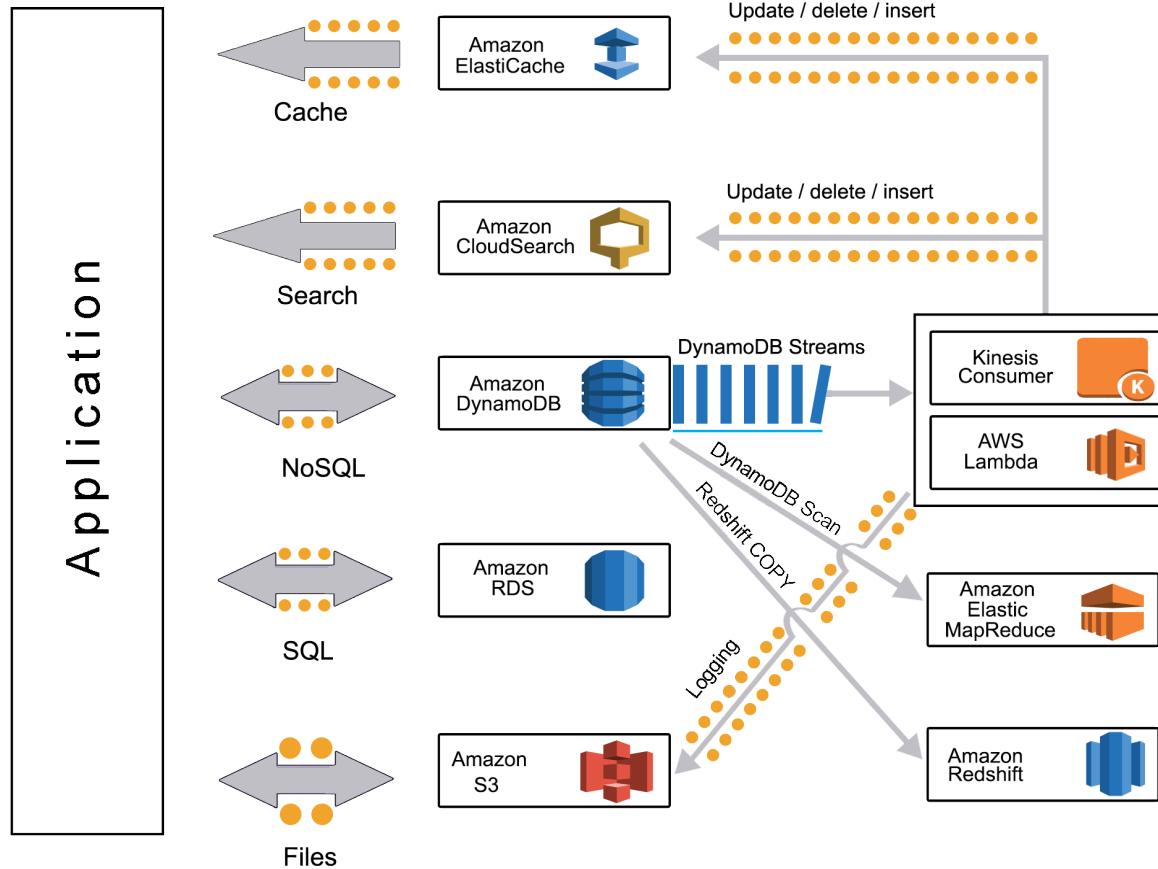
Aggregate data in-memory
and flush periodically

Important when: Performing
real-time aggregation
and analytics

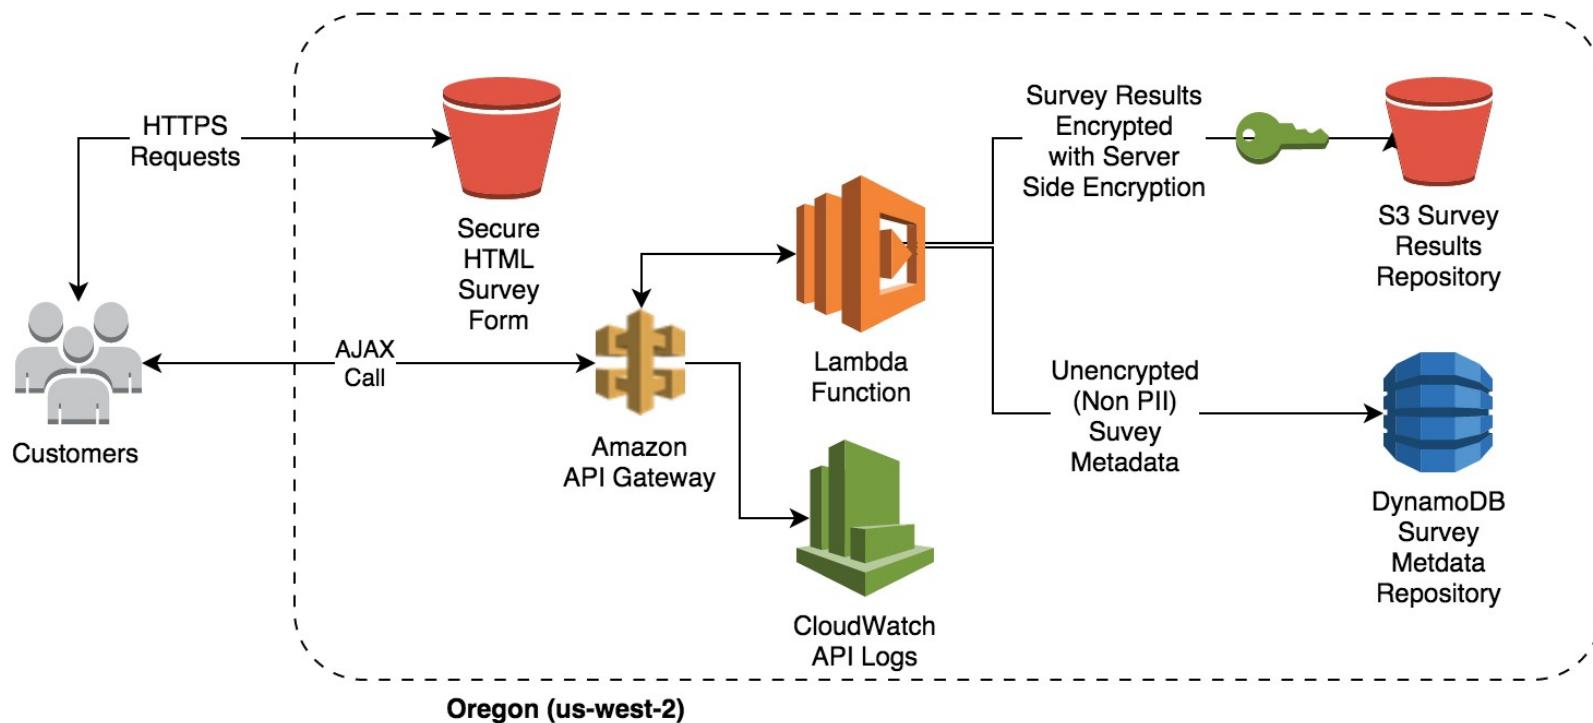


Architecture

Reference Architecture



Elastic Event Driven Applications



Thank you!

bit.ly/NoSQLDesignPatterns

Free Access to All DynamoDB Self-paced Labs: bit.ly/DynamoDBLabs
Normally 10 USD per lab. Promotion expires March 31st, 2016.



Chicago – April 18-19

AWS Summit – Chicago: An exciting, free cloud conference designed to educate and inform new customers about the AWS platform, best practices and new cloud services.

Details

- April 18-19, 2016
- Chicago, Illinois
- @ McCormick Place

Featuring

- New product launches
- 50+ sessions, labs, and bootcamps
- Executive and partner networking



Register Now

- Go to aws.amazon.com/summits
- Click on The AWS Summit - Chicago ... then register.
- Come and see what AWS and the cloud can do for you.