


CARLSON SCHOOL
OF MANAGEMENT

UNIVERSITY OF MINNESOTA



Introduction to NoSQL Data Stores

MSBA 6330 Prof Liu

1

Carlson School of Management

Learning Objectives

- Understand strengths and limitations of traditional SQL databases
- Understand common characteristics of NoSQL databases
- Understand the trade-offs facing distributed databases in light of the CAP theorem
- Distinguish between ACID and BASE consistency models
- Be familiar with different types of NoSQL / NewSQL databases and their use cases
- Understand the complexity of choosing different database technologies for big data applications

2

Carlson School of Management

Topics

- Limitations of RDBMS for the Web
- What is NoSQL?
- NoSQL and the CAP theorem
- Different kinds of NoSQL data stores
- NewSQL
- How to choose?

3

Cartoon School of Management

Introduction to NoSQL

LIMITATIONS OF RDBMS FOR THE WEB

4

Cartoon School of Management

RDBMS has been a dominant data storage model for a few decades

- Relational databases have been a mainstay of how data is stored since 1970s (and still are)
 - Shines with complicated queries
 - Declarative queries (SQL) rather than programing
 - Fast query & analysis, but not necessarily on large datasets
- Strong support for transactions – ACID Properties
 - **Atomic** – All of the work in a transaction completes (commit) or none of it
 - **Consistent** – A transaction transforms the database from one consistent state to another consistent state. Consistency is defined in terms of constraints.
 - **Isolated** – The results of any changes made during a transaction are not visible until the transaction has committed.
 - **Durable** – The results of a committed transaction survive failures

5

Cartoon School of Management

But there are limits in scaling RDBMSs

- Best way to provide ACID and a rich query model is to have the dataset on a single machine.
- However, there are limits to scaling up (vertical scaling) as the dataset gets big
- Past a certain point, an organization will find it is cheaper and more feasible to scale out (horizontal scaling) than scaling up
- DBAs began to look at ways of scaling out
 - *master-slave*: all writes go to the master. all reads can perform against a fleet of replicated slave databases (inconsistency; do not scale well)
 - *sharding*: partition database based on value range, hash-key, or dictionaries (unable to join across partitions; no referential integrity)

http://adam.herokuapp.com/past/2009/7/6/sql_databases_dont_scale/

6

Inadequacy of RDMS for the web

- Hooking your RDBMS to a web-based application was a recipe for headaches
 - "Slashdot effect": when a popular website links to a smaller website, causing a massive spike in traffic
- Sharding is like putting a bandage on a broken leg
 - It does not scale to handle hundreds of thousands of visitors in a short-time span
 - People with the largest datasets (terabyte/petabyte), e.g. Google, Facebook, Twitter, began to explore alternative ways to store data, especially around 2008/2009.

7

Introduction to NoSQL

WHAT IS NOSQL?

8

Drivers of the NoSQL

- Two papers were the seeds of the NoSQL movement
 - [BigTable](#) (Google, 2006):
 - Wide-column data model
 - Consistent, fault-tolerant, scale to large clusters (1000+ nodes)
 - [Dynamo](#) (Amazon, 2007):
 - Distributed key-value data store
 - Eventual consistency, always writable, fault-tolerant
- In addition,
 - The advent of Amazon S3 signals to others it was okay to look at alternative storage solutions
 - Many NoSQL solutions are open source software


9

Cartoon School of Management

What is NoSQL?

- Stands for **Not Only SQL** (coined by Eric Evans)
 - "to solve a problem that relational databases are a bad fit for."
- (No standard definition) nosql-database.org:
 - Next Generation Databases mostly addressing some of the points: being non-relational, distributed, open-source and horizontal scalable.

NoSQL-Databases.org:
Current list has over 150 NoSQL systems



10

Cartoon School of Management

Common NoSQL Characteristics

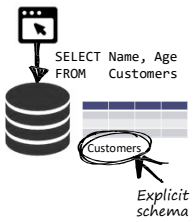
- don't use the relational data model
- optimized for running on a cluster
- open source
- schema-less (schema-free)
- tunable consistency

11

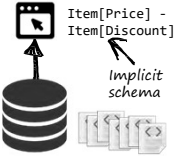
Cartoon School of Management

Schema-free Data Modeling

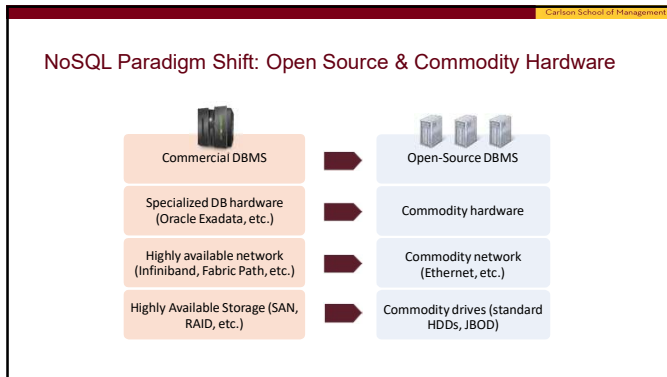
RDBMS:



NoSQL DB:



12



13

Carleton School of Management

NoSQL Pros and Cons

- Pros:
 - Massive scalability
 - Flexible schema
 - Quicker/cheaper to set up
 - Relaxed consistency (NO ACID) → higher performance & availability
- Cons:
 - No declarative query language (SQL) → more programming
 - May not integrate well with other applications that support SQL
 - May lose powerful SQL queries e.g., joins, group by, order by
 - Relaxed consistency → fewer guarantees (NO ACID)

14

Carleton School of Management

Challenge: Coordination

- The solution to availability and scalability is to decentralize and replicate functions and data...but how do we coordinate the nodes?
 - data consistency
 - update propagation
 - mutual exclusion
 - consistent global states
 - group membership
 - group communication
 - event ordering
 - distributed consensus
 - quorum consensus

15

Cartoon School of Management

Introduction to NoSQL

NOSQL AND THE CAP THEOREM

16

Cartoon School of Management

CAP Theorem

- Proposed as a conjecture by Eric Brewer (Berkeley, 1999)
 - Subsequently proved by Gilbert and Lynch (MIT) in 2002 as a theorem.
- In a distributed data store you can satisfy **at most 2 out of the 3 guarantees**:
 - Consistency**: Every read receives the most recent write or an error
 - Availability**: Every request receives a (non-error) response, without the guarantee that it contains the most recent write
 - Partition-tolerance**: the system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes
- Proof: When a network partition failure happens, should we
 - Cancel the operation and thus decrease the availability but ensure consistency?
 - Proceed with the operation and thus provide availability but risk inconsistency?

17

Cartoon School of Management

Fox&Brewer "CAP Theorem":
C-A-P: choose two.

Consistency C

Availability A

Partition-resilience P

Claim: every distributed system is on one side of the triangle.

CA: available, and consistent, unless there is a partition.
Do not scale well!!

CP: always consistent, even in a partition, but a reachable replica may deny service without agreement of the others (e.g., quorum).
Not always available

AP: a reachable replica provides service even in a partition, but may be inconsistent if there is a failure.
Not always consistent

18

18

Availability

- Traditionally, thought of as the server/process available five 9's (99.999 %).
- However, for large node system, at almost any point in time there's a good chance that a node is either down or there is a network disruption among the nodes.
 - Want a system that is resilient in the face of network disruption

19

Consistency Model

- A consistency model determines rules for visibility and apparent order of updates.
- For example:
 - Row X is replicated on nodes M and N
 - Client A writes row X to node N
 - Some period of time t elapses.
 - Client B reads row X from node M
 - Does client B see the write from client A?
- Consistency is a continuum with tradeoffs
 - CAP Theorem states: Strict Consistency can't be achieved at the same time as availability and partition-tolerance.

20

Eventual Consistency

- When no updates occur for a long period of time, eventually all updates will propagate through the system and all the nodes will be consistent
- For a given accepted update and a given node, eventually either the update reaches the node or the node is removed from service
- Known as **BASE (Basically Available, Soft state, Eventual consistency)**, as opposed to ACID
 - *Basically Available*: possibilities of faults but not a fault of the whole system
 - *Soft state*: copies of a data item may be inconsistent
 - *Eventually Consistent*: copies becomes consistent at some later time if there are no more updates to that data item

21

Cartoon School of Management

Implications of BASE consistency model

- Characteristics
 - Weak consistency – stale data OK
 - Availability first
 - Best effort
 - Approximate answers OK
 - Aggressive (optimistic)
 - Simpler and faster

22

Cartoon School of Management

Introduction to NoSQL

DIFFERENT KINDS OF NOSQL DATA STORES

23

Cartoon School of Management

NoSQL System Classification

- two common criteria

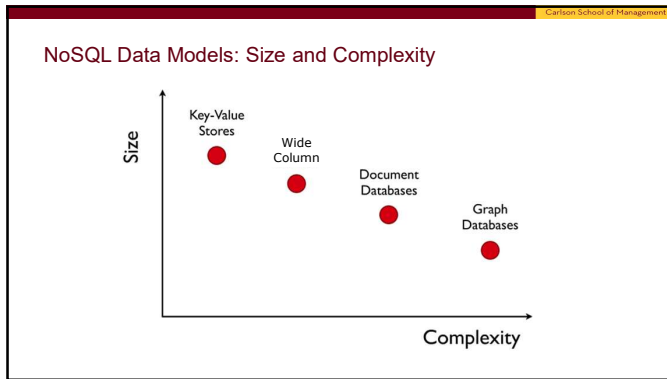
Data Model

- Key-Value
- Wide-Column
- Document
- Graph

Consistency/Availability Trade-Off

- AP:** Available & Partition Tolerant
- CP:** Consistent & Partition Tolerant
- CA:** Not Partition Tolerant

24



25

Cartoon School of Management

Introduction to NoSQL

KEY-VALUE STORES

26

Cartoon School of Management

Key-Value Stores

- Data model: (key) → value
- Interface: CRUD (Create, Read, Update, Delete)

Key

Value: An opaque blob

- Examples: Amazon Dynamo (AP), Riak (AP), Redis (CP)

High availability but not always consistent

Strict consistency but not always available during partitions

27

Key-Value Store Pros and Cons

- Pros:
 - very fast
 - very scalable
 - simple model
 - able to distribute horizontally
- Cons:
 - many data structures (objects) can't be easily modeled as key value pairs
 - Only primary index: lookup by key

28

Key-value Store Use Cases

- Key-value data stores are ideal for storing user profiles, blog comments, product recommendations (user preferences), session information, shopping cart data.
 - Twitter uses **Redis** to deliver your [Twitter timeline](#)
 - [Pinterest](#) uses **Redis** to store lists of users, followers, unfollowers, boards, and more
 - [Quora](#) uses **Memcached** to cache results from slower, persistent databases
- avoid if we need to
 - query the database by specific data value
 - need relationships between data values
 - need to operate on multiple unique keys

29

Introduction to NoSQL

WIDE-COLUMN STORES

30

Carleton School of Management

Wide-Column Stores

- Data model: (rowkey, column, timestamp) -> value
 - can be interpreted as two/three dimensional key-value store
 - You can group related columns into **column families**.
- Interface: CRUD, Scan

- Examples: Cassandra (AP), Google BigTable (CP), HBase (CP)
 - High availability but not always consistent
 - Strict consistency but not always available during partitions

31

Carleton School of Management

Wide-column Store Advantages & Use Cases

- Column stores offer very high performance and a highly scalable architecture. They are fast to load and query
 - Ideal for real-time data logging, analysis, and random read/write access of huge amounts of data.
 - But does not support transactions
- Popular among companies and organizations dealing with big data, IoT, and user recommendation and personalization engines.
 - [Spotify](#) uses **Cassandra** to store user profile attributes and metadata about artists, songs, etc. for better personalization
 - [Facebook](#) built its revamped Messages on top of **HBase**, but is now also used for Nearby Friends feature and search
 - [Outbrain](#) uses **Cassandra** to serve over 190 billion personalized content recommendations each month

32

Carleton School of Management


Introduction to NoSQL

DOCUMENT STORES

33

Document Stores

- Data model: (collection, key) -> document (JSON, XML, BSON)
 - can index common columns for faster performance
- Interface: CRUD, Queries, Map-Reduce



- Examples: CouchDB (AP), RethinkDB (CP), MongoDB (CP)

34

Document Databases Advantages & Use Cases

- Can store data with flexible and evolving schema, scale out well, and perform fast ad hoc queries using secondary indexes.
- Good for content management systems, blogging platforms, analytics platforms, e-commerce platforms
 - SEGA uses cloud hosted MongoDB for handling 11 million in-game accounts
 - Aer Lingus uses MongoDB with Studio 3T to handle ticketing and internal apps
 - Built on MongoDB, The Weather Channel's iOS and Android apps deliver weather alerts to 40 million users in real-time
- avoid if one must run complex search queries or require complex transactions

35


Introduction to NoSQL

GRAPH DATABASES

36

Graph Databases


- Some data are more naturally represented as a graph



37

Graph Databases

- Data model: $G = (V, E)$: Graph-Property Model
- Interface: Traversal algorithms, queries, transactions



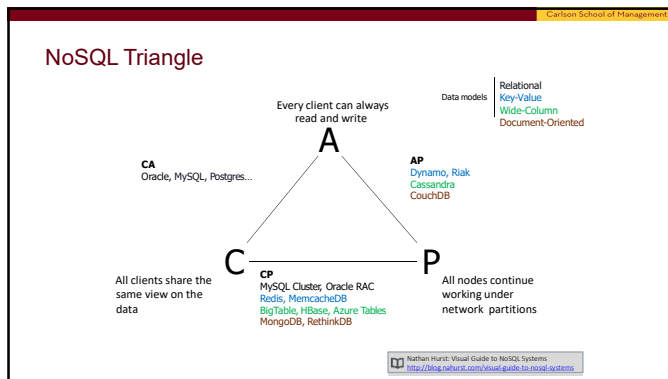
- Examples: Neo4j (CA), InfiniteGraph (CA), OrientDB (CA)

38


Graph Database Advantages and Use Cases

- Graph databases are great for establishing data relationships especially when dealing with large data sets.
 - Fault detection, real-time recommendation engines, master data management, network and IT operations, identity and access management.
- They offer blazing fast query performance and flexibility that relational databases cannot compete with, especially as data grows much deeper.
 - [Walmart](#) uses Neo4j to provide customers personalized, relevant product recommendations and promotions
 - [Medium](#) uses Neo4j to build their social graph to enhance content personalization
 - [Cisco](#) uses Neo4j to mine customer support cases to anticipate bugs

39



40



41

Cartoon School of Management

NewSQL

- NewSQL refers to a new class of databases providing transactional consistency and massive scalability
 - Transactional support of SQL + NoSQL massive scalability
 - Some guarantee stronger consistency, others (MemSQL) provide tunable consistency
 - NewSQL favors consistency over availability (CP)
 - They are cloud-first databases
 - lock-free concurrency control + share nothing architecture
- Examples
 - SAP HANA, Nuodb, MemSQL, VoltDB, Google Spanner, CockroachDB

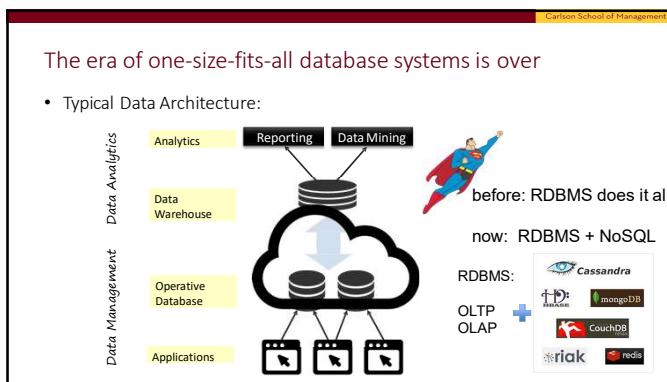
42

Cartoon School of Management

Introduction to NoSQL

HOW TO CHOOSE?

43



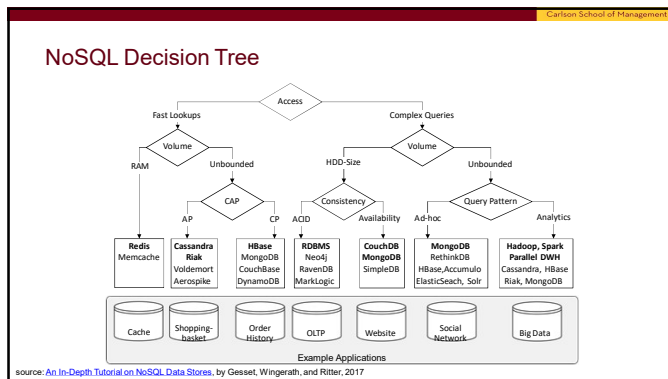
44

Cartoon School of Management

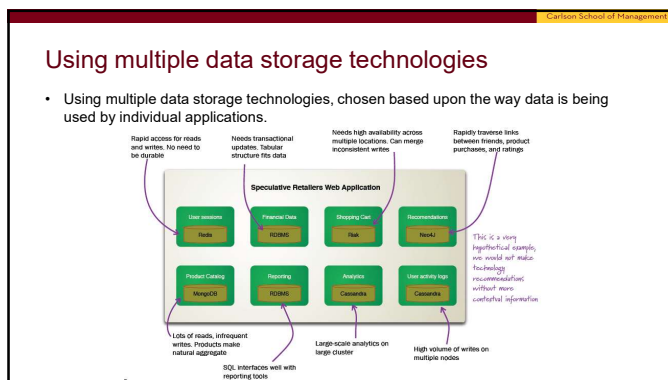
Which database to choose? Why?

- Large part of the choice depends on the business use-case and tradeoffs between scale, consistency, and availability
 - The CAP Theorem
- It also depends on the type of queries you plan to perform on the data?
 - Are they simple look ups or more complex queries?
 - Different SQL/NoSQL has different
- One need also take into account what is the cost of not having ACID?

45



46



47

NoSQL - Summary

- NoSQL was initially motivated by inability of traditional RDBMS in dealing with web scale data request
- NoSQL databases reject:
 - Overhead of ACID transactions
 - "Complexity" of SQL
 - Burden of up-front schema design
 - Declarative query expression
- But CAP theorem dictates that there is no one size-fit all solution
 - NewSQL embraces ACID and scalability at the cost of availability
- Choices of form of databases depend on business use-case and nature of data

48